

The background of the slide features a complex, abstract network graph. It consists of numerous small, glowing green and yellow dots representing nodes, connected by a dense web of thin, translucent lines representing edges. The graph is set against a dark, almost black, background, creating a sense of depth and connectivity. The overall aesthetic is modern and technological.

# THE CUDA C++ DEVELOPER'S TOOLBOX

Bryce Adelstein Lelbach

Principal Architect

✉ bryce@lelbach.com  @blelbach



The background of the slide features a complex, abstract network graph. It consists of numerous small, glowing green and yellow dots representing nodes, connected by a dense web of thin, translucent lines representing edges. The graph is highly interconnected, with clusters of nodes appearing in various parts of the frame. The overall effect is one of data flow, computation, or connectivity.

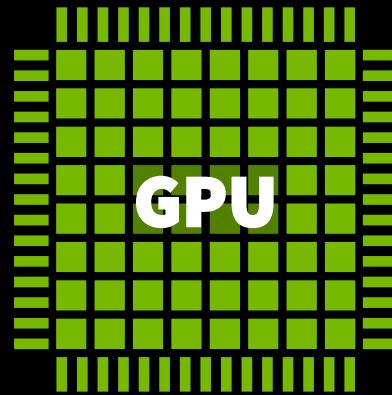
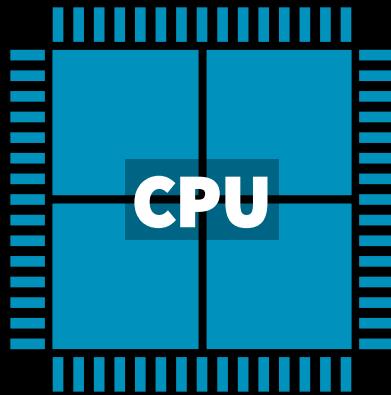
# THE CUDA C++ DEVELOPER'S TOOLBOX

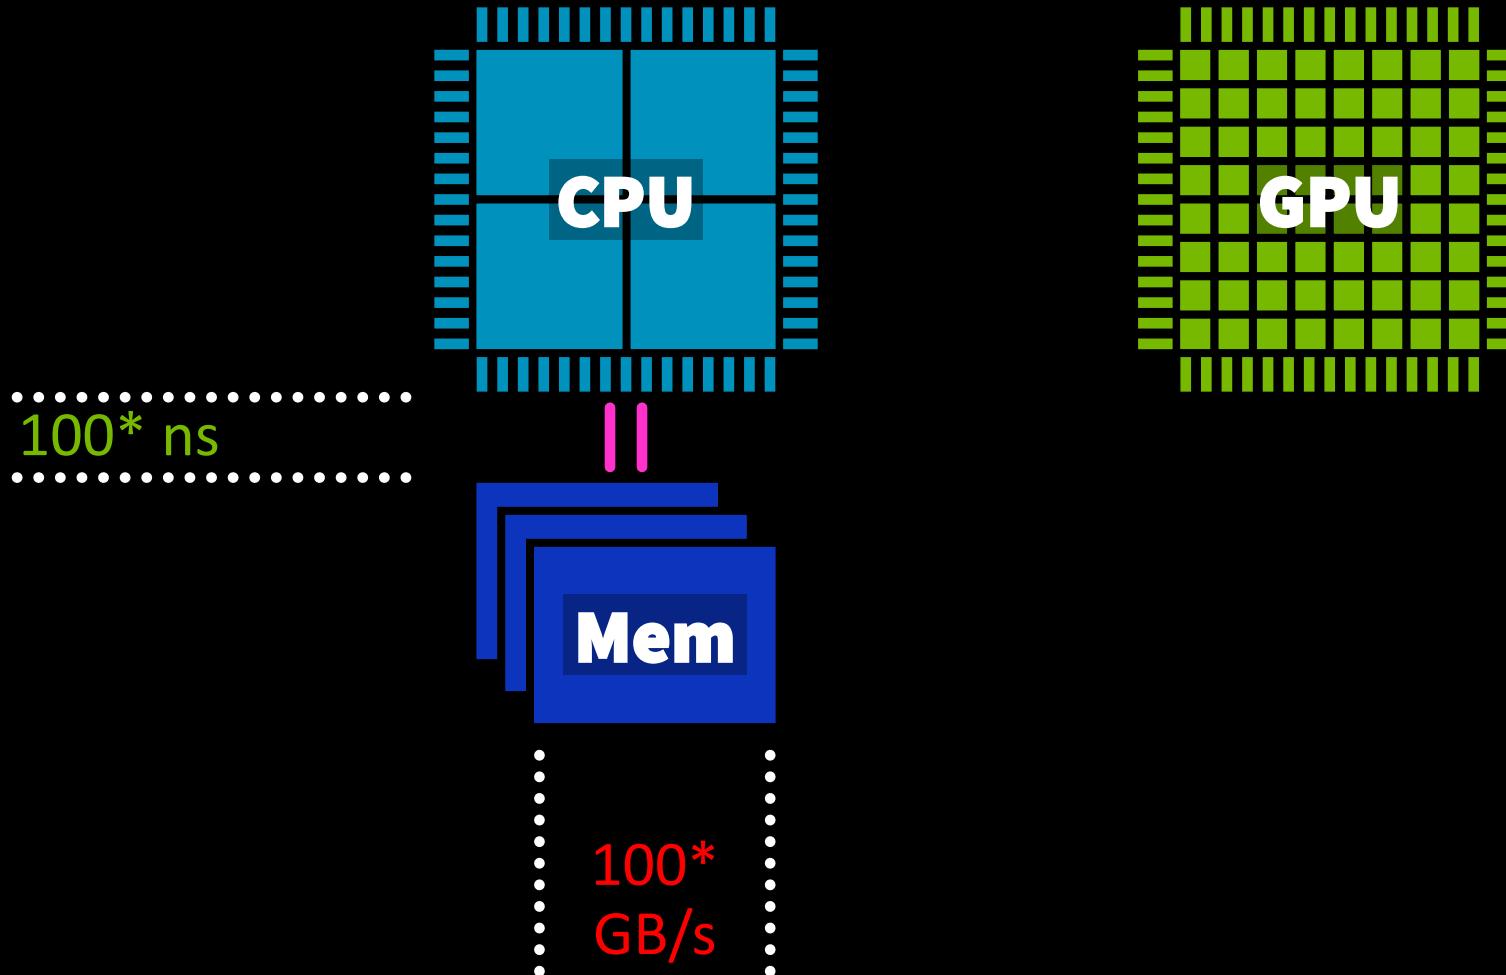
Bryce Adelstein Lelbach

Principal Architect

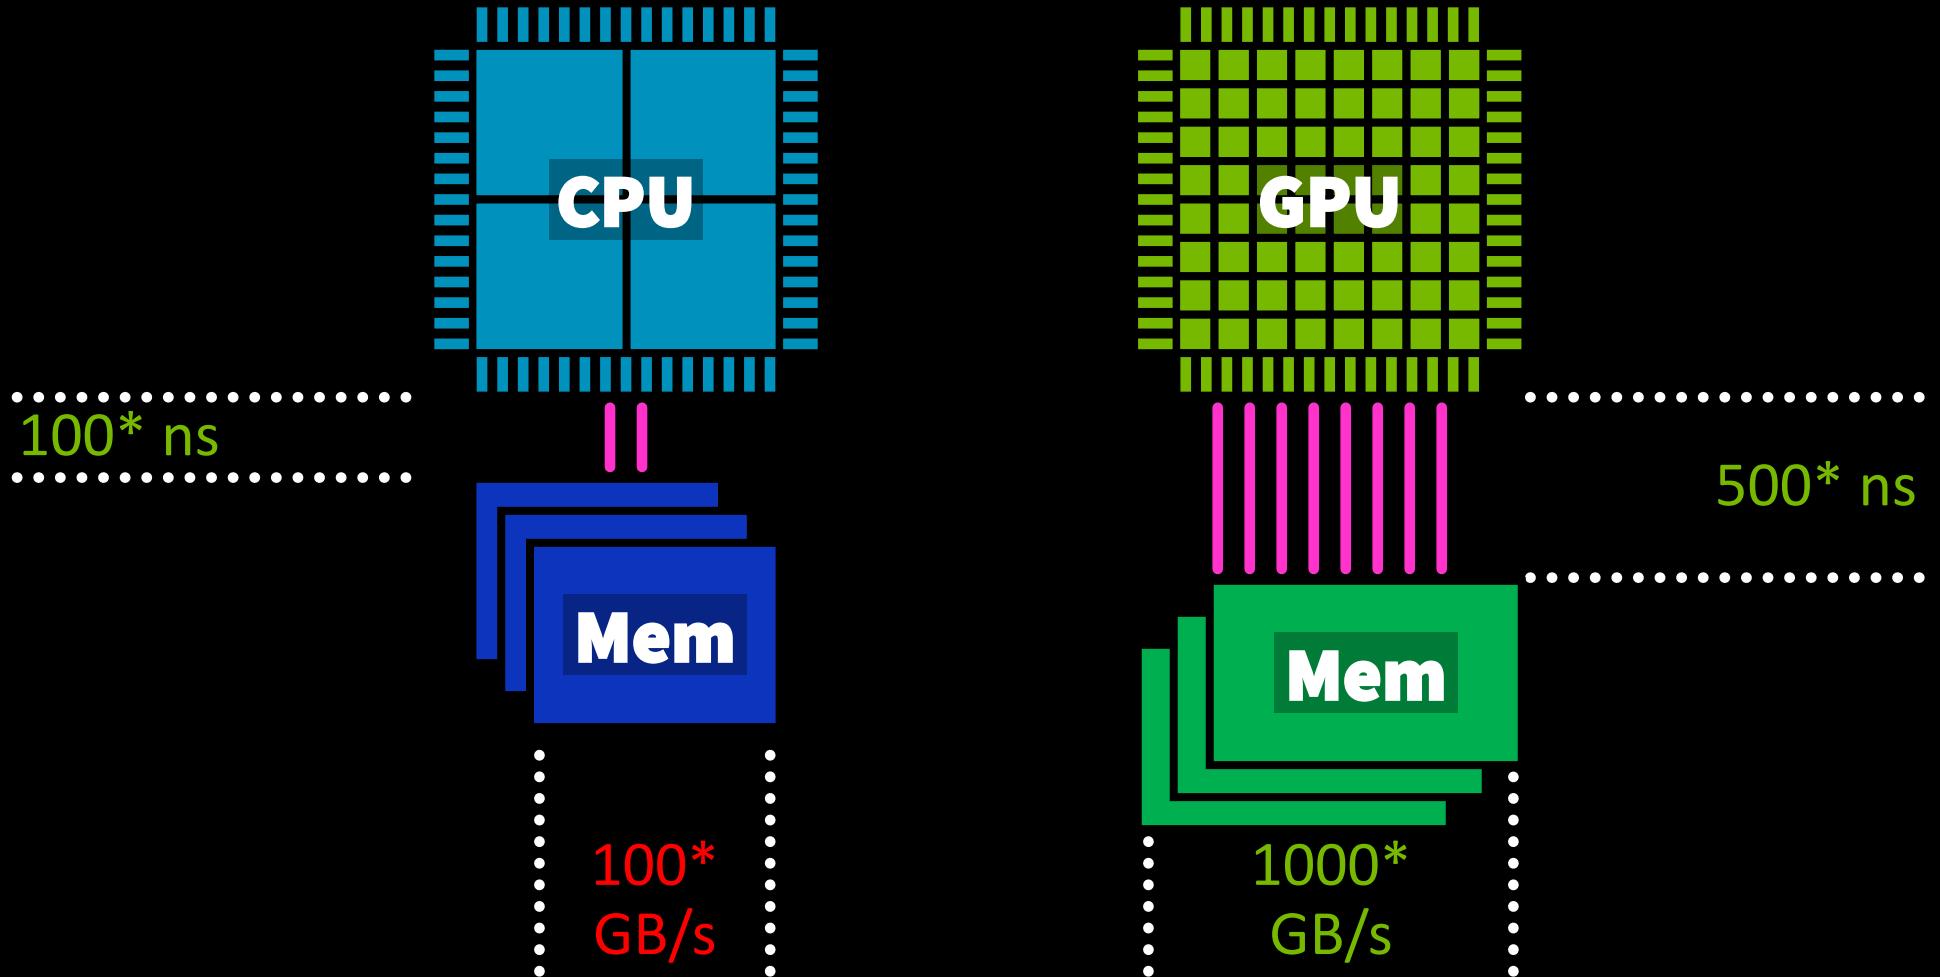
✉ bryce@lelbach.com  @blelbach







\* Numbers are made up and for expository purposes only.



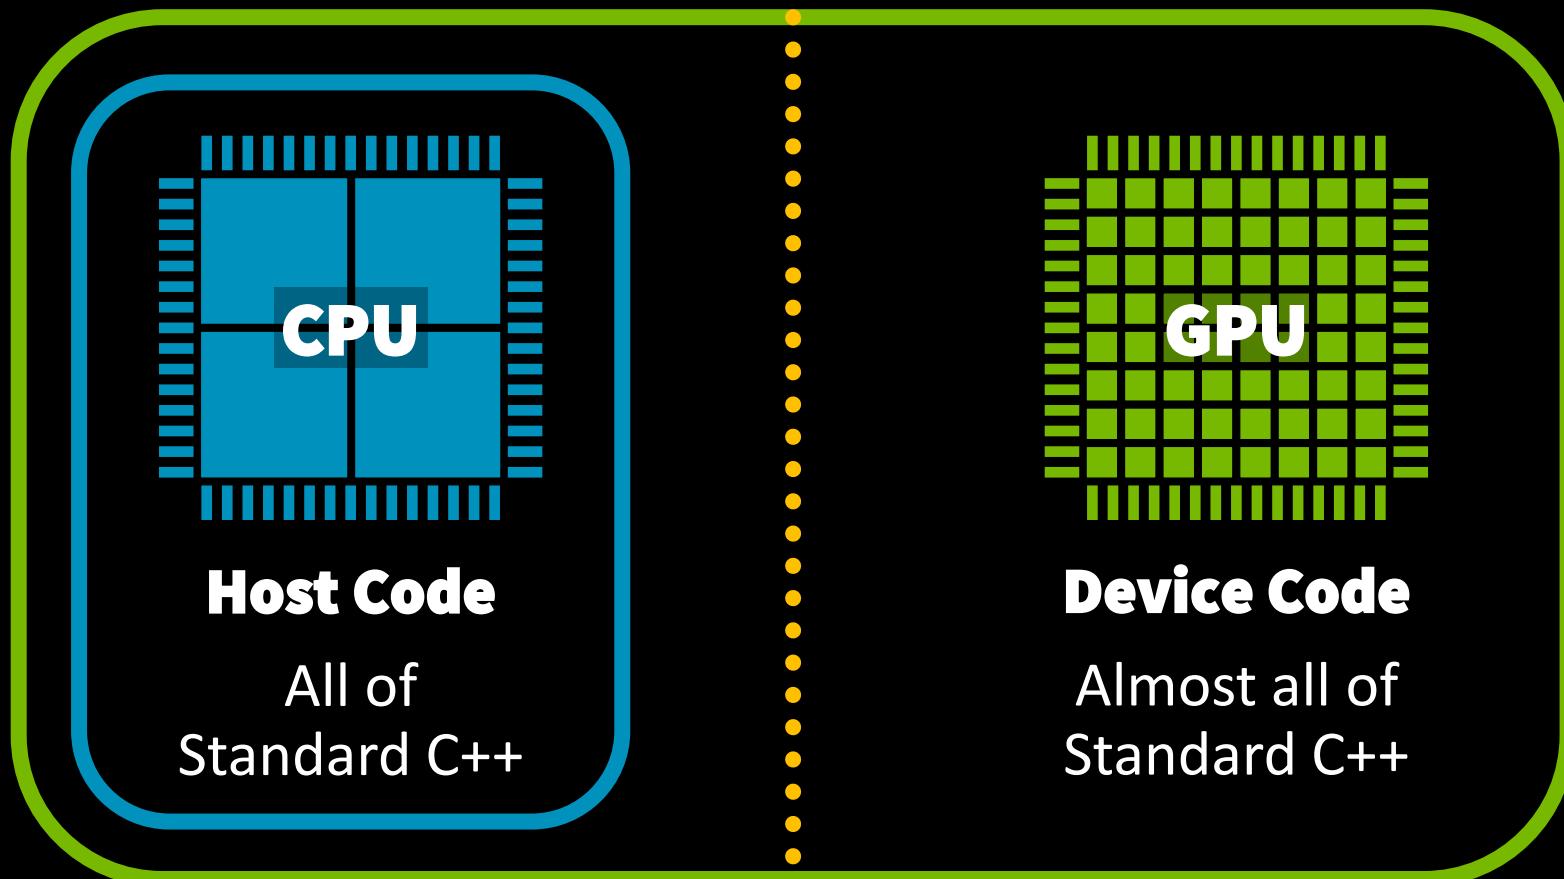
\* Numbers are made up and for expository purposes only.

# How do we program a GPU in C++?

# **How do we program a GPU in C++?**

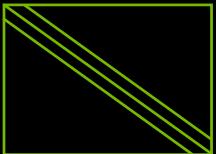
## **CUDA C++ and Accelerated Libraries**

**CUDA C++ is an extension of Standard C++ for writing programs that simultaneously run on CPUs (hosts) & GPUs (devices).**





**cuBLAS**



**cuSPARSE**



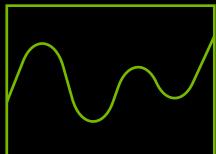
**cuDNN**



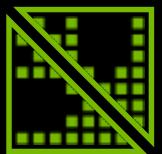
**cuFile  
(GDS)**



**NVBENCH**

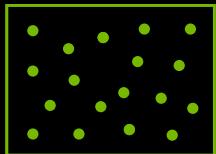


**cuFFT**

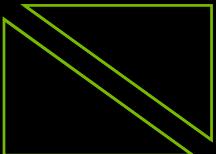


**cuDSS**

**CV  
CUDA**



**cuRAND**

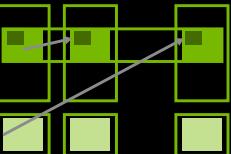


**cuSOLVER**

**RAPIDS**

**cuLitho**

**Cooperative  
Groups**

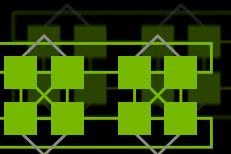


**NVSHMEM**

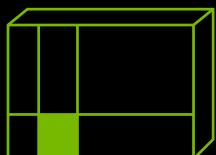


**CUDA-Q**

**CUTLASS**



**NCCL**



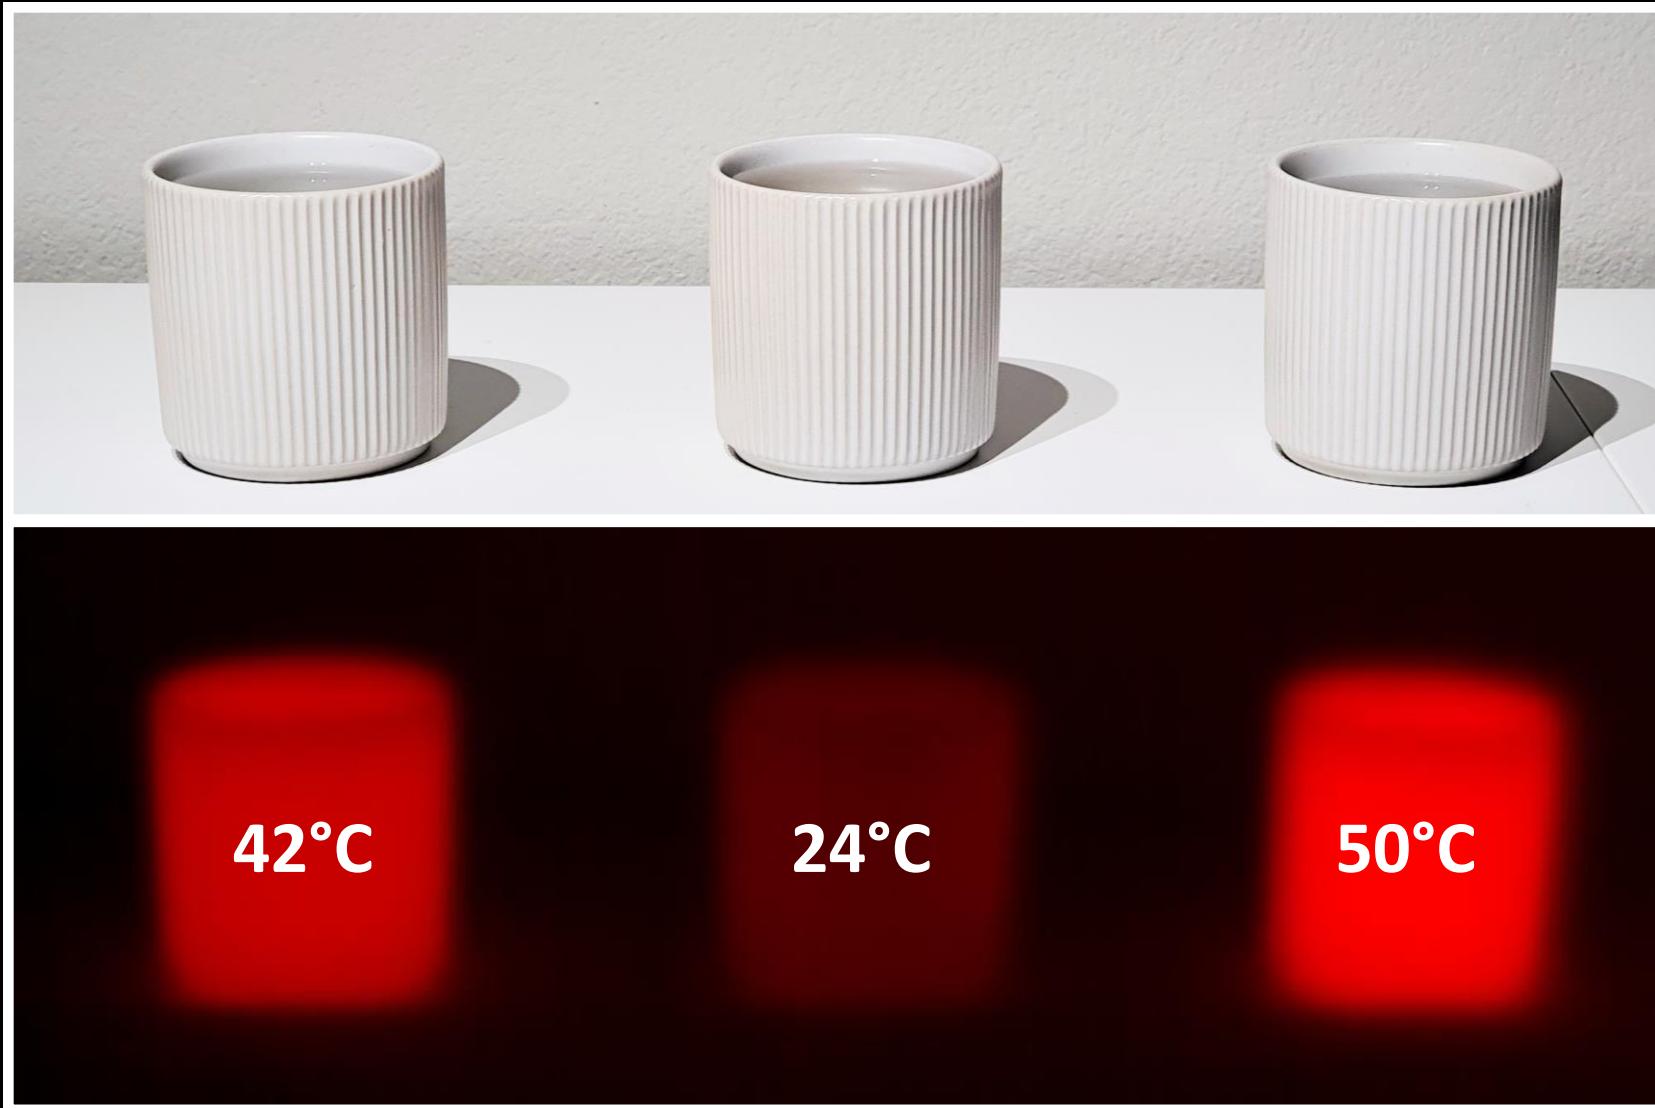
**cuTENSOR**



**CUDA Math**

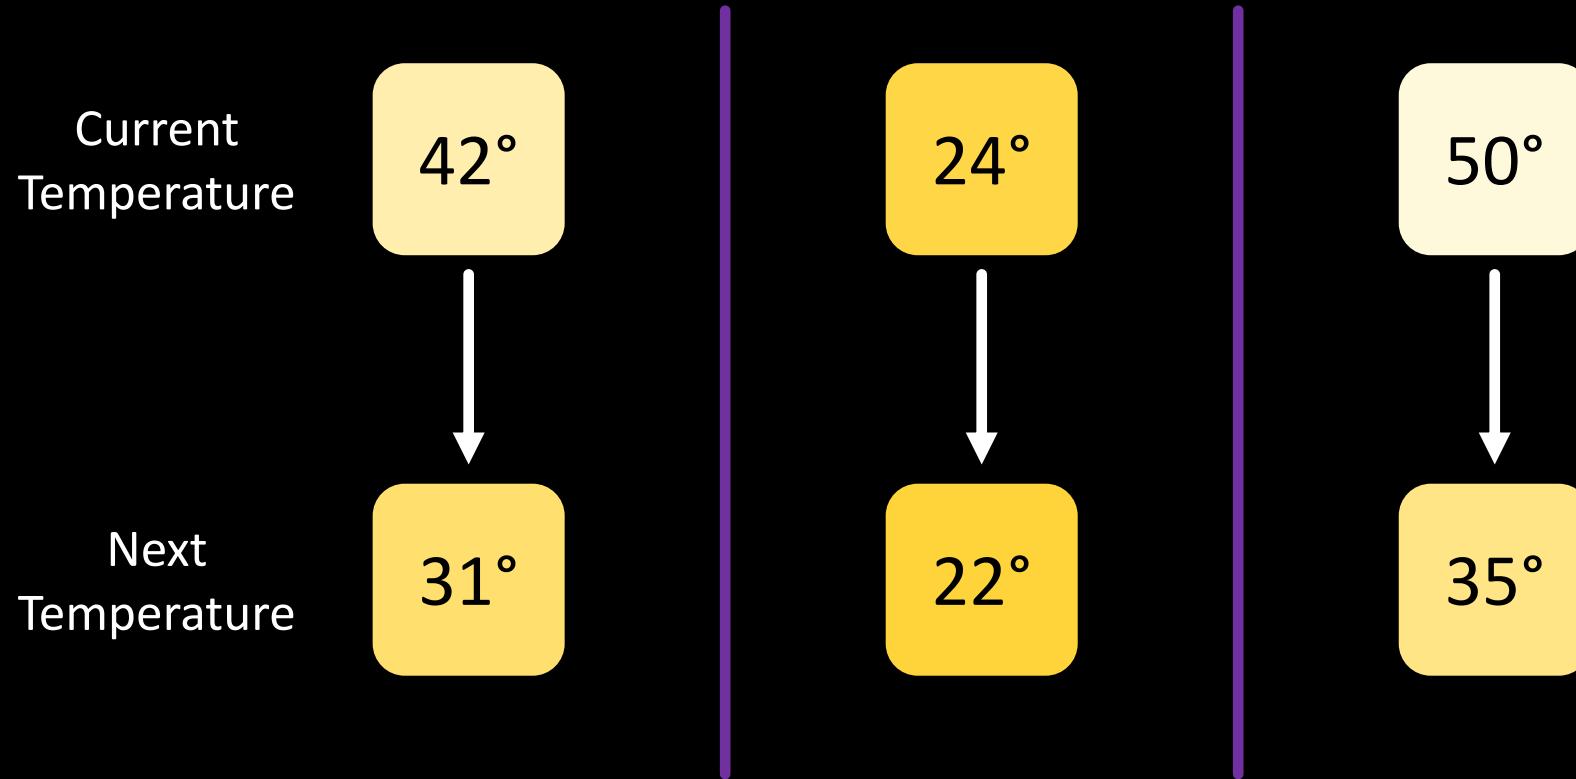


**CUDA  
Runtime**



Ambient Temperature =  $20^{\circ}\text{C}$

Heat Transfer Coefficient ( $k$ ) = 0.5



$$\text{Next Temp} = \text{Current Temp} + \text{Heat Transfer Coefficient} * (\text{Ambient Temp} - \text{Current Temp})$$

```
int steps = 3;
float k = 0.5;
float ambient_temp = 20;
std::vector<float> cups{42, 24, 50};
```

```
int steps = 3;
float k = 0.5;
float ambient_temp = 20;
std::vector<float> cups{42, 24, 50};

auto op = [=] (float t) {
    float diff = ambient_temp - t;
    return t + k * diff;
};
```

```
int steps = 3;
float k = 0.5;
float ambient_temp = 20;
std::vector<float> cups{42, 24, 50};

auto op = [=] (float t) {
    float diff = ambient_temp - t;
    return t + k * diff;
};

for (int step : std::views::iota(0, steps))
{
    ...
}
```

```

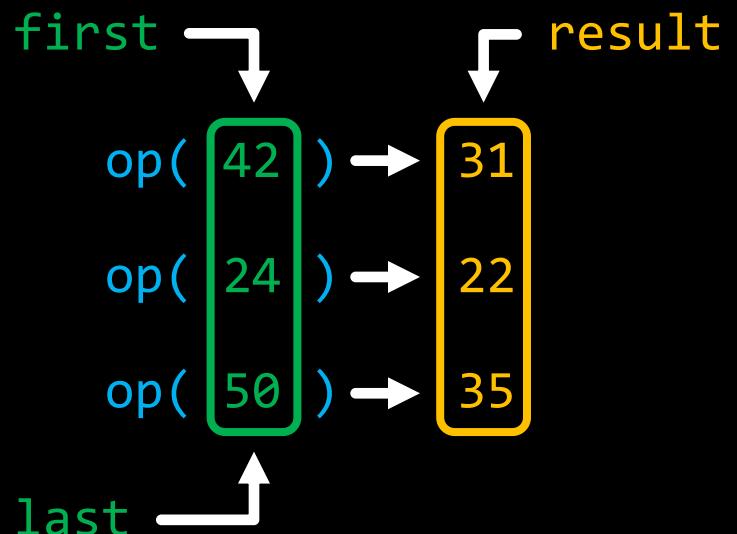
int steps = 3;
float k = 0.5;
float ambient_temp = 20;
std::vector<float> cups{42, 24, 50};

auto op = [=] (float t) {
    float diff = ambient_temp - t;
    return t + k * diff;
};

for (int step : std::views::iota(0, steps))
{
    std::println("{} {}", step, cups);
    std::transform(cups.begin(), cups.end(),
                  cups.begin(), op);
}

```

`std::transform(first, last,  
result, op);`





$t + k * \text{diff}$



```
g++ main.cpp -o a.out
```

Host code  
executable  
by CPU

vmla.f32

**CUDA  
C++**

$t + k * \text{diff}$

**NVCC**

```
nvcc main.cpp -o a.out
```

Host code  
executable  
by CPU

vmla.f32

fma.rn.f32

Device code  
executable  
by GPU

```
int steps = 3;
float k = 0.5;
float ambient_temp = 20;
std::vector<float> cups{42, 24, 50};

auto op = [=] (float t) {
    float diff = ambient_temp - t;
    return t + k * diff;
};

for (int step : std::views::iota(0, steps))
{
    std::println("{} {}", step, cups);
    std::transform(cups.begin(), cups.end(),
                  cups.begin(), op);
}
```

```
int steps = 3;
float k = 0.5;
float ambient_temp = 20;
std::vector<float> cups{42, 24, 50};

auto op = [=] __host__ __device__ (float t) {
    float diff = ambient_temp - t;
    return t + k * diff;
};

for (int step : std::views::iota(0, steps))
{
    std::println("{} {}", step, cups);
    std::transform(cups.begin(), cups.end(),
                  cups.begin(), op);
}
```

```
int steps = 3;
float k = 0.5;
float ambient_temp = 20;
std::vector<float> cups{42, 24, 50};

auto op = [=] __host__ __device__ (float t) {
    float diff = ambient_temp - t;
    return t + k * diff;
};

for (int step : std::views::iota(0, steps))
{
    std::println("{} {}", step, cups);
    std::transform(cups.begin(), cups.end(),
                  cups.begin(), op);
}
```

```
void a();
__host__
void b();
```

Only executable by the CPU.  
Includes all unannotated functions.

```
int steps = 3;
float k = 0.5;
float ambient_temp = 20;
std::vector<float> cups{42, 24, 50};

auto op = [=] __host__ __device__ (float t) {
    float diff = ambient_temp - t;
    return t + k * diff;
};

for (int step : std::views::iota(0, steps))
{
    std::println("{} {}", step, cups);
    std::transform(cups.begin(), cups.end(),
                  cups.begin(), op);
}
```

void a();  
\_\_host\_\_  
void b();

Only executable  
by the CPU.  
Includes all  
unannotated  
functions.

\_\_device\_\_  
void c();

Only executable  
by the GPU.

```

int steps = 3;
float k = 0.5;
float ambient_temp = 20;
std::vector<float> cups{42, 24, 50};

auto op = [=] __host__ __device__ (float t) {
    float diff = ambient_temp - t;
    return t + k * diff;
};

for (int step : std::views::iota(0, steps))
{
    std::println("{} {}", step, cups);
    std::transform(cups.begin(), cups.end(),
                  cups.begin(), op);
}

```

`void a();  
__host__  
void b();`

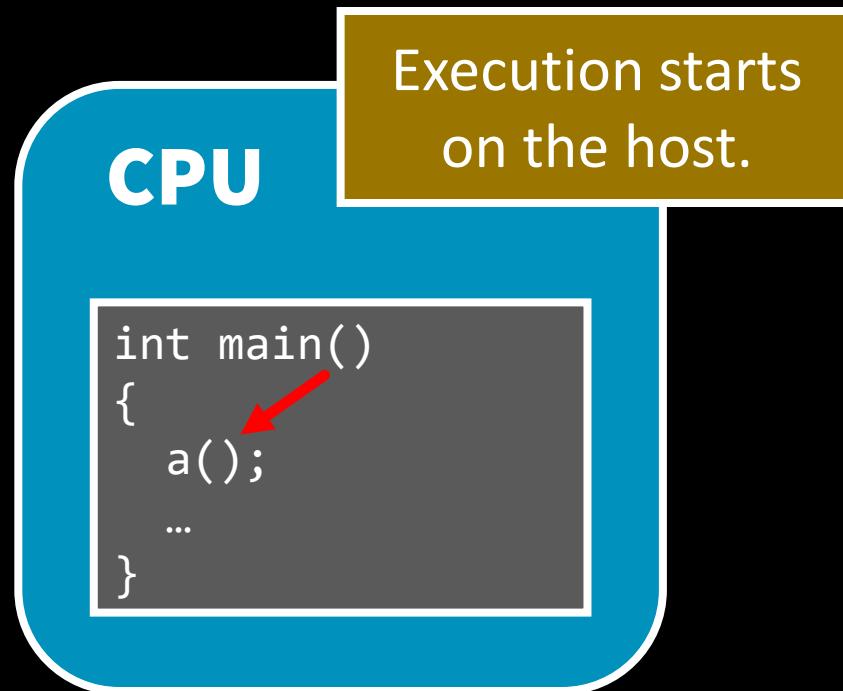
Only executable by the CPU.  
Includes all unannotated functions.

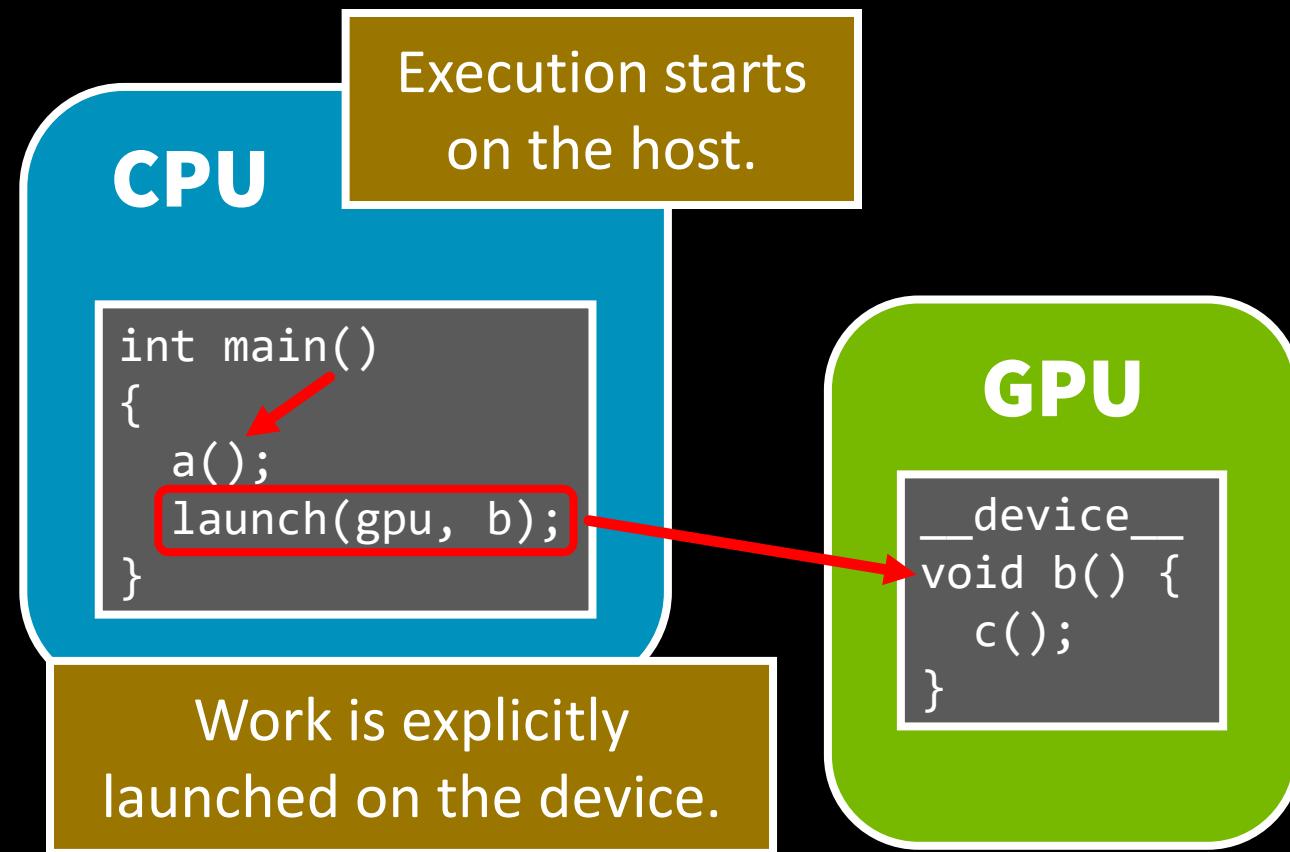
`__device__  
void c();`

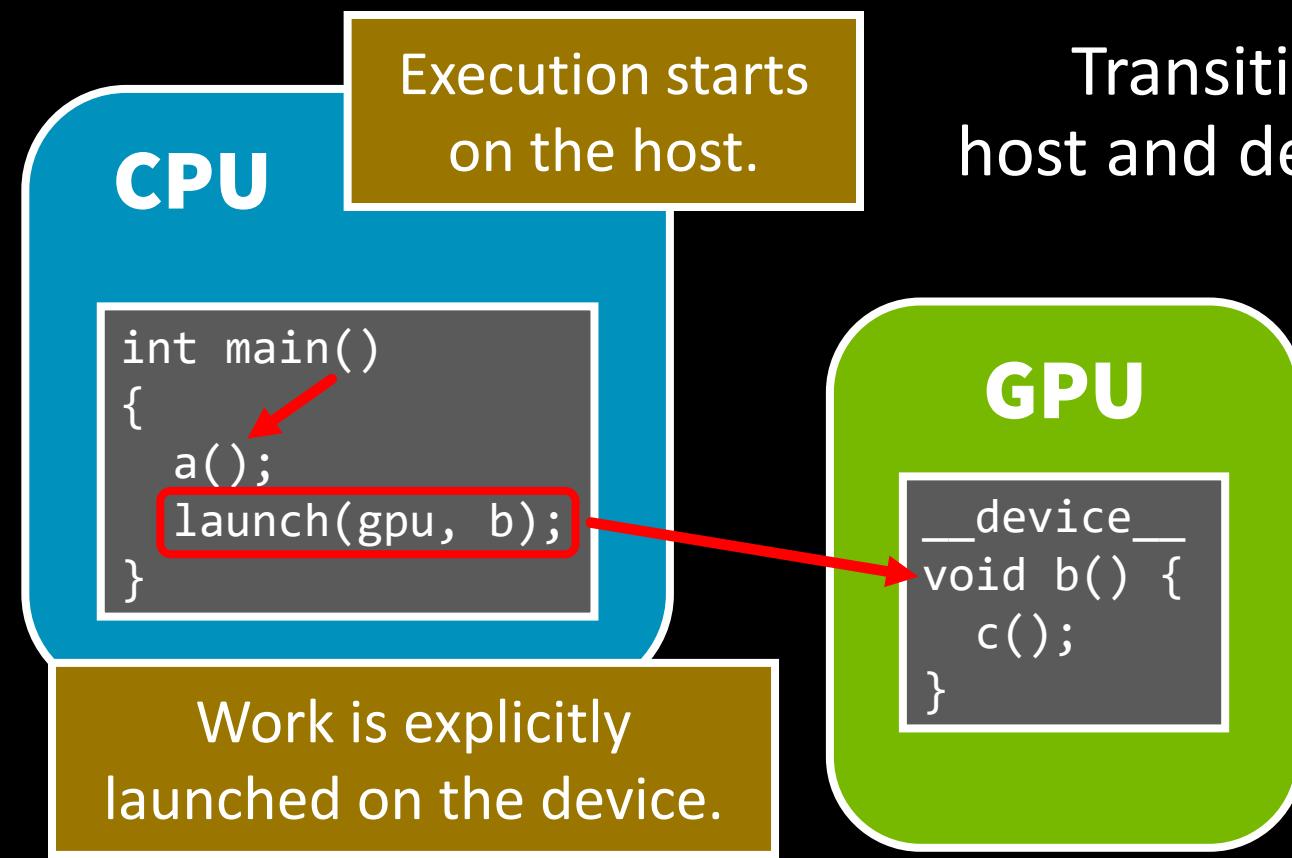
Only executable by the GPU.

`__host__  
__device__  
void d();`

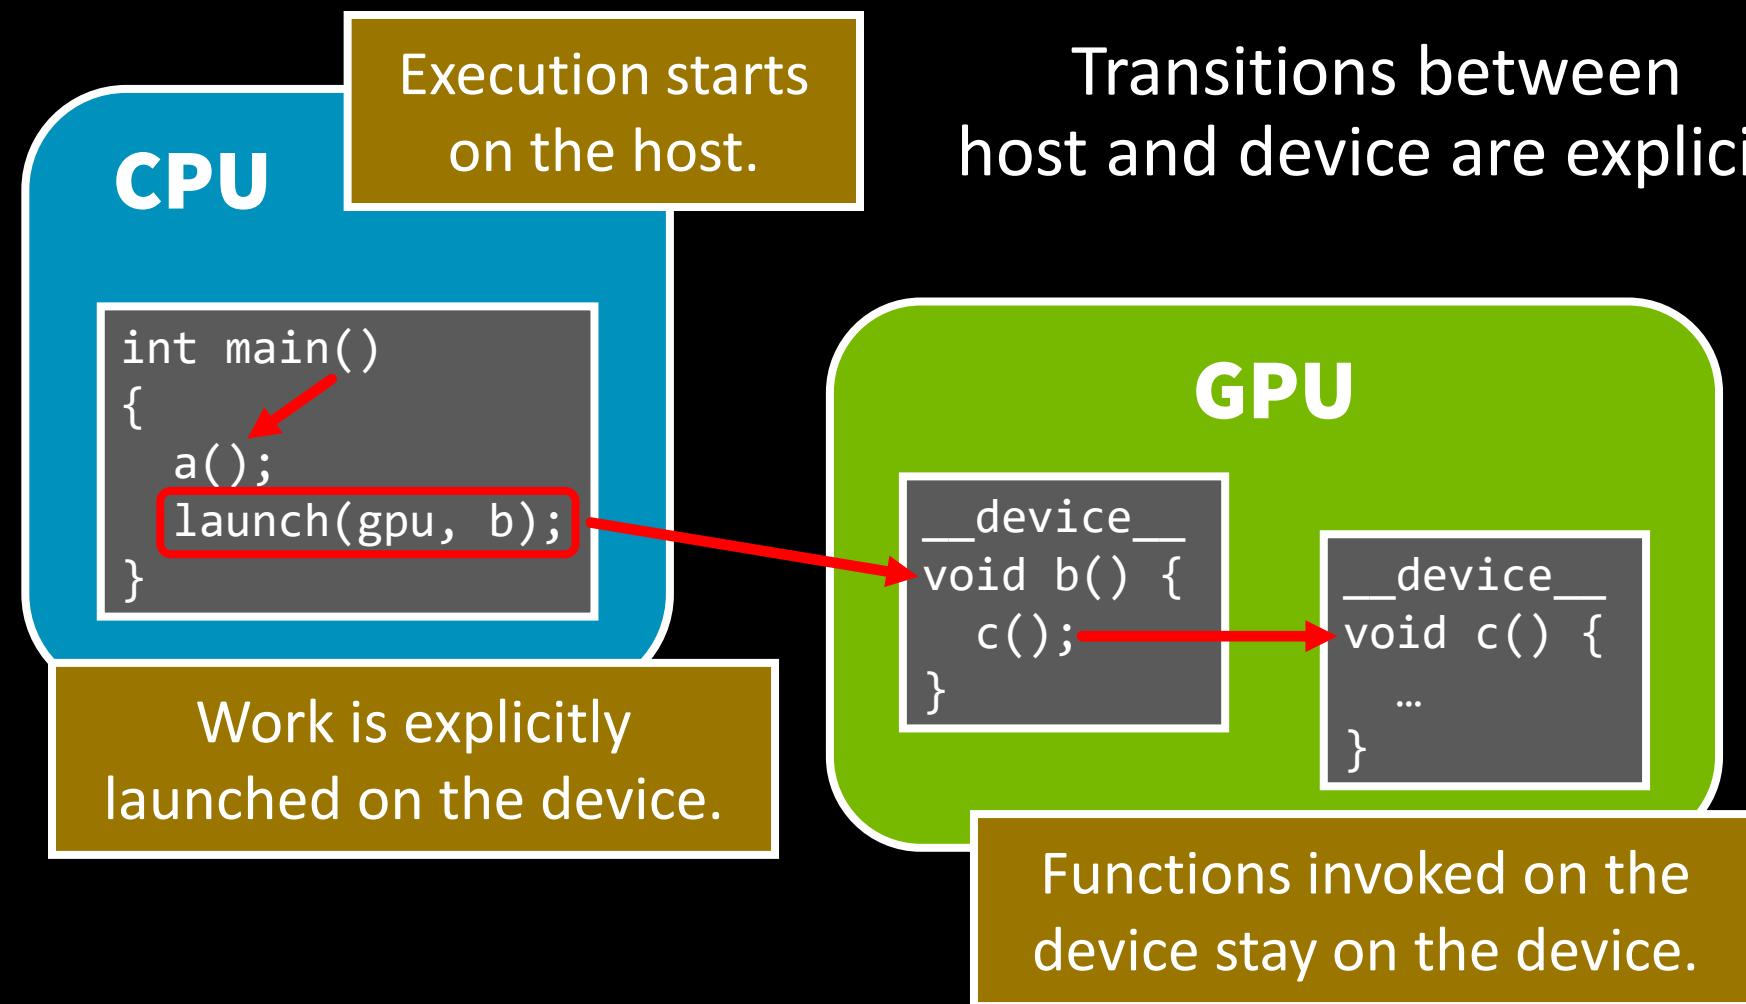
Executable by CPU and GPU.







Transitions between host and device are explicit.



```
int steps = 3;
float k = 0.5;
float ambient_temp = 20;
std::vector<float> cups{42, 24, 50};

auto op = [=] __host__ __device__ (float t) {
    float diff = ambient_temp - t;
    return t + k * diff;
};

for (int step : std::views::iota(0, steps))
{
    std::println("{} {}", step, cups);
    std::transform(cups.begin(), cups.end(),
                  cups.begin(), op);
}
```

```
int steps = 3;
float k = 0.5;
float ambient_temp = 20;
std::vector<float> cups{42, 24, 50};

auto op = [=] __host__ __device__ (float t) {
    float diff = ambient_temp - t;
    return t + k * diff;
};

for (int step : std::views::iota(0, steps))
{
    std::println("{} {}", step, cups);
    thrust::transform(thrust::cuda::par, cups.begin(), cups.end(),
                     cups.begin(), op);
}
```

```
int steps = 3;
float k = 0.5;
float ambient_temp = 20;
thrust::universal_vector<float> cups{42, 24, 50};

auto op = [=] __host__ __device__ (float t) {
    float diff = ambient_temp - t;
    return t + k * diff;
};

for (int step : std::views::iota(0, steps))
{
    std::println("{} {}", step, cups);
    thrust::transform(thrust::cuda::par, cups.begin(), cups.end(),
                      cups.begin(), op);
}
```



# The C++ Parallel Algorithms Library

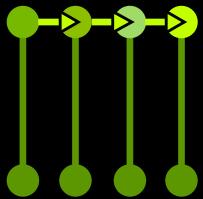
<https://nvidia.github.io/cccl/thrust>



# The C++ Parallel Algorithms Library

<https://nvidia.github.io/cccl/thrust>

## Standard Algorithms



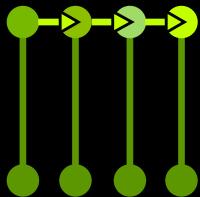
- `thrust::transform_reduce`
- `thrust::inclusive_scan`
- `thrust::sort`
- `thrust::copy`
- ...



# The C++ Parallel Algorithms Library

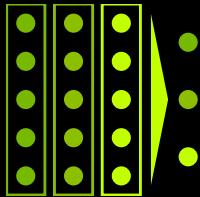
<https://nvidia.github.io/cccl/thrust>

## Standard Algorithms



- `thrust::transform_reduce`
- `thrust::inclusive_scan`
- `thrust::sort`
- `thrust::copy`
- ...

## Extended Algorithms



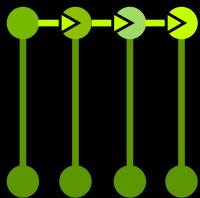
- `thrust::reduce_by_key`
- `thrust::sort_by_key`
- `thrust::tabulate`
- `thrust::gather`
- ...



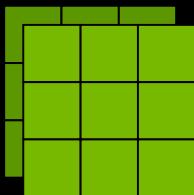
# The C++ Parallel Algorithms Library

<https://nvidia.github.io/cccl/thrust>

## Standard Algorithms



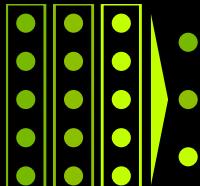
- `thrust::transform_reduce`
- `thrust::inclusive_scan`
- `thrust::sort`
- `thrust::copy`
- ...



## Containers

- `thrust::device_vector`
- `thrust::host_vector`
- `thrust::universal_vector`
- `thrust::allocate_unique`
- ...

## Extended Algorithms



- `thrust::reduce_by_key`
- `thrust::sort_by_key`
- `thrust::tabulate`
- `thrust::gather`
- ...

```
std::vector<T> h0(1 << 30);  
std::vector<T> h1(1 << 30);
```

```
h0 = h1;
```

`std::vector`'s contents might  
only be accessible in *host code*.

Construction and assignment  
is *serial*.

```
std::vector<T> h0(1 << 30);  
std::vector<T> h1(1 << 30);
```

```
h0 = h1;
```

```
thrust::universal_vector<T> u0(1 << 30);  
thrust::universal_vector<T> u1(1 << 30);
```

```
u0 = u1;
```

`std::vector`'s contents might  
only be accessible in *host code*.

Construction and assignment  
is *serial*.

`thrust::universal_vector`'s  
contents are accessible in  
host and device code.

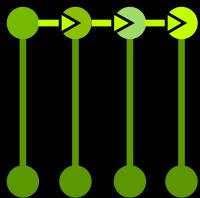
Construction and assignment  
is parallel.



# The C++ Parallel Algorithms Library

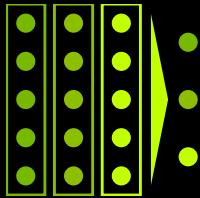
<https://nvidia.github.io/cccl/thrust>

## Standard Algorithms

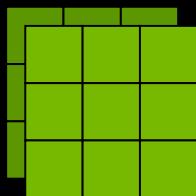


- `thrust::transform_reduce`
- `thrust::inclusive_scan`
- `thrust::sort`
- `thrust::copy`
- ...

## Extended Algorithms

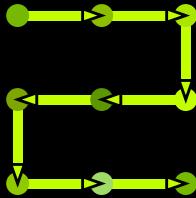


- `thrust::reduce_by_key`
- `thrust::sort_by_key`
- `thrust::tabulate`
- `thrust::gather`
- ...



## Containers

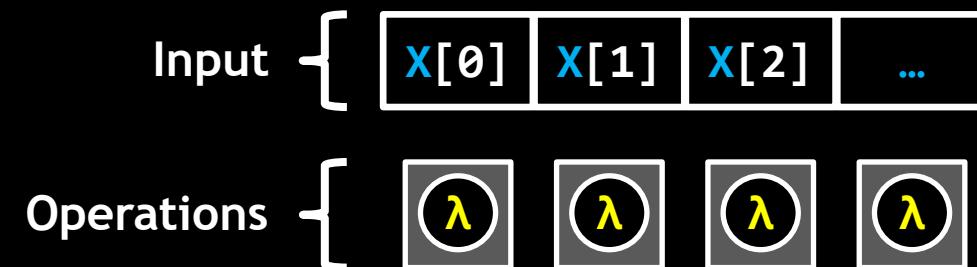
- `thrust::device_vector`
- `thrust::host_vector`
- `thrust::universal_vector`
- `thrust::allocate_unique`
- ...



## Iterators

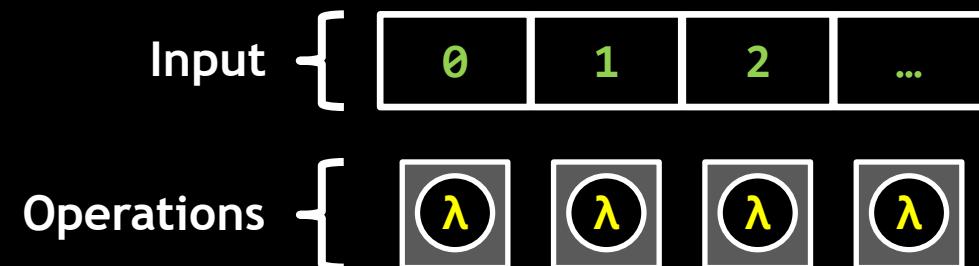
- `thrust::counting_iterator`
- `thrust::constant_iterator`
- `thrust::transform_iterator`
- `thrust::zip_iterator`
- ...

```
thrust::universal_vector X(N);  
  
thrust::for_each_n(  
    thrust::cuda::par, X.begin(), N,  
    [...] __host__ __device__ (auto& obj) { ... });
```



```
auto i = thrust::make_counting_iterator(0);

thrust::for_each_n(
    thrust::cuda::par, i, N,
    [...] __host__ __device__ (auto idx) { ... });
```

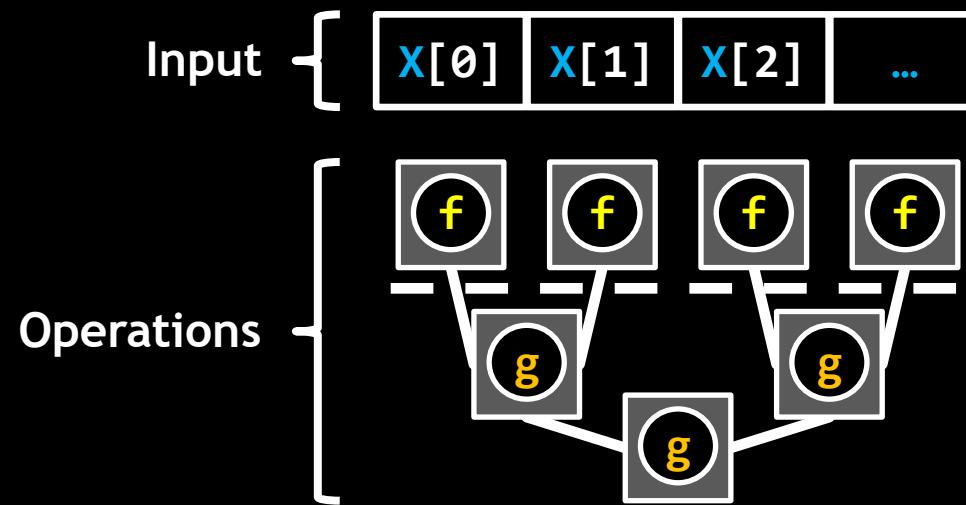


```
auto i = thrust::make_constant_iterator(42);

thrust::for_each_n(
    thrust::cuda::par, i, N,
    [...] __host__ __device__ (auto idx) { ... });
```



```
thrust::universal_vector X(N), tmp(N);  
  
thrust::transform(thrust::cuda::par,  
    X.begin(), X.end(), tmp.begin(), f);  
  
auto r = thrust::reduce(thrust::cuda::par,  
    tmp.begin(), tmp.end(), T{}, g);
```



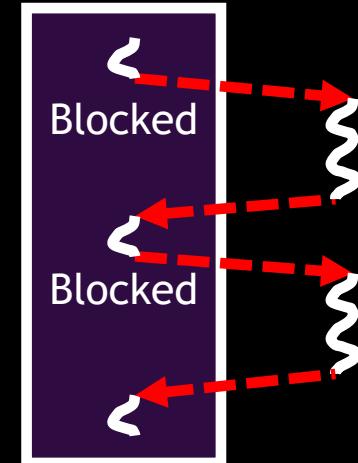
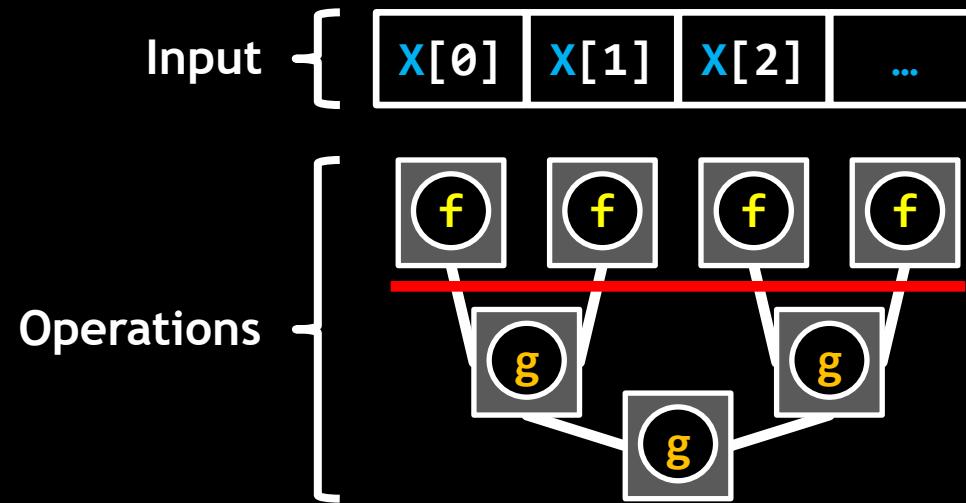
```

thrust::universal_vector X(N), tmp(N);

thrust::transform(thrust::cuda::par,
    X.begin(), X.end(), tmp.begin(), f);

auto r = thrust::reduce(thrust::cuda::par,
    tmp.begin(), tmp.end(), T{}, g);

```

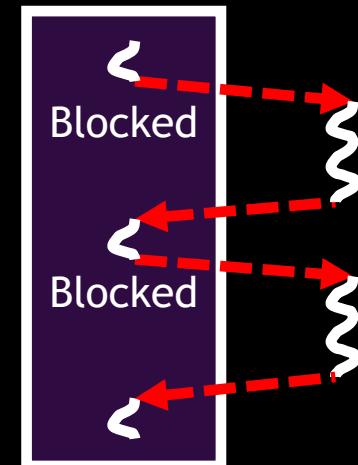
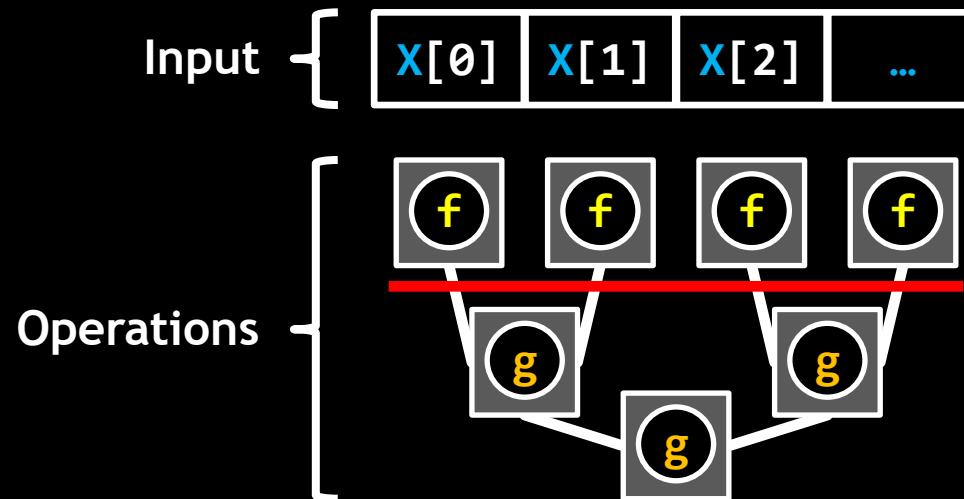


```
thrust::universal_vector X(N), tmp(N);
```

```
thrust::transform(thrust::cuda::par,  
X.begin(), X.end(), tmp.begin(), f);
```

```
auto r = thrust::reduce(thrust::cuda::par,  
tmp.begin(), tmp.end(), T{}, g);
```

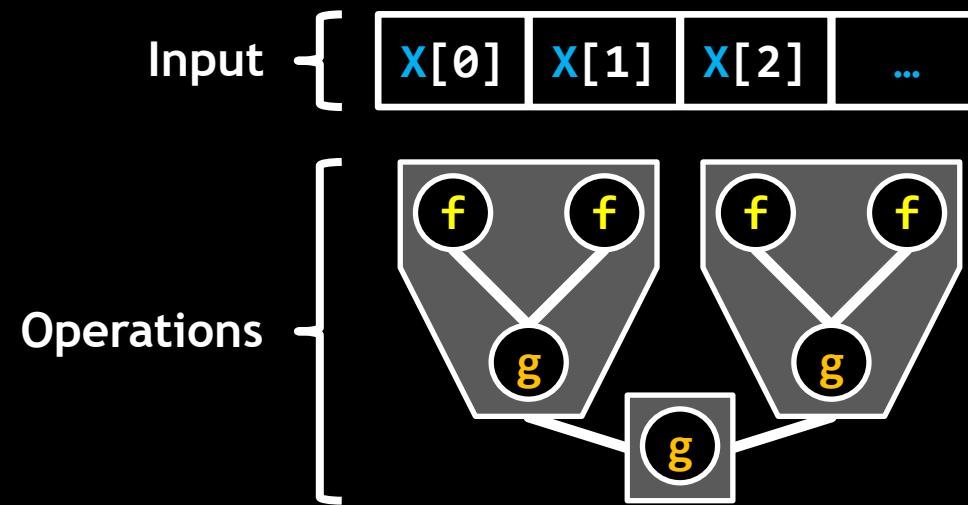
Materializes O(N)  
temporary storage  
in precious device  
memory.



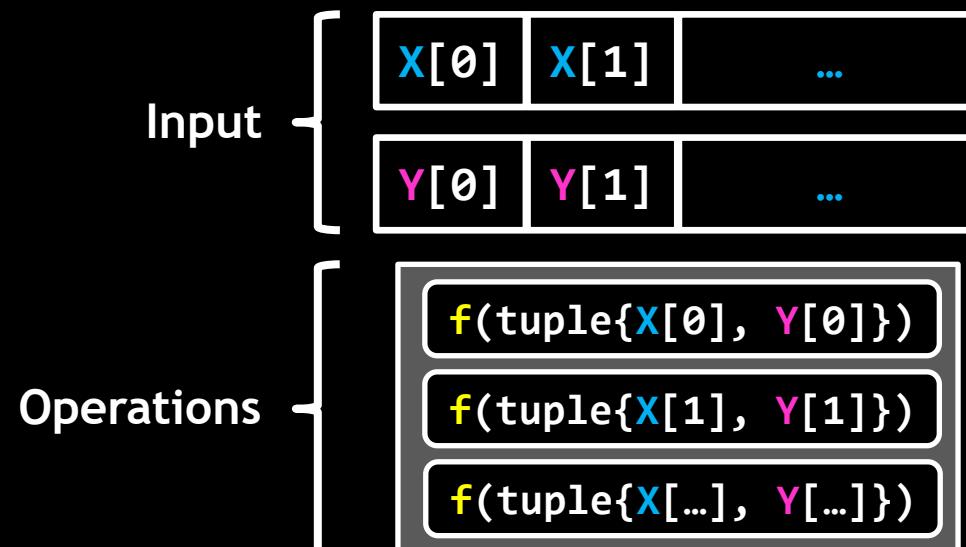
```
thrust::universal_vector X(N);

auto tmp = thrust::make_transform_iterator(
    X.begin(), f);

auto r = thrust::reduce(thrust::cuda::par,
    tmp, tmp + N, T{}, g);
```



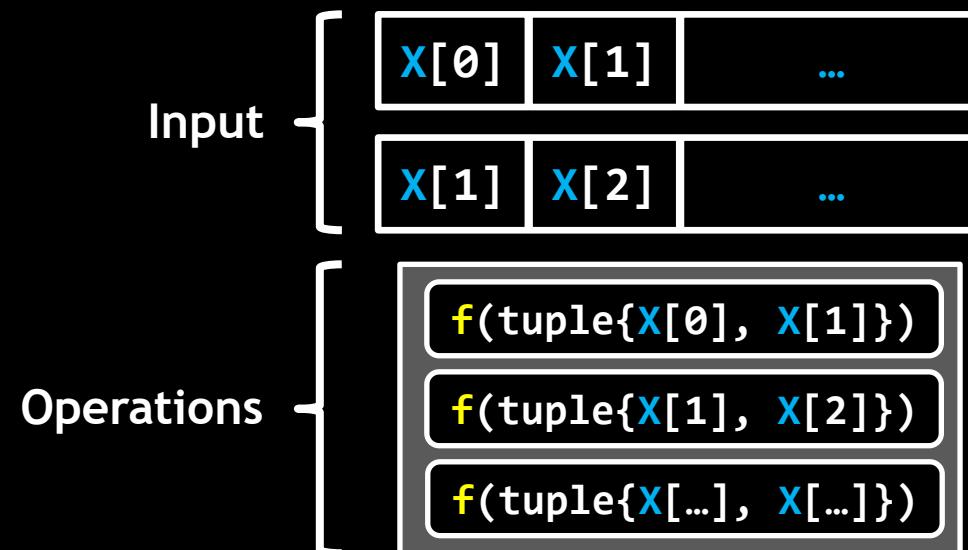
```
thrust::universal_vector X(N), Y(N);  
  
auto XY = thrust::make_zip_iterator(  
    thrust::cuda::par, X.begin(), Y.begin());  
  
thrust::for_each_n(  
    thrust::cuda::par, XY, N, f);
```



```
thrust::universal_vector X(N);

auto adj = thrust::make_zip_iterator(
    thrust::cuda::par, X.begin(), X.begin() + 1);

thrust::for_each_n(
    thrust::cuda::par, adj, N - 1, f);
```



42      31

42      -      31      11

24      22      →      24      -      22      →      2      →      15

50      35

50      -      35      15

A      B

A      -      B      max

```
thrust::universal_vector<int> A(...);  
thrust::universal_vector<int> B(...);
```

0

1

2

3

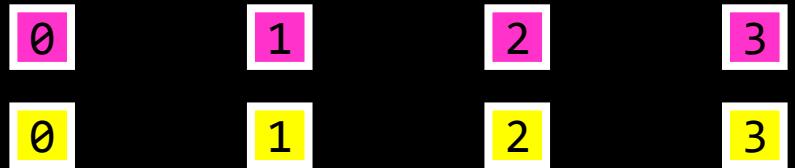
0

1

2

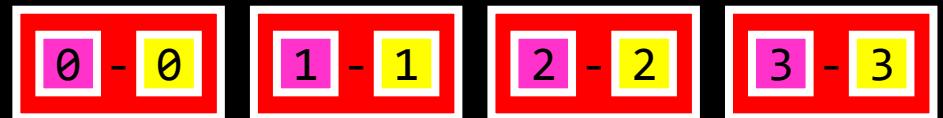
3

```
thrust::universal_vector<int> A(...);  
thrust::universal_vector<int> B(...);
```



```
thrust::universal_vector<int> diffs(A.size());
```

```
thrust::transform(thrust::cuda::par,  
    A.begin(), A.end(), B.begin(), diffs.begin(),  
    [] __host__ __device__ (int a, int b)  
    { return abs(a - b); });
```



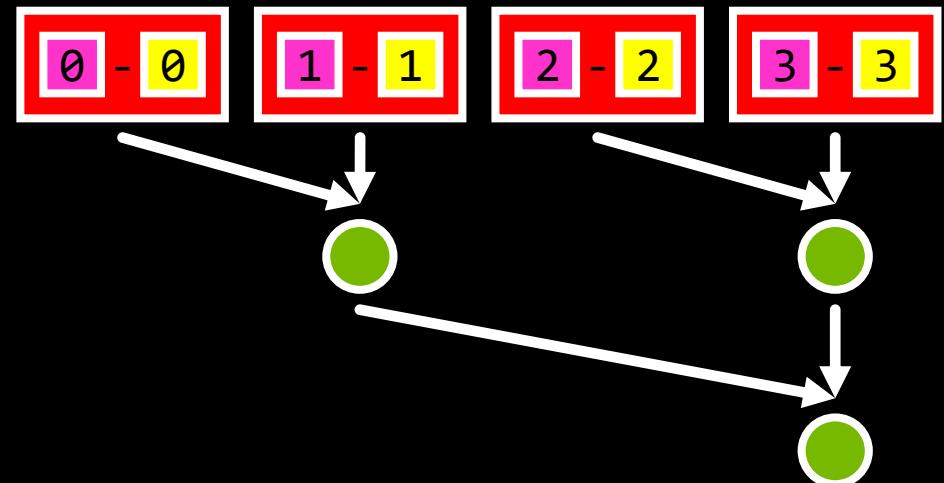
```
thrust::universal_vector<int> A(...);  
thrust::universal_vector<int> B(...);
```



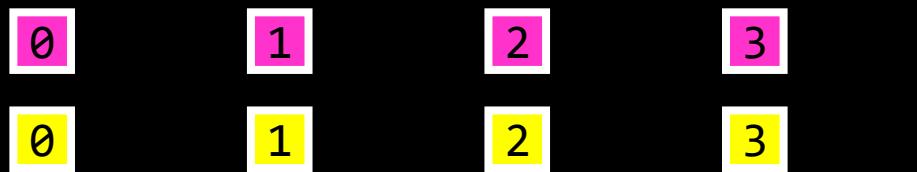
```
thrust::universal_vector<int> diffs(A.size());
```

```
thrust::transform(thrust::cuda::par,  
    A.begin(), A.end(), B.begin(), diffs.begin(),  
    [] __host__ __device__ (int a, int b)  
    { return abs(a - b); });
```

```
auto max_diff = thrust::reduce(  
    thrust::cuda::par,  
    diffs.begin(), diffs.end(),  
    0,  
    cuda::maximum{});
```



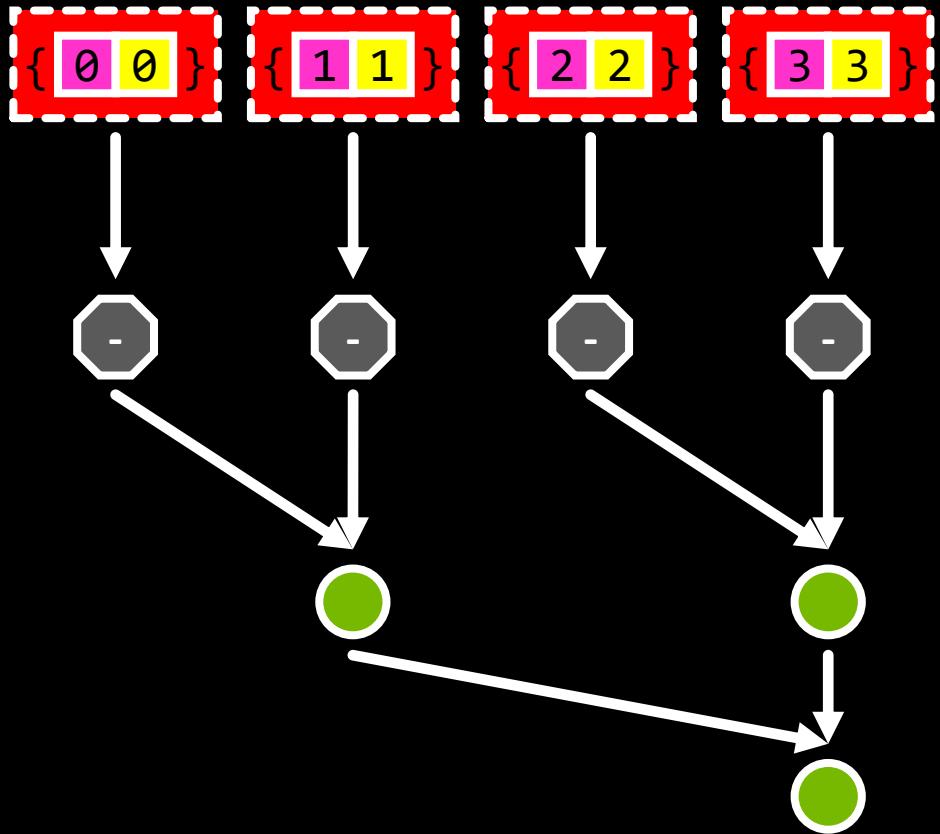
```
thrust::universal_vector<int> A(...);  
thrust::universal_vector<int> B(...);
```



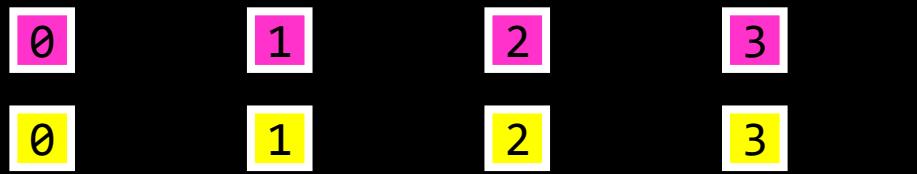
```
auto AB = thrust::make_zip_iterator(  
    A.begin(), B.begin());
```

```
auto diffs = thrust::make_transform_iterator(  
    AB.begin(),  
    [] __host__ __device__  
    (cuda::std::tuple<int, int> ab) {  
        auto [a, b] = ab;  
        return abs(a - b)  
    });
```

```
auto max_diff = thrust::reduce(  
    thrust::cuda::par,  
    diffs, diffs + A.size(),  
    0,  
    cuda::maximum{});
```



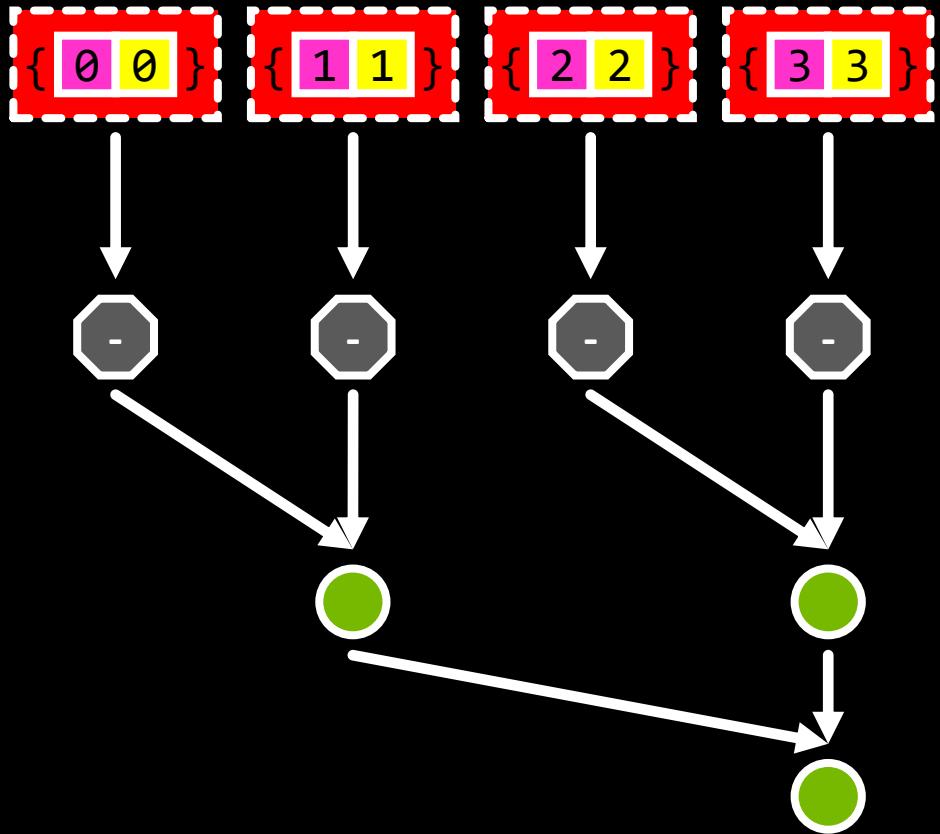
```
thrust::universal_vector<int> A(...);  
thrust::universal_vector<int> B(...);
```



```
auto AB = thrust::make_zip_iterator(  
    A.begin(), B.begin());
```

```
auto diffs = thrust::make_transform_iterator(  
    AB.begin(),  
    [] __host__ __device__  
    (cuda::std::tuple<int, int> ab) {  
        auto [a, b] = ab;  
        return abs(a - b)  
    });
```

```
auto max_diff = thrust::reduce(  
    thrust::cuda::par,  
    diffs, diffs + A.size(),  
    0,  
    cuda::maximum{});
```



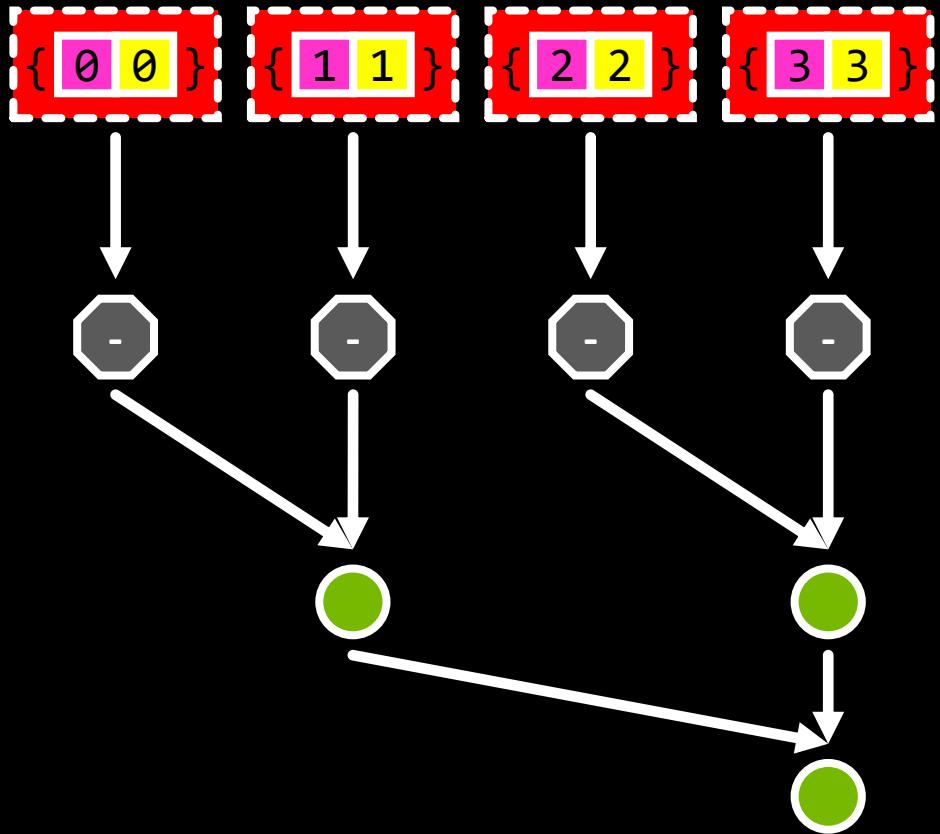
```
thrust::universal_vector<int> A(...);  
thrust::universal_vector<int> B(...);
```



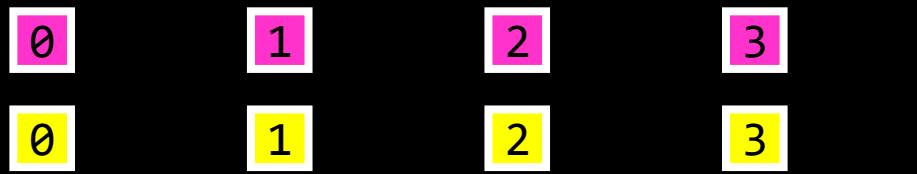
```
auto AB = thrust::make_zip_iterator(  
    A.begin(), B.begin());
```

```
auto diffs = thrust::make_transform_iterator(  
    AB.begin(),  
    [] __host__ __device__  
    (cuda::std::tuple<int, int> ab) {  
        auto [a, b] = ab;  
        return abs(a - b)  
    });
```

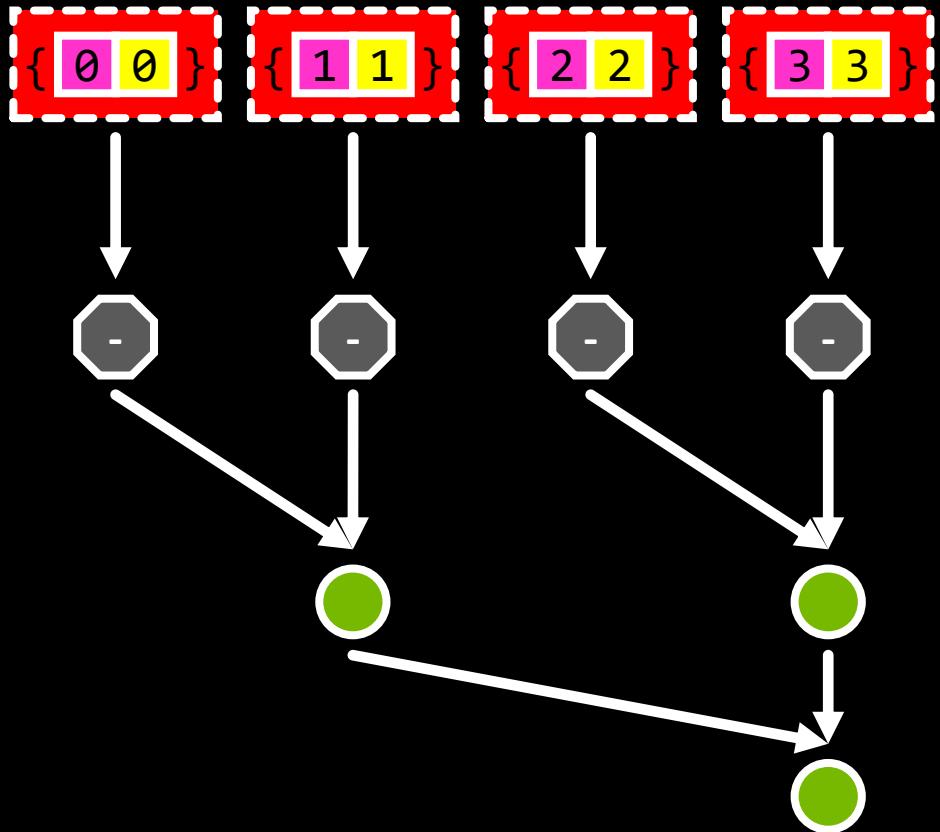
```
auto max_diff = thrust::reduce(  
    thrust::cuda::par,  
    diffs, diffs + A.size(),  
    0,  
    cuda::maximum{});
```



```
thrust::universal_vector<int> A(...);  
thrust::universal_vector<int> B(...);
```



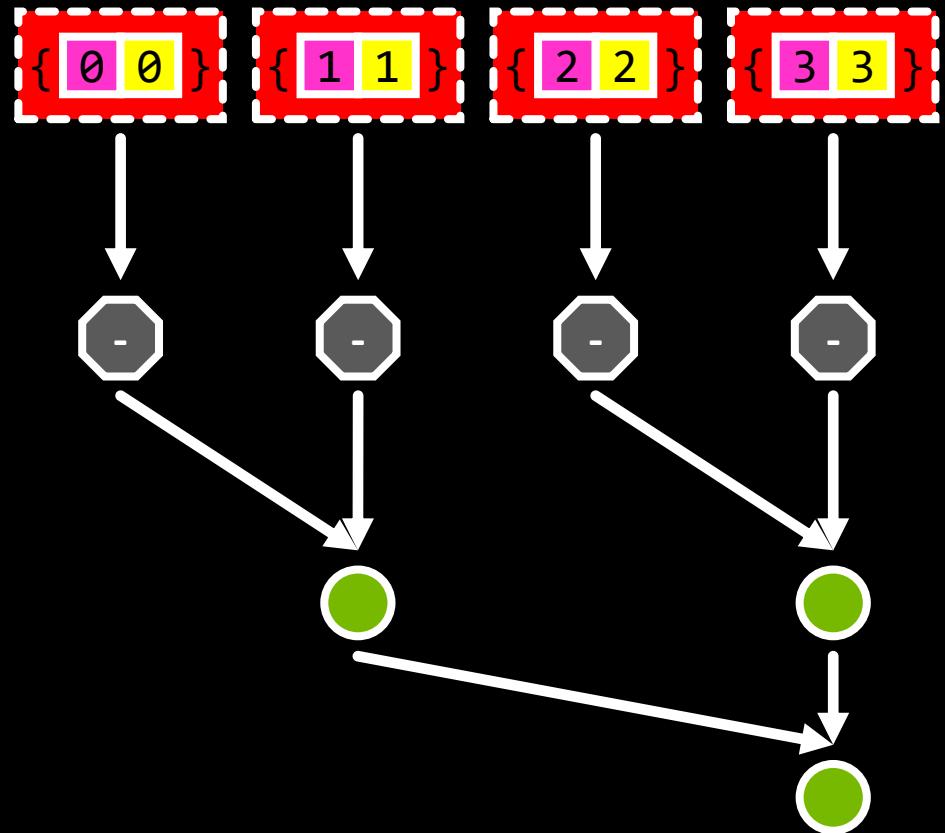
```
auto AB = thrust::make_zip_iterator(  
    A.begin(), B.begin());  
  
auto max_diff = thrust::transform_reduce(  
    thrust::cuda::par,  
    AB, AB + A.size(),  
    [] __host__ __device__  
    (cuda::std::tuple<int, int> ab) {  
        auto [a, b] = ab;  
        return abs(a - b)  
    },  
    0,  
    cuda::maximum{});
```

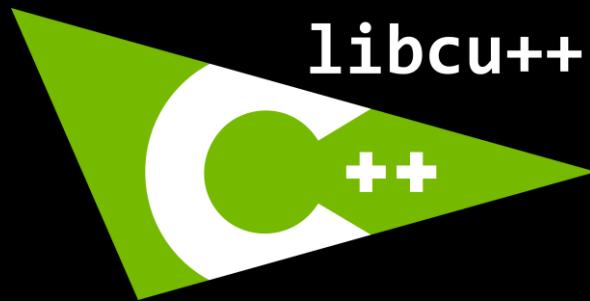


```
thrust::universal_vector<int> A(...);  
thrust::universal_vector<int> B(...);
```



```
auto AB = thrust::make_zip_iterator(  
    A.begin(), B.begin());  
  
auto max_diff = thrust::transform_reduce(  
    thrust::cuda::par,  
    AB, AB + A.size(),  
    [] __host__ __device__  
    (cuda::std::tuple<int, int> ab) {  
        auto [a, b] = ab;  
        return abs(a - b)  
    },  
    0,  
    cuda::maximum{});
```

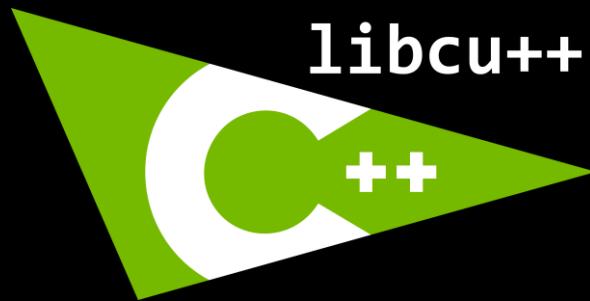




# The CUDA C++ Foundational Library

<https://nvidia.github.io/cccl/libcudacxx>

Learn More: [S72575 How You Should Write a CUDA C++ Kernel](#)



# The CUDA C++ Foundational Library

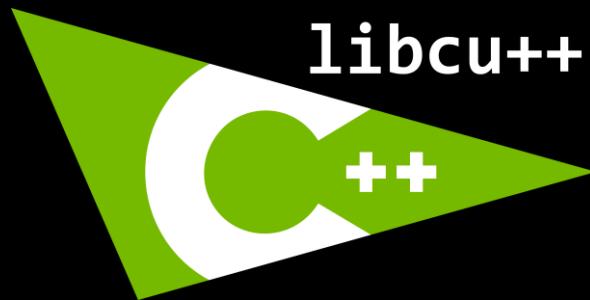
<https://nvidia.github.io/cccl/libcudacxx>

Learn More: [S72575 How You Should Write a CUDA C++ Kernel](#)

## Host Compiler's Standard Library (GCC, LLVM, MSVC, etc)

```
#include <...>
std:::
```

Standard C++, `__host__` only.  
Complete, strictly conforming to Standard C++.



# The CUDA C++ Foundational Library

<https://nvidia.github.io/cccl/libcudacxx>

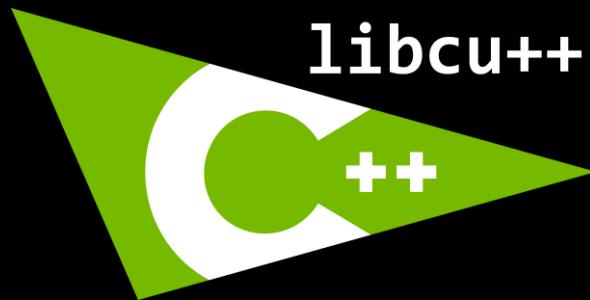
Learn More: [S72575 How You Should Write a CUDA C++ Kernel](#)

## Host Compiler's Standard Library (GCC, LLVM, MSVC, etc)

#include <...> std:::	Standard C++, <code>__host__</code> only. Complete, strictly conforming to Standard C++.
--------------------------	---

#include <cuda/std/> <b>cuda::std:::</b>	CUDA C++, <code>__host__ __device__</code> . Subset, strictly conforming to Standard C++.
---	--

**libcu++**



# The CUDA C++ Foundational Library

<https://nvidia.github.io/cccl/libcudacxx>

Learn More: [S72575 How You Should Write a CUDA C++ Kernel](#)

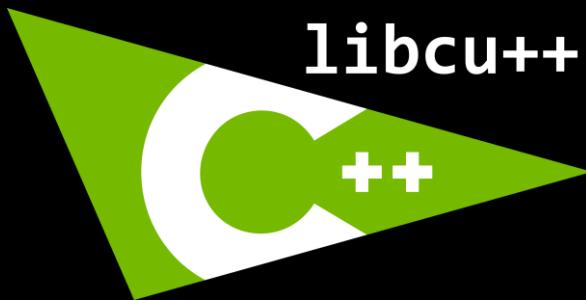
## Host Compiler's Standard Library (GCC, LLVM, MSVC, etc)

#include <...> std:::	Standard C++, <code>__host__</code> only. Complete, strictly conforming to Standard C++.
--------------------------	---

#include <cuda/std/> <b>cuda::std::</b> :	CUDA C++, <code>__host__ __device__</code> . Subset, strictly conforming to Standard C++.
--	--

#include <cuda/> <b>cuda::</b>	CUDA C++, <code>__host__</code> and/or <code>__device__</code> . Modern C++ API for CUDA & Standard C++ extensions.
-----------------------------------	--

**libcu++**



# The CUDA C++ Foundational Library

<https://nvidia.github.io/cccl/libcudacxx>

Learn More: [S72575 How You Should Write a CUDA C++ Kernel](#)

## Host Compiler's Standard Library (GCC, LLVM, MSVC, etc)

```
#include <...>  
std:::
```

Standard C++, `__host__` only.  
Complete, strictly conforming to Standard C++.

```
#include <cuda/std/...>  
cuda::std:::
```

CUDA C++, `__host__ __device__`.  
Subset, strictly conforming to Standard C++.

```
#include <cuda/...>  
cuda:::
```

CUDA C++, `__host__` and/or `__device__`.  
Modern C++ API for CUDA & Standard C++ extensions.

```
#include <cuda/experimental/...>  
cuda::experimental:: (cudax::)
```

CUDA C++, `__host__` and/or `__device__`.  
Beta features.

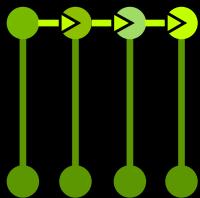
libcu++



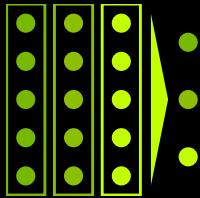
# The C++ Parallel Algorithms Library

<https://nvidia.github.io/cccl/thrust>

## Standard Algorithms

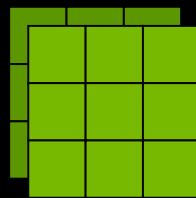


- `thrust::transform_reduce`
- `thrust::inclusive_scan`
- `thrust::sort`
- `thrust::copy`
- ...



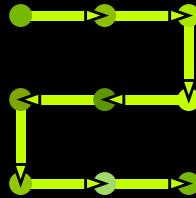
## Extended Algorithms

- `thrust::reduce_by_key`
- `thrust::sort_by_key`
- `thrust::tabulate`
- `thrust::gather`
- ...



## Containers

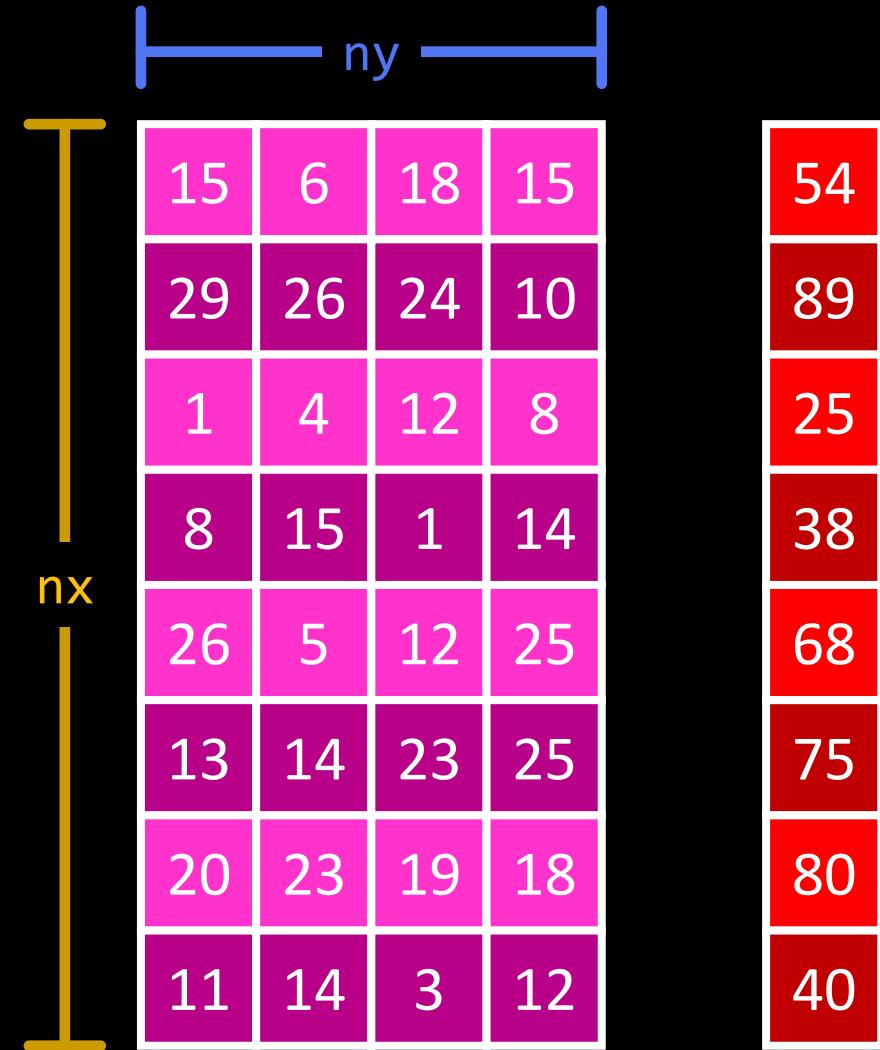
- `thrust::device_vector`
- `thrust::host_vector`
- `thrust::universal_vector`
- `thrust::allocate_unique`
- ...



## Iterators

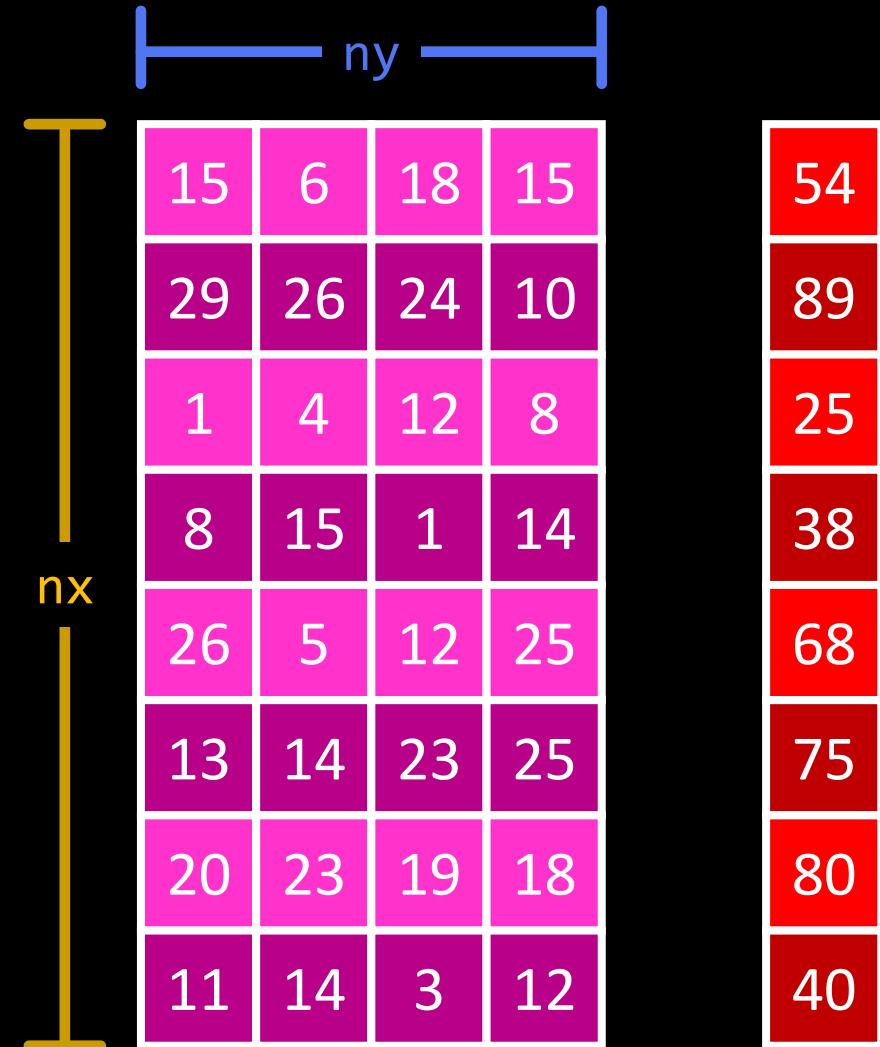
- `thrust::counting_iterator`
- `thrust::constant_iterator`
- `thrust::transform_iterator`
- `thrust::zip_iterator`
- ...

```
int nx(...), ny(...);  
thrust::universal_vector<float> M(nx * ny);
```



```
int nx(...), ny(...);  
thrust::universal_vector<float> M(nx * ny);
```

Each set of consecutive keys that compare equal defines a group.



```
int nx(...), ny(...);  
thrust::universal_vector<float> M(nx * ny);
```

Each set of consecutive keys that compare equal defines a group.

ny			
nx	0	1	2
0	0	0	0
1	1	1	1
2	2	2	2
3	3	3	3
4	4	4	4
5	5	5	5
6	6	6	6
7	7	7	7

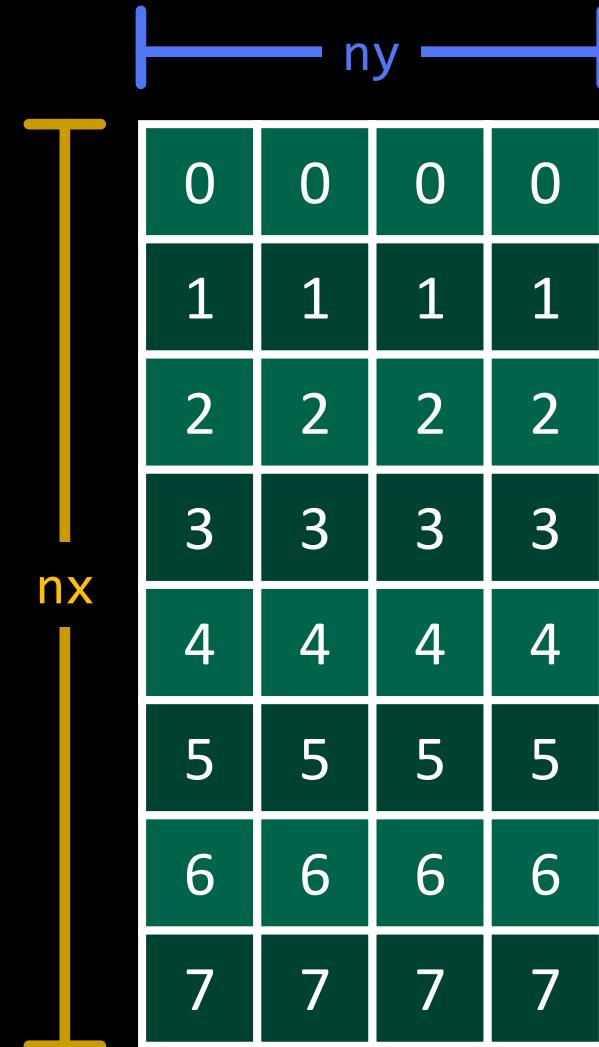
```

int nx(...), ny(...);
thrust::universal_vector<float> M(nx * ny);

auto row_ids_begin =
    thrust::make_transform_iterator(
        thrust::make_counting_iterator(0),
        [=] __host__ __device__ (int x) {
            return x / ny;
        });
auto row_ids_end = row_ids_begin + nx * ny;

```

Each set of consecutive keys that compare equal defines a group.



0	0	0	0
1	1	1	1
2	2	2	2
3	3	3	3
4	4	4	4
5	5	5	5
6	6	6	6
7	7	7	7

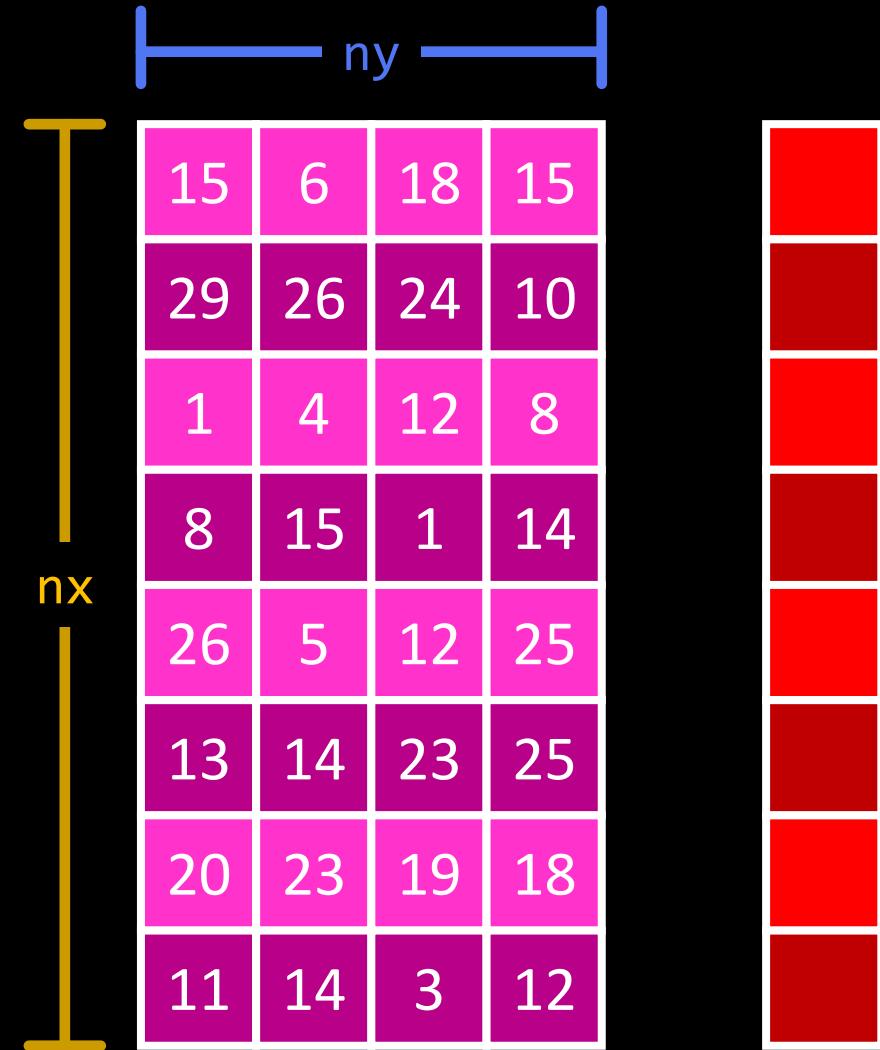
```

int nx(...), ny(...);
thrust::universal_vector<float> M(nx * ny);

auto row_ids_begin =
    thrust::make_transform_iterator(
        thrust::make_counting_iterator(0),
        [=] __host__ __device__ (int x) {
            return x / ny;
        });
auto row_ids_end = row_ids_begin + nx * ny;

thrust::universal_vector<float> sums(nx);

```



```

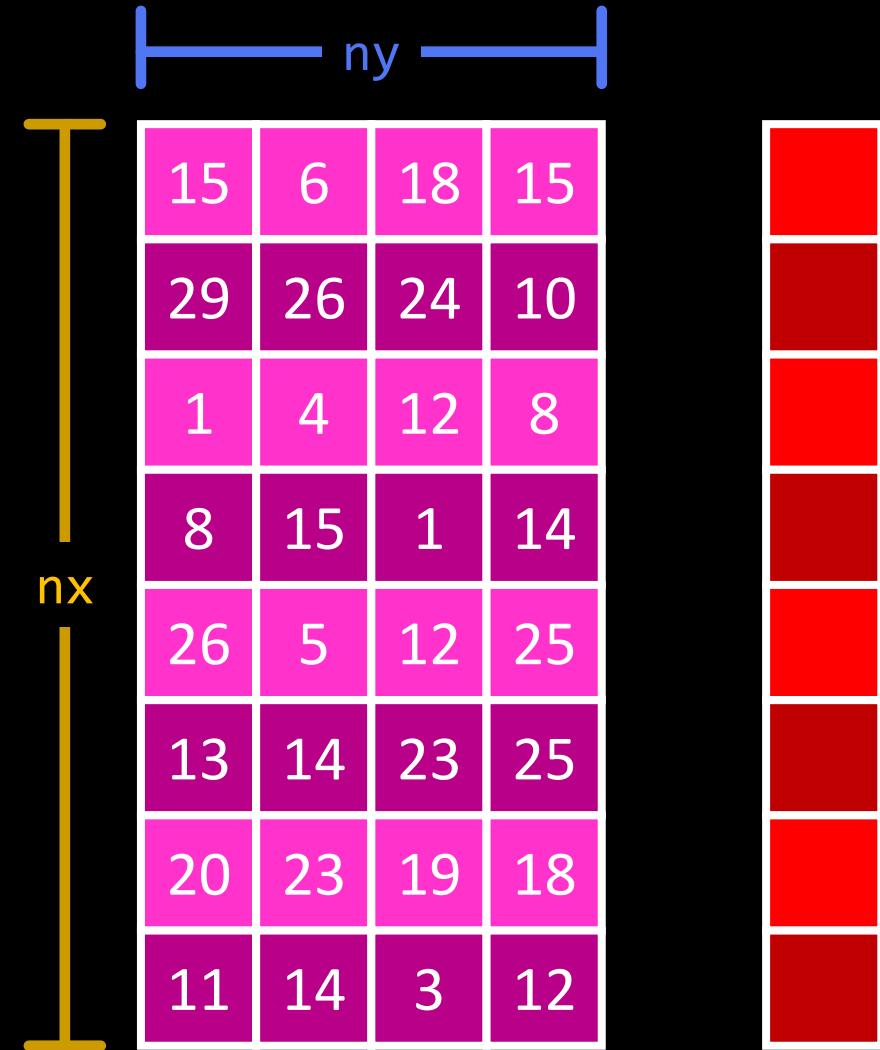
int nx(...), ny(...);
thrust::universal_vector<float> M(nx * ny);

auto row_ids_begin =
    thrust::make_transform_iterator(
        thrust::make_counting_iterator(0),
        [=] __host__ __device__ (int x) {
            return x / ny;
        });
auto row_ids_end = row_ids_begin + nx * ny;

thrust::universal_vector<float> sums(nx);

thrust::reduce_by_key(thrust::cuda::par,
...);

```



```

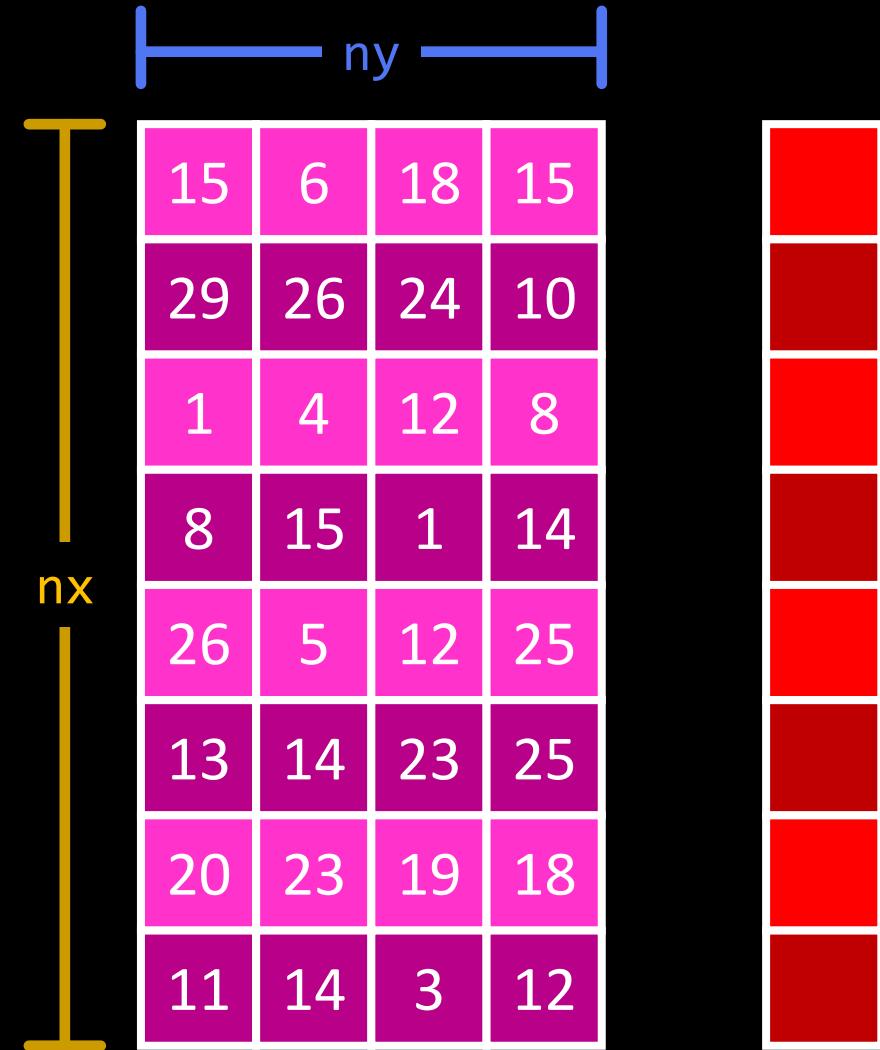
int nx(...), ny(...);
thrust::universal_vector<float> M(nx * ny);

auto row_ids_begin =
    thrust::make_transform_iterator(
        thrust::make_counting_iterator(0),
        [=] __host__ __device__ (int x) {
            return x / ny;
        });
auto row_ids_end = row_ids_begin + nx * ny;

thrust::universal_vector<float> sums(nx);

thrust::reduce_by_key(thrust::cuda::par,
// Keys input.
...);

```



```

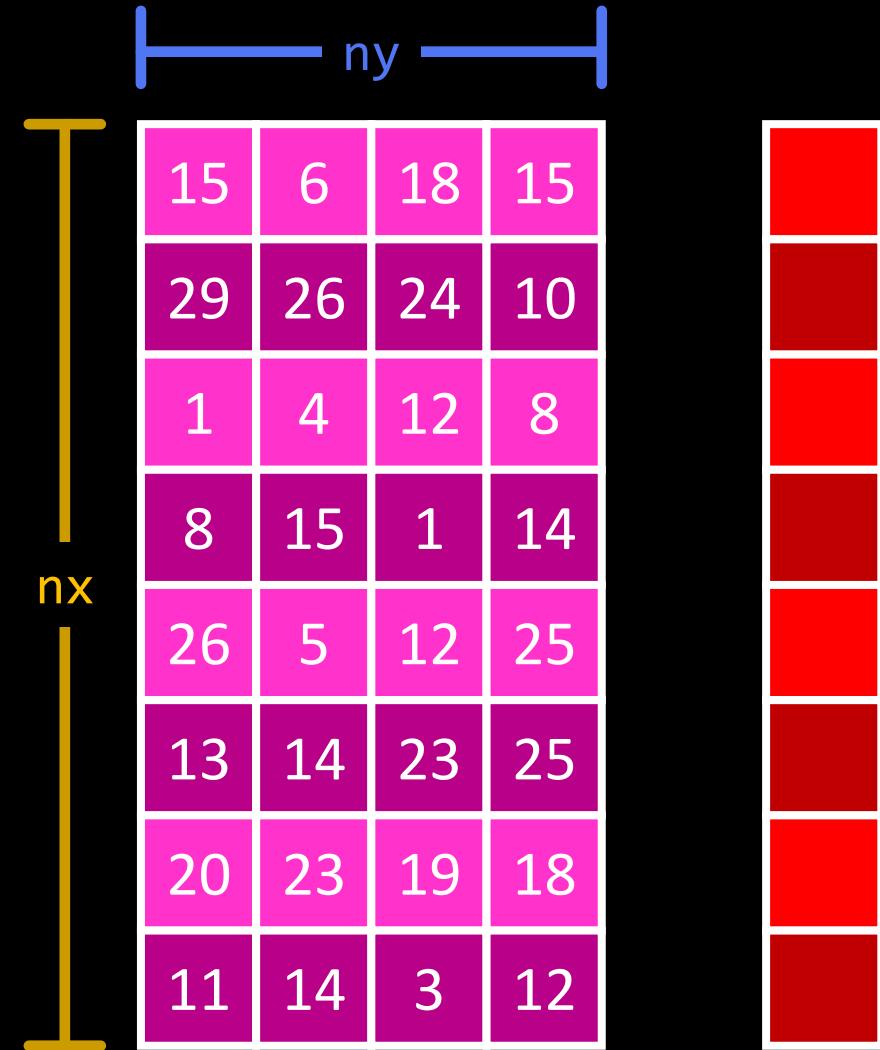
int nx(...), ny(...);
thrust::universal_vector<float> M(nx * ny);

auto row_ids_begin =
    thrust::make_transform_iterator(
        thrust::make_counting_iterator(0),
        [=] __host__ __device__ (int x) {
            return x / ny;
        });
auto row_ids_end = row_ids_begin + nx * ny;

thrust::universal_vector<float> sums(nx);

thrust::reduce_by_key(thrust::cuda::par,
    row_ids_begin, row_ids_end,
    // Values input.
    ...);

```



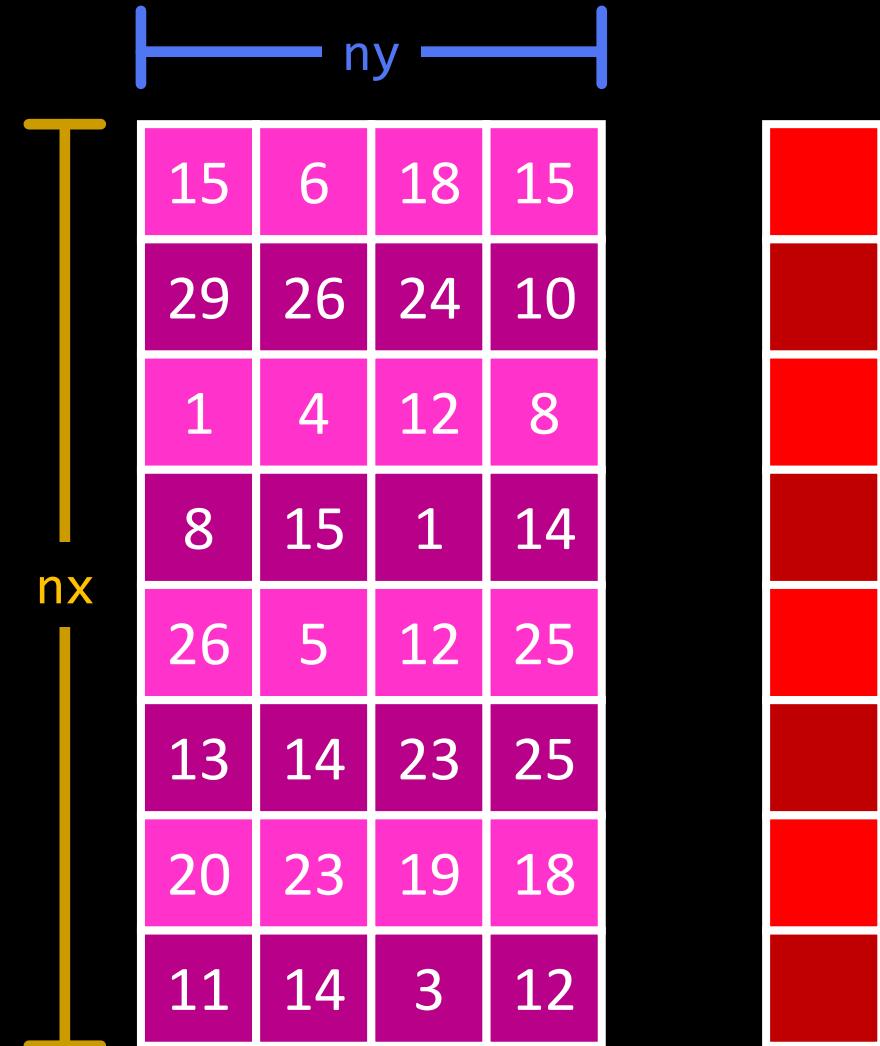
```
int nx(...), ny(...);  
thrust::universal_vector<float> M(nx * ny);
```

```
auto row_ids_begin =  
    thrust::make_transform_iterator(  
        thrust::make_counting_iterator(0),  
        [=] __host__ __device__ (int x) {  
            return x / ny;  
        });
```

```
auto row_ids_end =  
    thrust::make_transform_iterator(  
        thrust::make_counting_iterator(nx),  
        [=] __host__ __device__ (int x) {  
            return x / ny;  
        });
```

The underlying algorithm is applied  
to the values in each group.

```
thrust::reduce_by_key(thrust::cuda::par,  
    row_ids_begin, row_ids_end,  
    M.begin(),  
    ...);
```



```

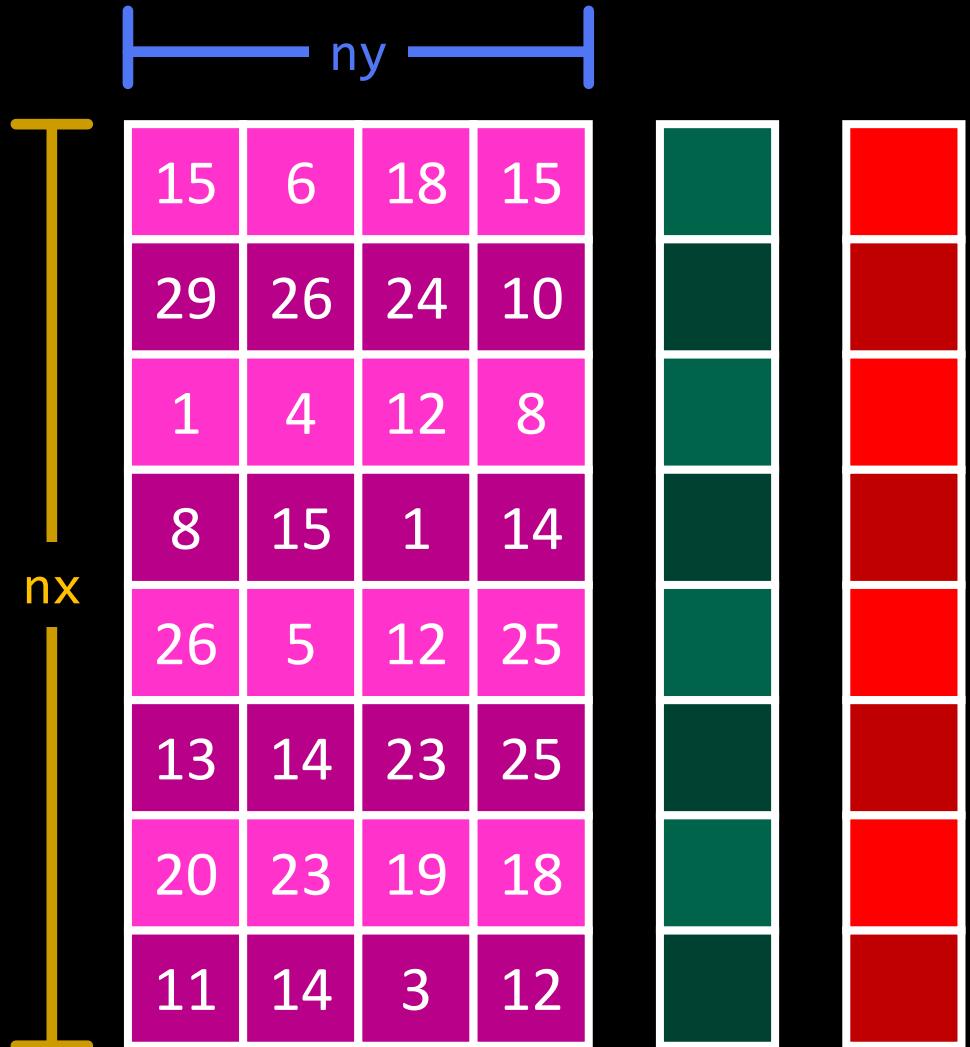
int nx(...), ny(...);
thrust::universal_vector<float> M(nx * ny);

auto row_ids_begin =
    thrust::make_transform_iterator(
        thrust::make_counting_iterator(0),
        [=] __host__ __device__ (int x) {
            return x / ny;
        });
auto row_ids_end = row_ids_begin + nx * ny;

thrust::universal_vector<float> sums(nx);

thrust::reduce_by_key(thrust::cuda::par,
    row_ids_begin, row_ids_end,
    M.begin(),
    // Keys output.
    ...);

```



```

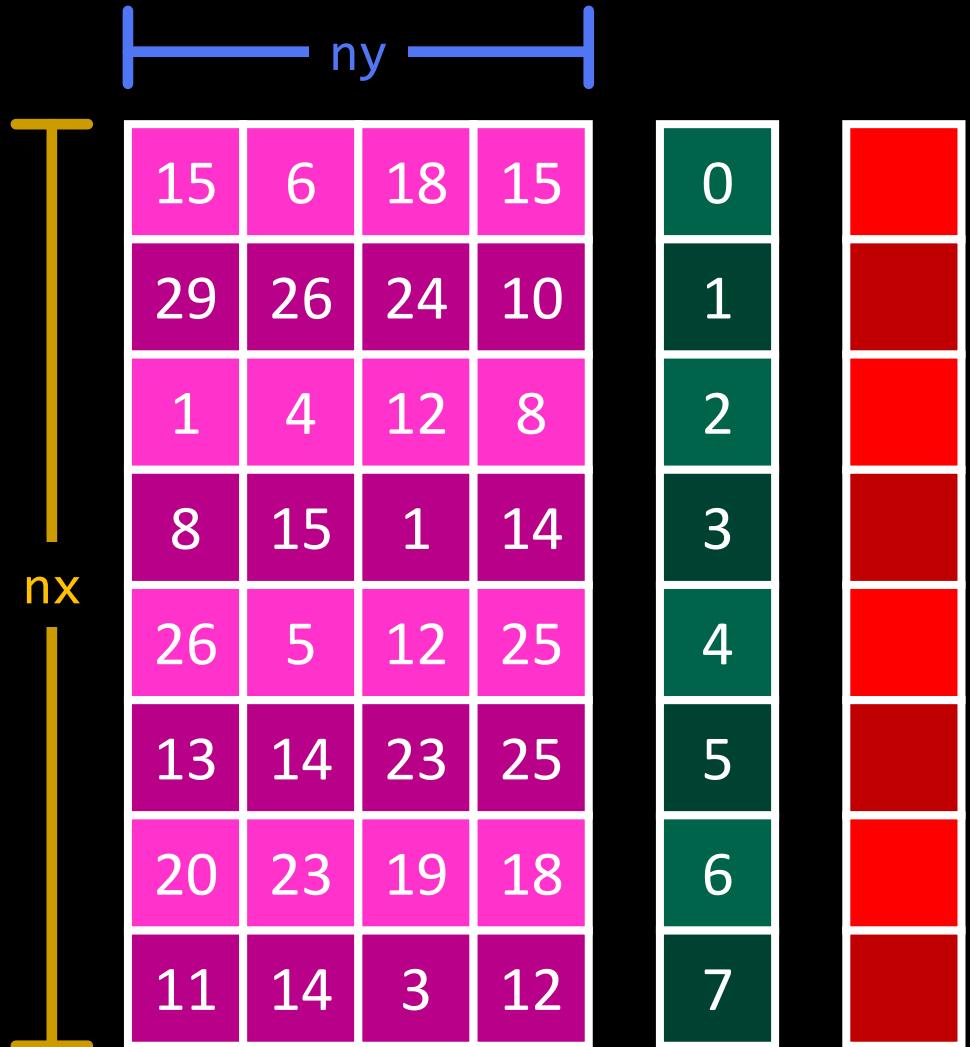
int nx(...), ny(...);
thrust::universal_vector<float> M(nx * ny);

auto row_ids_begin =
    thrust::make_transform_iterator(
        thrust::make_counting_iterator(0),
        [=] __host__ __device__ (int x) {
            return x / ny;
        });
auto ro = thrust::make_transform_iterator(
    row_ids_begin,
    [=] __host__ __device__ (int x) {
        return M[x];
    });

The first key in each group is copied
to the key output range.

thrust::reduce_by_key(thrust::cuda::par,
    row_ids_begin, row_ids_end,
    M.begin(),
    // Keys output.
    ...);

```



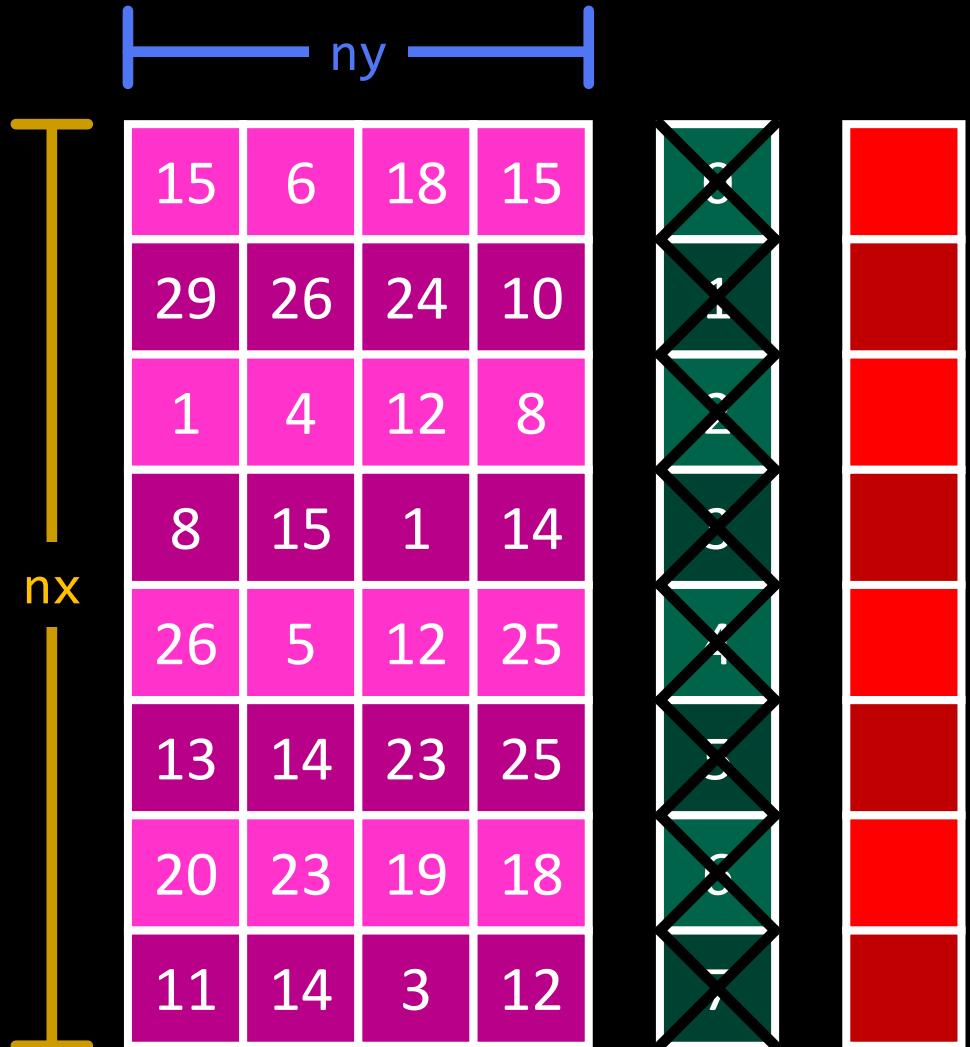
```

int nx(...), ny(...);
thrust::universal_vector<float> M(nx * ny);

auto row_ids_begin =
    thrust::make_transform_iterator(
        thrust::make_counting_iterator(0),
        [=] __host__ __device__ (int x) {
            return x / ny;
        });
auto ro = thrust::make_discard_iterator();
thrust::reduce_by_key(thrust::cuda::par,
    row_ids_begin, row_ids_end,
    M.begin(),
    thrust::make_discard_iterator(),
    ...);

```

The first key in each group is copied to the key output range.



```

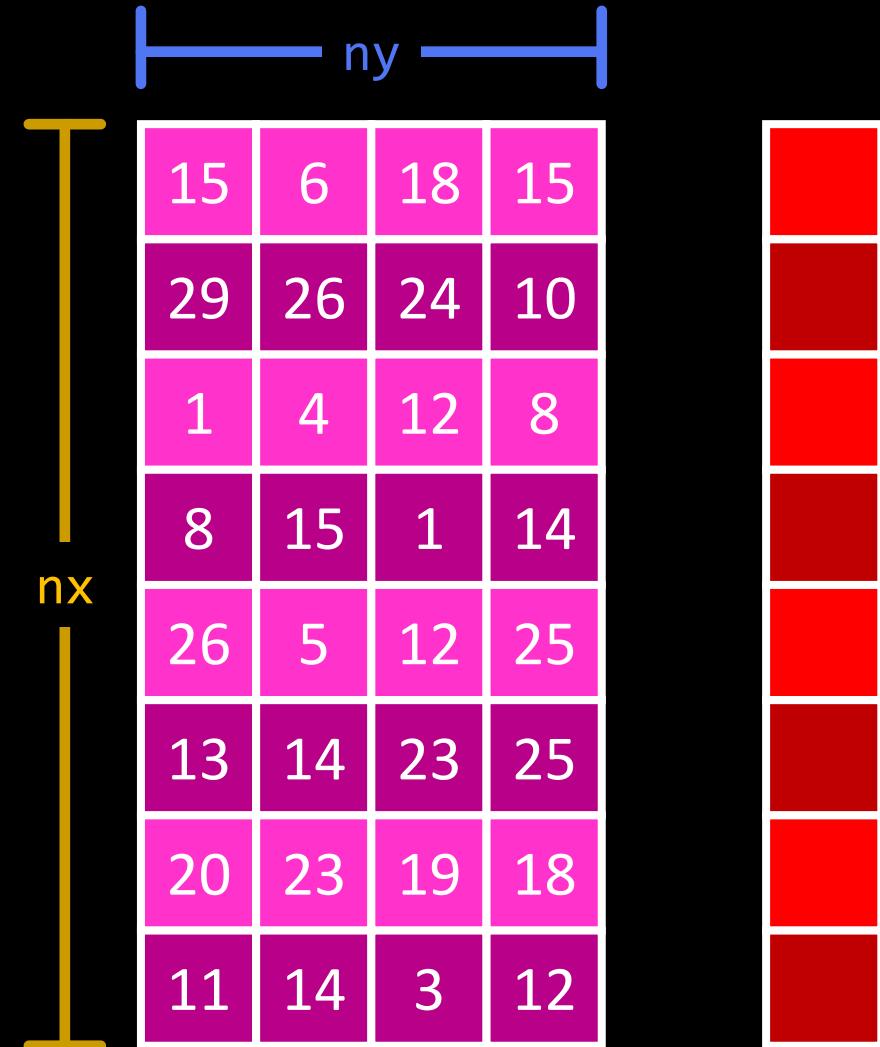
int nx(...), ny(...);
thrust::universal_vector<float> M(nx * ny);

auto row_ids_begin =
    thrust::make_transform_iterator(
        thrust::make_counting_iterator(0),
        [=] __host__ __device__ (int x) {
            return x / ny;
        });
auto row_ids_end = row_ids_begin + nx * ny;

thrust::universal_vector<float> sums(nx);

thrust::reduce_by_key(thrust::cuda::par,
    row_ids_begin, row_ids_end,
    M.begin(),
    thrust::make_discard_iterator(),
    // Values output.);

```

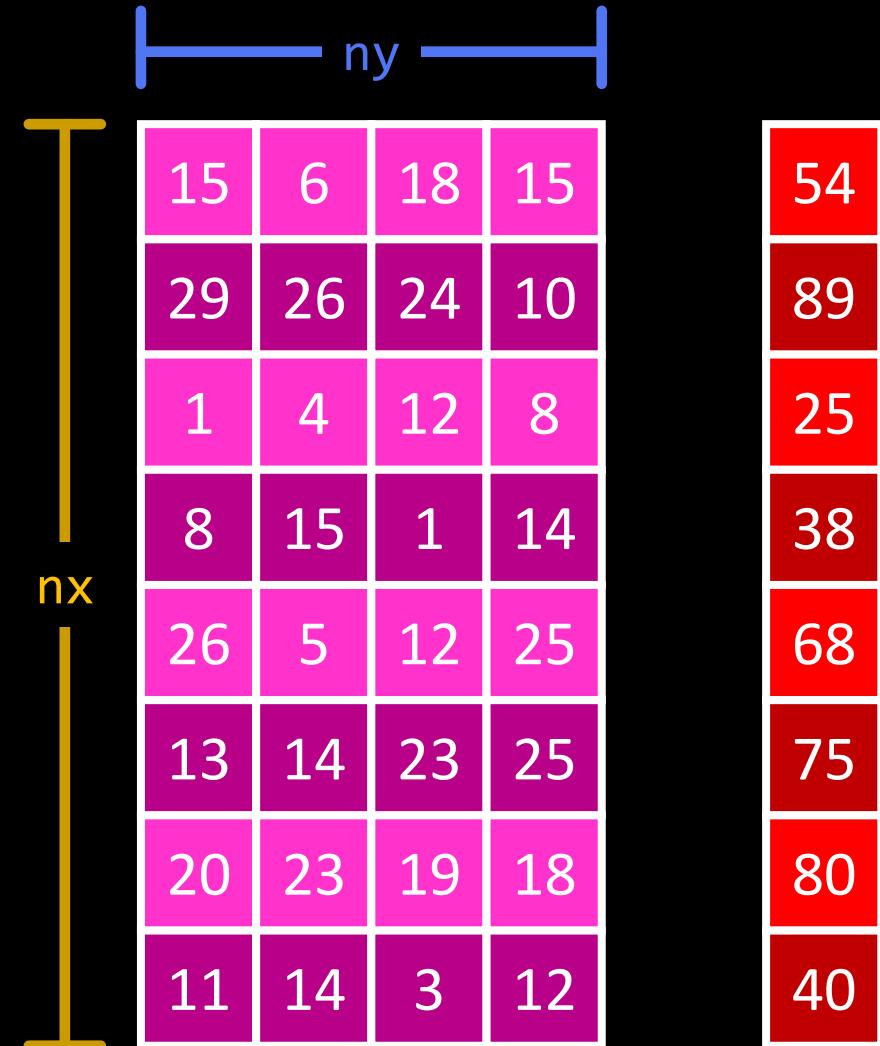


```
int nx(...), ny(...);  
thrust::universal_vector<float> M(nx * ny);
```

```
auto row_ids_begin =  
    thrust::make_transform_iterator(  
        thrust::make_counting_iterator(0),  
        [=] __host__ __device__ (int x) {  
            return x / ny;  
        });
```

```
auto ro [The algorithm result for each group  
is written to the values output.]  
thrust:
```

```
thrust::reduce_by_key(thrust::cuda::par,  
                     row_ids_begin, row_ids_end,  
                     M.begin(),  
                     thrust::make_discard_iterator(),  
                     sums.begin());
```



```

int nx(...), ny(...);
thrust::universal_vector<float> M(nx * ny);

auto row_ids_begin =
    thrust::make_transform_iterator(
        thrust::make_counting_iterator(0),
        [=] __host__ __device__ (int x) {
            return x / ny;
        });
auto row_ids_end = row_ids_begin + nx * ny;

thrust::universal_vector<float> sums(nx);

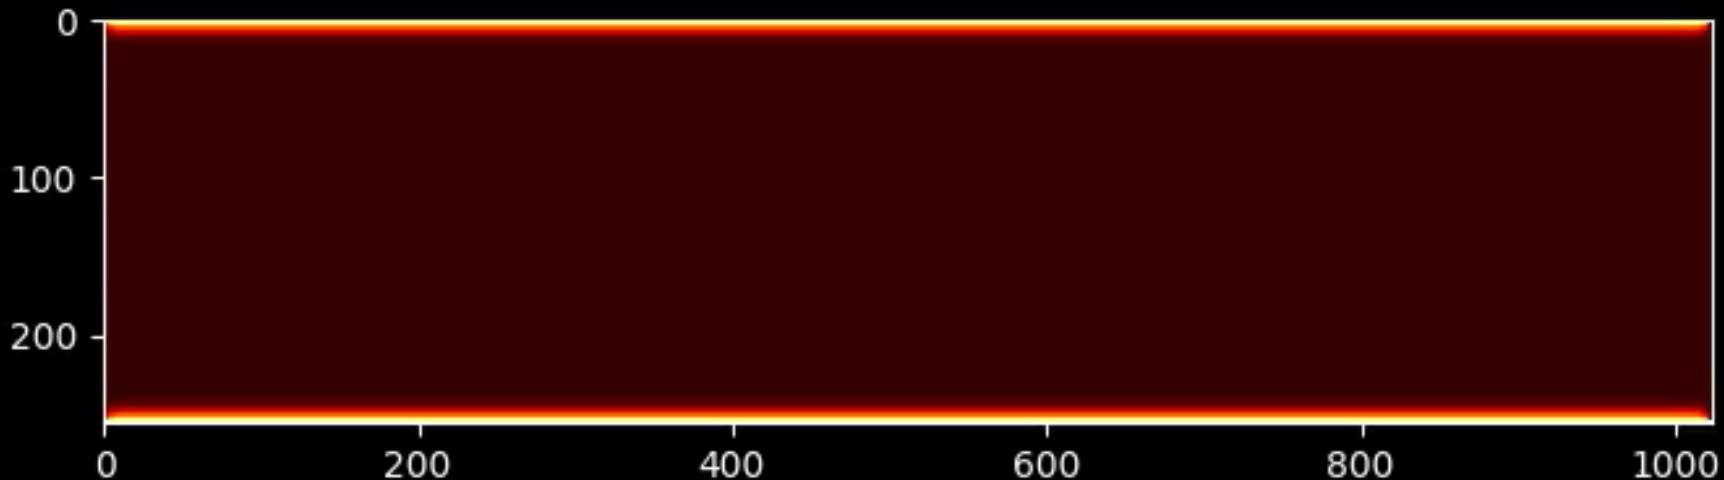
thrust::reduce_by_key(thrust::cuda::par,
    row_ids_begin, row_ids_end,
    M.begin(),
    thrust::make_discard_iterator(),
    sums.begin());

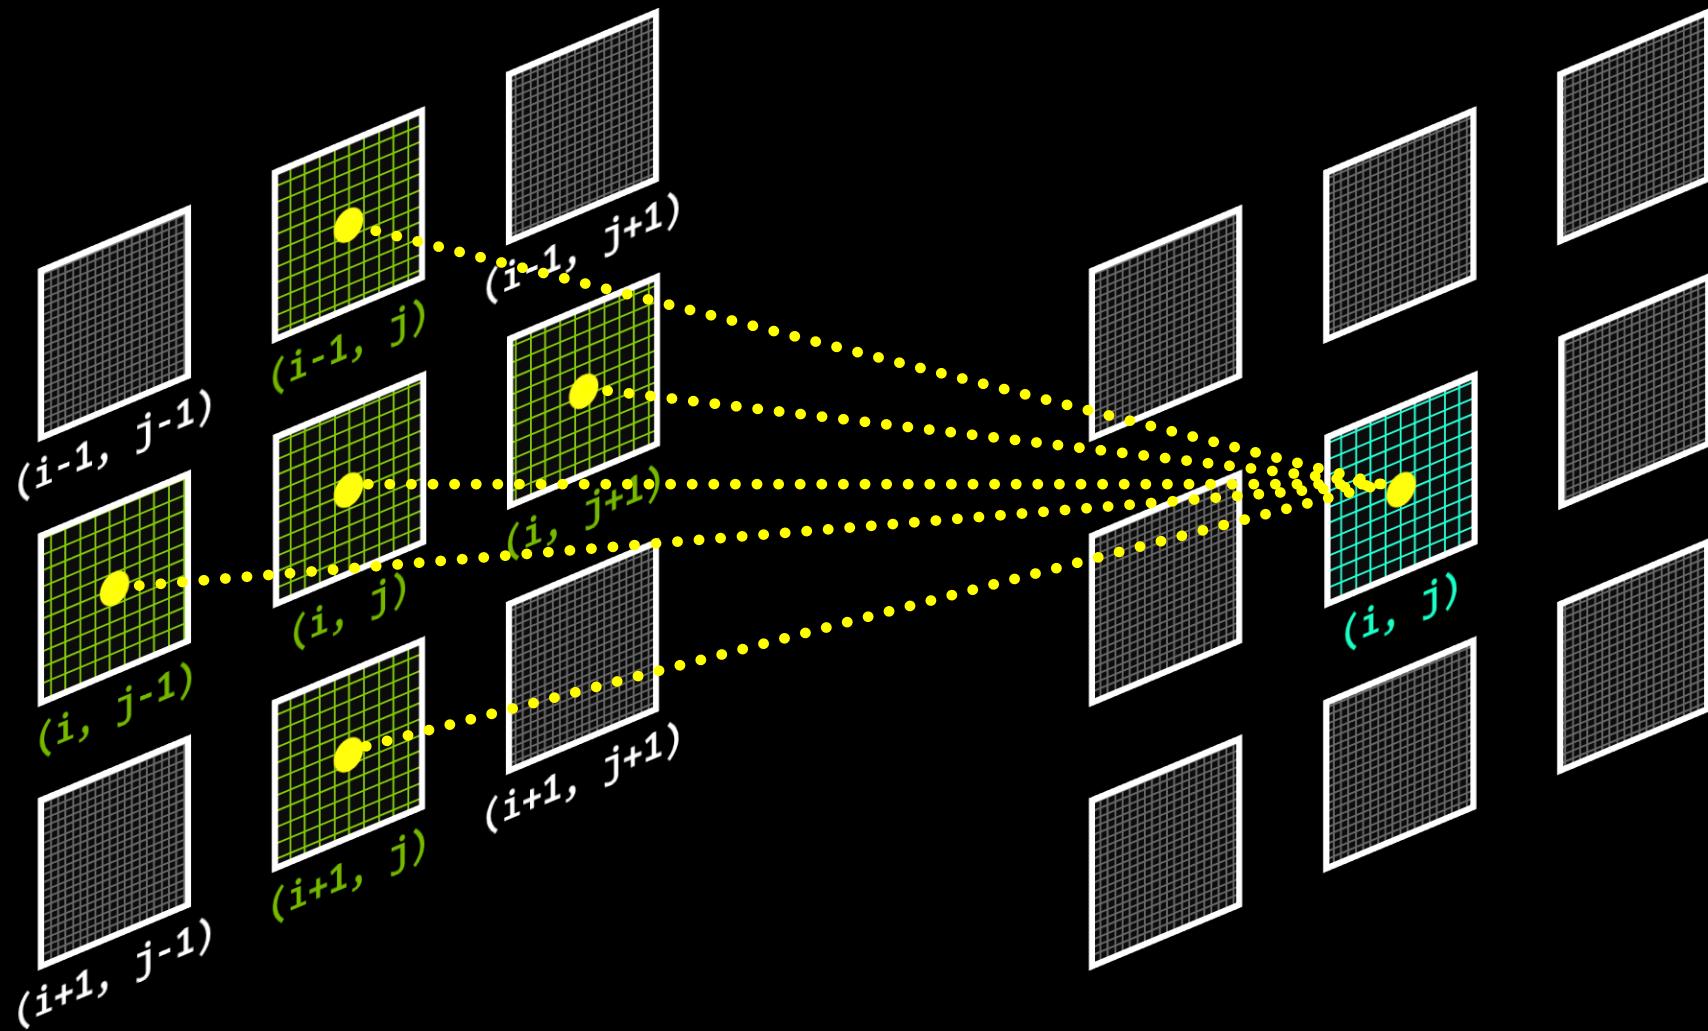
```

15	6	18	15	54
29	26	24	10	89
1	4	12	8	25
8	15	1	14	38
26	5	12	25	68
13	14	23	25	75
20	23	19	18	80
11	14	3	12	40

# By Key Algorithms

- `reduce_by_key`
- `(inclusive|exclusive)_scan_by_key`
- `sort_by_key`
- `unique_by_key`
- `merge_by_key`
- `set_(union|difference|intersection)_by_key`





```
void heat_equation(auto policy, int nx, int ny,
                   const thrust::universal_vector<float>& U_data,
                   thrust::universal_vector<float>& V_data);
```

```
void heat_equation(auto policy, int nx, int ny,
                   const thrust::universal_vector<float>& U_data,
                   thrust::universal_vector<float>& V_data) {
    ...
    thrust::for_each_n(policy, thrust::make_counting_iterator(0), U.size(), ...);
}
```

```
void heat_equation(auto policy, int nx, int ny,
                    const thrust::universal_vector<float>& U_data,
                    thrust::universal_vector<float>& V_data) {
    ...
    thrust::for_each_n(policy, thrust::make_counting_iterator(0), U.size(),
                      [?, nx, ny] __host__ __device__ (int xy) {
    ...
});
```

```
void heat_equation(auto policy, int nx, int ny,
                    const thrust::universal_vector<float>& U_data,
                    thrust::universal_vector<float>& V_data) {
    const float* U = thrust::raw_pointer_cast(U_data.data());
    const float* V = thrust::raw_pointer_cast(V_data.data());
    thrust::for_each_n(policy, thrust::make_counting_iterator(0), U.size(),
                      [U, V, nx, ny] __host__ __device__ (int xy) {
    ...
});
```

```
void heat_equation(auto policy, int nx, int ny,
                    const thrust::universal_vector<float>& U_data,
                    thrust::universal_vector<float>& V_data) {
    const float* U = thrust::raw_pointer_cast(U_data.data());
    const float* V = thrust::raw_pointer_cast(V_data.data());
    thrust::for_each_n(policy, thrust::make_counting_iterator(0), U.size(),
                      [U, V, nx, ny] __host__ __device__ (int xy) {
        int x = xy / ny;
        int y = xy % ny;

        ...
    });
}
```

```
void heat_equation(auto policy, int nx, int ny,
                    const thrust::universal_vector<float>& U_data,
                    thrust::universal_vector<float>& V_data) {
    const float* U = thrust::raw_pointer_cast(U_data.data());
    const float* V = thrust::raw_pointer_cast(V_data.data());
    thrust::for_each_n(policy, thrust::make_counting_iterator(0), U.size(),
                      [U, V, nx, ny] __host__ __device__ (int xy) {
        int x = xy / ny;
        int y = xy % ny;

        if (x > 0 && y > 0 && x < nx - 1 && y < ny - 1) {
            ...
        } else ...
    });
}
```

```

void heat_equation(auto policy, int nx, int ny,
                   const thrust::universal_vector<float>& U_data,
                   thrust::universal_vector<float>& V_data) {
    const float* U = thrust::raw_pointer_cast(U_data.data());
    const float* V = thrust::raw_pointer_cast(V_data.data());
    thrust::for_each_n(policy, thrust::make_counting_iterator(0), U.size(),
                      [U, V, nx, ny] __host__ __device__ (int xy) {
        int x = xy / ny;
        int y = xy % ny;

        if (x > 0 && y > 0 && x < nx - 1 && y < ny - 1) {
            auto d2tdx2 = U[x*ny + y - 1] - U[x*ny + y] * 2 + U[x*ny + y + 1];
            auto d2tdy2 = U[(x - 1)*ny + y] - U[x*ny + y] * 2 + U[(x + 1)*ny + y];

            V[x*ny + y] = U[x*ny + y] + 0.2f * (d2tdx2 + d2tdy2);
        } else ...
    });
}

```

```

void heat_equation(auto policy, int nx, int ny,
                   const thrust::universal_vector<float>& U_data,
                   thrust::universal_vector<float>& V_data) {
    const float* U = thrust::raw_pointer_cast(U_data.data());
    const float* V = thrust::raw_pointer_cast(V_data.data());
    thrust::for_each_n(policy, thrust::make_counting_iterator(0), U.size(),
                      [U, V, nx, ny] __host__ __device__ (int xy) {
        int x = xy / ny;
        int y = xy % ny;

        if (x > 0 && y > 0 && x < nx - 1 && y < ny - 1) {
            auto d2tdx2 = U[x*ny + y - 1] - U[x*ny + y] * 2 + U[x*ny + y + 1];
            auto d2tdy2 = U[(x - 1)*ny + y] - U[x*ny + y] * 2 + U[(x + 1)*ny + y];

            V[x*ny + y] = U[x*ny + y] + 0.2f * (d2tdx2 + d2tdy2);
        } else {
            V[x*ny + y] = U[x*ny + y];
        }
    });
}

```

# `mdspan`: A non-owning handle to multidimensional data

```
template <class T, class Extents, class LayoutPolicy = ..., class AccessorPolicy = ...>
class std::mdspan;
```

```
mdspan A(data, N, M}; // Row-major: C/NumPy.           // Column-major: Fortran/MATLAB.
mdspan A(data, layout_right::mapping(N, M));          mdspan B(data, layout_left::mapping(N, M));
```

```
A(i, j)      == data(i * M + j)
A.stride(0)  == M
A.stride(1)  == 1
```

```
B(i, j)      == data(i + j * N)
B.stride(0)  == 1
B.stride(1)  == N
```

Location	Element
0	$a_{00}$
1	$a_{01}$
2	$a_{10}$
3	$a_{11}$

$$\begin{bmatrix} a_{00} & a_{01} \\ a_{10} & a_{11} \end{bmatrix}$$

Location	Element
0	$a_{00}$
1	$a_{10}$
2	$a_{01}$
3	$a_{11}$

Learn More: [S51755](#) C++ Standard Parallelism (2023)

```

void heat_equation(auto policy, int nx, int ny,
                    const thrust::universal_vector<float>& U_data,
                    thrust::universal_vector<float>& V_data) {
    const float* U = thrust::raw_pointer_cast(U_data.data());
    const float* V = thrust::raw_pointer_cast(V_data.data());
    thrust::for_each_n(policy, thrust::make_counting_iterator(0), U.size(),
                      [U, V, nx, ny] __host__ __device__ (int xy) {
        int x = xy / ny;
        int y = xy % ny;

        if (x > 0 && y > 0 && x < nx - 1 && y < ny - 1) {
            auto d2tdx2 = U[x*ny + y - 1] - U[x*ny + y] * 2 + U[x*ny + y + 1];
            auto d2tdy2 = U[(x - 1)*ny + y] - U[x*ny + y] * 2 + U[(x + 1)*ny + y];

            V[x*ny + y] = U[x*ny + y] + 0.2f * (d2tdx2 + d2tdy2);
        } else {
            V[x*ny + y] = U[x*ny + y];
        }
    });
}

```

```

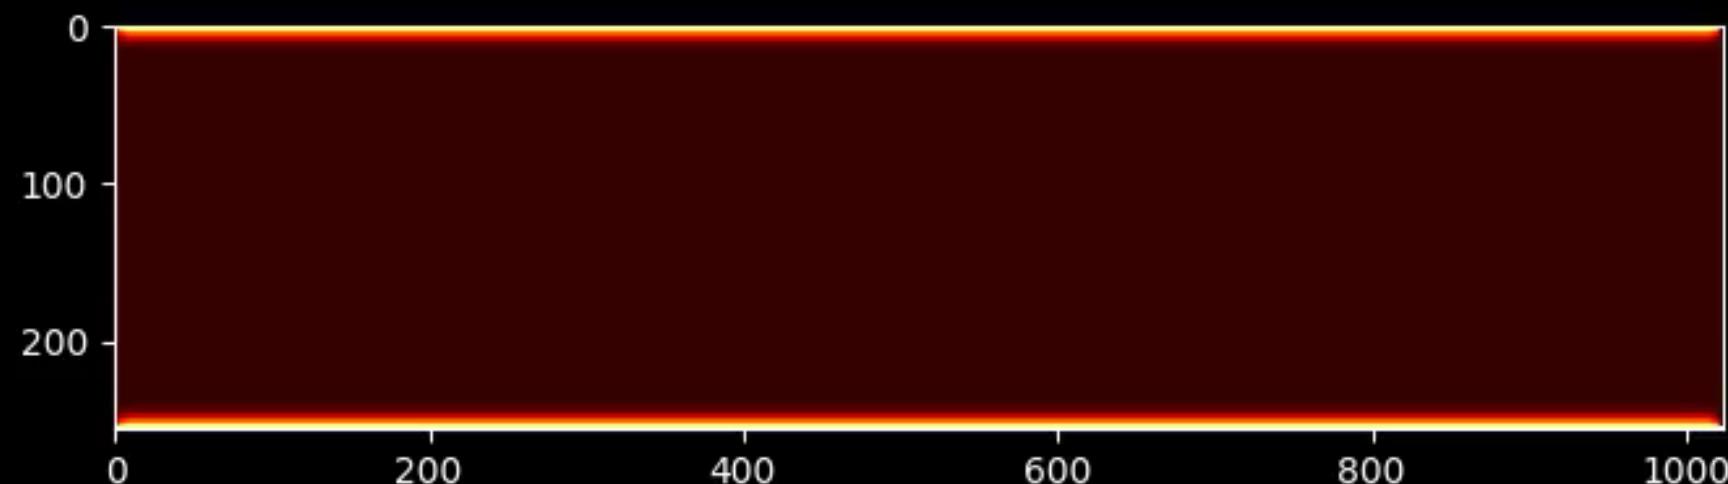
void heat_equation(auto policy,
                    cuda::std::mdspan<float const, cuda::std::dims<2>> U,
                    cuda::std::mdspan<float, cuda::std::dims<2>> V) {
    thrust::for_each_n(policy, thrust::make_counting_iterator(0), U.size(),
                      [U, V] __host__ __device__ (int xy) {
                        int x = xy / U.extent(1);
                        int y = xy % U.extent(1);

                        if (x > 0 && y > 0 && x < I.extent(0) - 1 && y < U.extent(1) - 1) {
                            auto d2tdx2 = U(x, y - 1) - U(x, y) * 2 + U(x, y + 1);
                            auto d2tdy2 = U(x - 1, y) - U(x, y) * 2 + U(x + 1, y);

                            V(x, y) = U(x, y) + 0.2f * (d2tdx2 + d2tdy2);
                        } else {
                            V(x, y) = U(x, y);
                        }
                    });
}

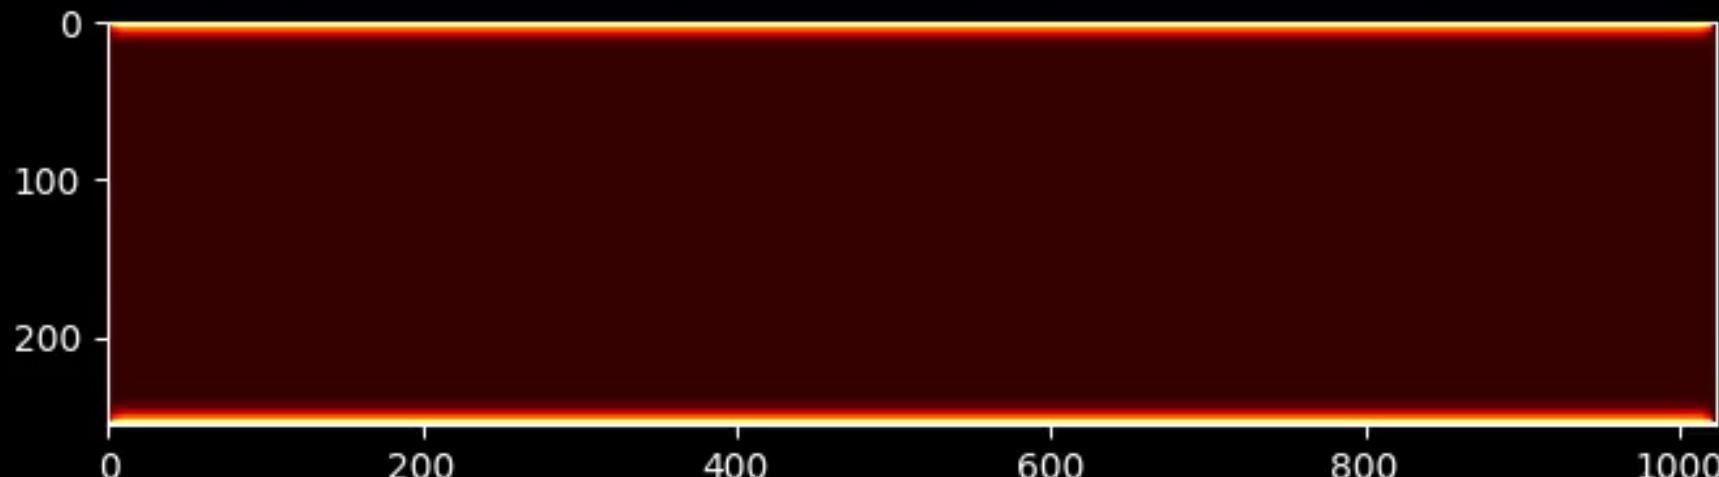
```

```
void initialize_oven(auto policy, cuda::std::mdspan<float, cuda::std::dims<2>> U);
```

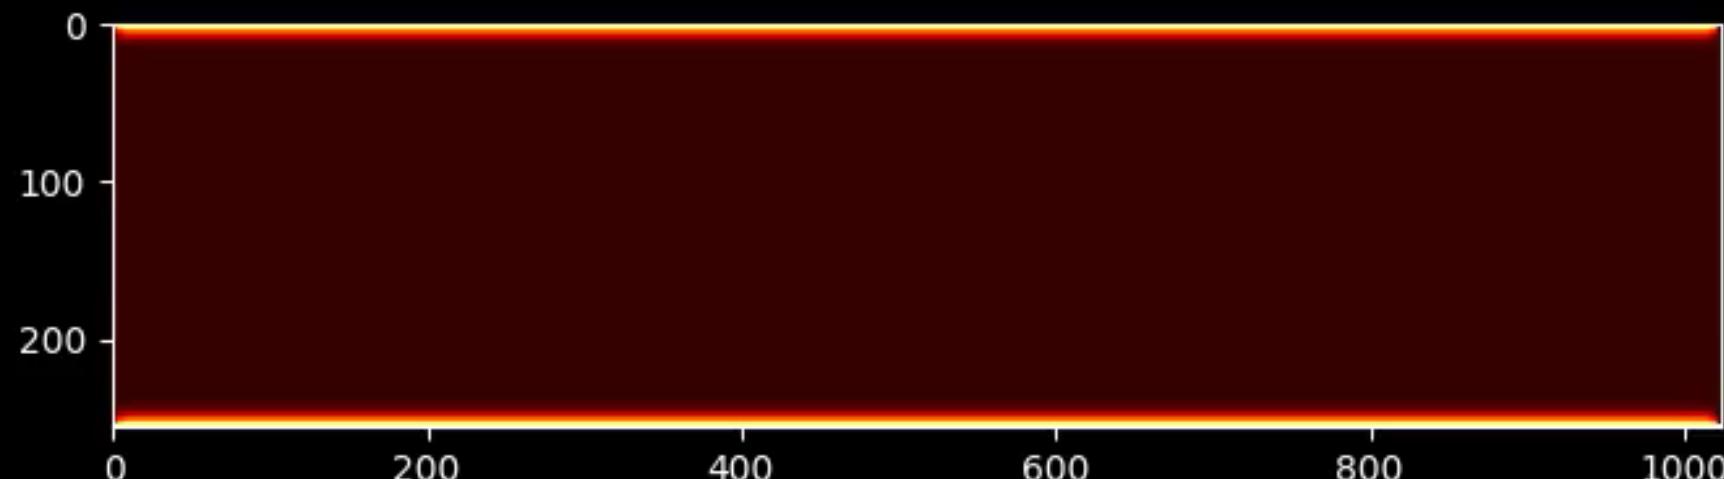


```
void initialize_oven(auto policy, cuda::std::mdspan<float, cuda::std::dims<2>> U) {
    auto nx = U.extent(0);
    auto top = cuda::std::submdspan(U, 0, cuda::std::full_extent);

    auto bot = cuda::std::submdspan(U, nx-1, cuda::std::full_extent);
}
```



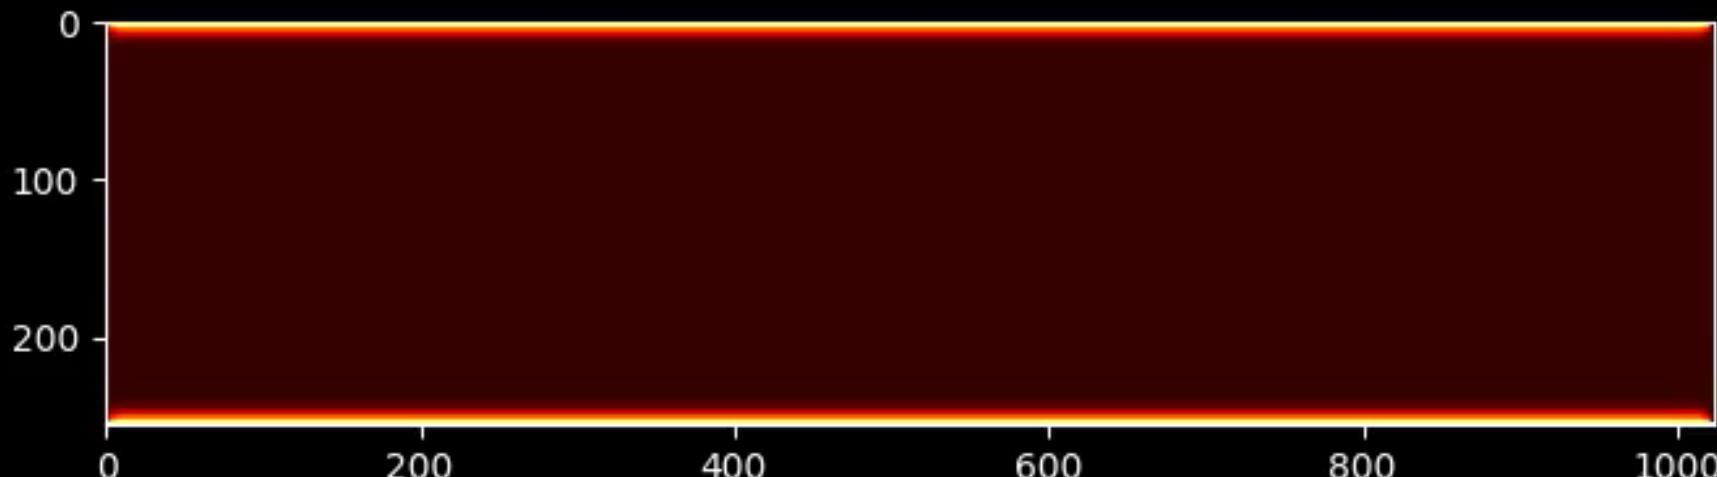
```
void initialize_oven(auto policy, cuda::std::mdspan<float, cuda::std::dims<2>> U) {
    auto nx = U.extent(0);
    auto top = cuda::std::submdspan(U, 0, cuda::std::full_extent);
    auto mid = cuda::std::submdspan(U, cuda::std::tuple(1, nx-2), cuda::std::full_extent);
    auto bot = cuda::std::submdspan(U, nx-1, cuda::std::full_extent);
}
```



```
void initialize_oven(auto policy, cuda::std::mdspan<float, cuda::std::dims<2>> U) {
    auto nx = U.extent(0);
    auto top = cuda::std::submdspan(U, 0, cuda::std::full_extent);
    auto mid = cuda::std::submdspan(U, cuda::std::tuple(1, nx-2), cuda::std::full_extent);
    auto bot = cuda::std::submdspan(U, nx-1, cuda::std::full_extent);

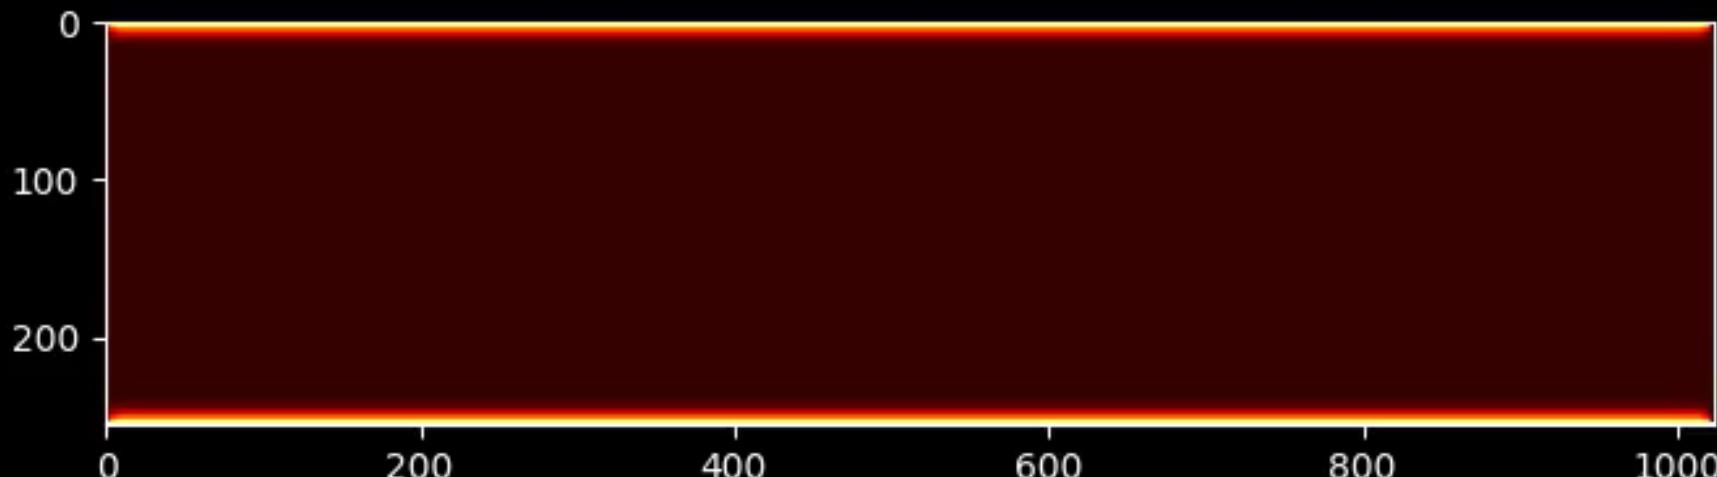
    thrust::fill_n(policy, top.data(), top.size(), 90.0);

    thrust::fill_n(policy, bot.data(), bot.size(), 90.0);
}
```



```
void initialize_oven(auto policy, cuda::std::mdspan<float, cuda::std::dims<2>> U) {
    auto nx = U.extent(0);
    auto top = cuda::std::submdspan(U, 0, cuda::std::full_extent);
    auto mid = cuda::std::submdspan(U, cuda::std::tuple(1, nx-2), cuda::std::full_extent);
    auto bot = cuda::std::submdspan(U, nx-1, cuda::std::full_extent);

    thrust::fill_n(policy, top.data(), top.size(), 90.0);
    thrust::fill_n(policy, mid.data(), mid.size(), 15.0);
    thrust::fill_n(policy, bot.data(), bot.size(), 90.0);
}
```



```
int nx(...), ny(...), write_steps(...), compute_steps(...);
```

```
int nx(...), ny(...), write_steps(...), compute_steps(...);

thrust::universal_vector<float> U_data(nx * ny), V_data(nx * ny);
cuda::std::mdspan U(thrust::raw_pointer_cast(U_data.data()), nx, ny);
cuda::std::mdspan V(thrust::raw_pointer_cast(V_data.data()), nx, ny);
```

```
int nx(...), ny(...), write_steps(...), compute_steps(...);

thrust::universal_vector<float> U_data(nx * ny), V_data(nx * ny);
cuda::std::mdspan U(thrust::raw_pointer_cast(U_data.data()), nx, ny);
cuda::std::mdspan V(thrust::raw_pointer_cast(V_data.data()), nx, ny);

initialize_oven(thrust::cuda::par, U);
```

```
int nx(...), ny(...), write_steps(...), compute_steps(...);

thrust::universal_vector<float> U_data(nx * ny), V_data(nx * ny);
cuda::std::mdspan U(thrust::raw_pointer_cast(U_data.data()), nx, ny);
cuda::std::mdspan V(thrust::raw_pointer_cast(V_data.data()), nx, ny);

initialize_oven(thrust::cuda::par, U);

for (auto write_step : std::views::iota(0, write_steps)) {
    ...
}
```

```
int nx(...), ny(...), write_steps(...), compute_steps(...);

thrust::universal_vector<float> U_data(nx * ny), V_data(nx * ny);
cuda::std::mdspan U(thrust::raw_pointer_cast(U_data.data()), nx, ny);
cuda::std::mdspan V(thrust::raw_pointer_cast(V_data.data()), nx, ny);

initialize_oven(thrust::cuda::par, U);

for (auto write_step : std::views::iota(0, write_steps)) {
    save_to_file(U);

    for (auto compute_step : std::views::iota(0, compute_steps)) {
        ...
    }
}
```

```
int nx(...), ny(...), write_steps(...), compute_steps(...);

thrust::universal_vector<float> U_data(nx * ny), V_data(nx * ny);
cuda::std::mdspan U(thrust::raw_pointer_cast(U_data.data()), nx, ny);
cuda::std::mdspan V(thrust::raw_pointer_cast(V_data.data()), nx, ny);

initialize_oven(thrust::cuda::par, U);

for (auto write_step : std::views::iota(0, write_steps)) {
    save_to_file(U);

    for (auto compute_step : std::views::iota(0, compute_steps)) {
        heat_equation(thrust::cuda::par, U, V);
        ...
    }
}
```

```
int nx(...), ny(...), write_steps(...), compute_steps(...);

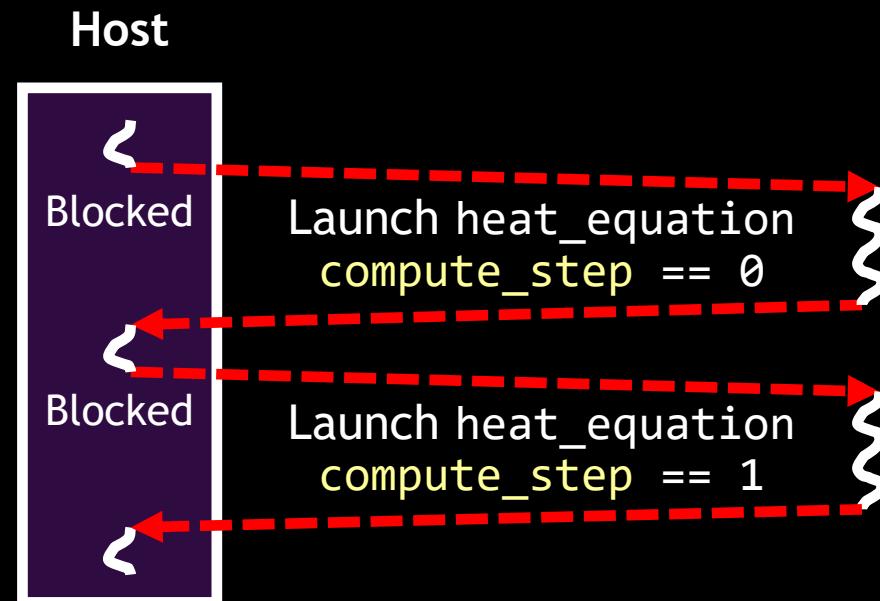
thrust::universal_vector<float> U_data(nx * ny), V_data(nx * ny);
cuda::std::mdspan U(thrust::raw_pointer_cast(U_data.data()), nx, ny);
cuda::std::mdspan V(thrust::raw_pointer_cast(V_data.data()), nx, ny);

initialize_oven(thrust::cuda::par, U);

for (auto write_step : std::views::iota(0, write_steps)) {
    save_to_file(U);

    for (auto compute_step : std::views::iota(0, compute_steps)) {
        heat_equation(thrust::cuda::par, U, V);
        U.swap(V);
    }
}
```

```
compute_step = 0;  
heat_equation(thrust::cuda::par, U, V);  
U.swap(V);  
compute_step = 1;  
heat_equation(thrust::cuda::par, U, V);  
U.swap(V);
```



```
int nx(...), ny(...), write_steps(...), compute_steps(...);

thrust::universal_vector<float> U_data(nx * ny), V_data(nx * ny);
cuda::std::mdspan U(thrust::raw_pointer_cast(U_data.data()), nx, ny);
cuda::std::mdspan V(thrust::raw_pointer_cast(V_data.data()), nx, ny);

initialize_oven(thrust::cuda::par, U);

for (auto write_step : std::views::iota(0, write_steps)) {
    save_to_file(U);

    for (auto compute_step : std::views::iota(0, compute_steps)) {
        heat_equation(thrust::cuda::par, U, V);
        U.swap(V);
    }
}
```

```
int nx(...), ny(...), write_steps(...), compute_steps(...);

thrust::universal_vector<float> U_data(nx * ny), V_data(nx * ny);
cuda::std::mdspan U(thrust::raw_pointer_cast(U_data.data()), nx, ny);
cuda::std::mdspan V(thrust::raw_pointer_cast(V_data.data()), nx, ny);

initialize_oven(thrust::cuda::par_nosync, U);

for (auto write_step : std::views::iota(0, write_steps)) {
    save_to_file(U);

    for (auto compute_step : std::views::iota(0, compute_steps)) {
        heat_equation(thrust::cuda::par_nosync, U, V);
        U.swap(V);
    }
}
```

```
int nx(...), ny(...), write_steps(...), compute_steps(...);

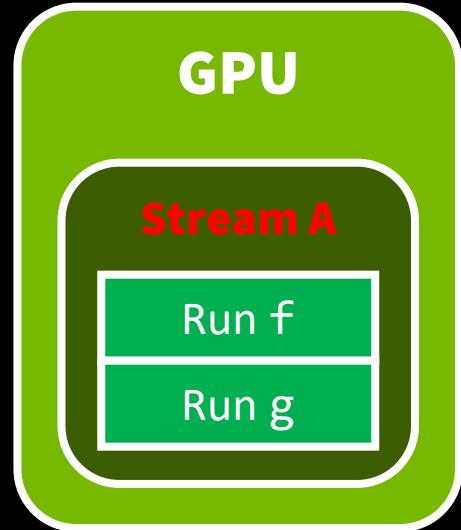
thrust::universal_vector<float> U_data(nx * ny), V_data(nx * ny);
cuda::std::mdspan U(thrust::raw_pointer_cast(U_data.data()), nx, ny);
cuda::std::mdspan V(thrust::raw_pointer_cast(V_data.data()), nx, ny);

initialize_oven(thrust::cuda::par_nosync, U);

for (auto write_step : std::views::iota(0, write_steps)) {
    cudaDeviceSynchronize();
    save_to_file(U);

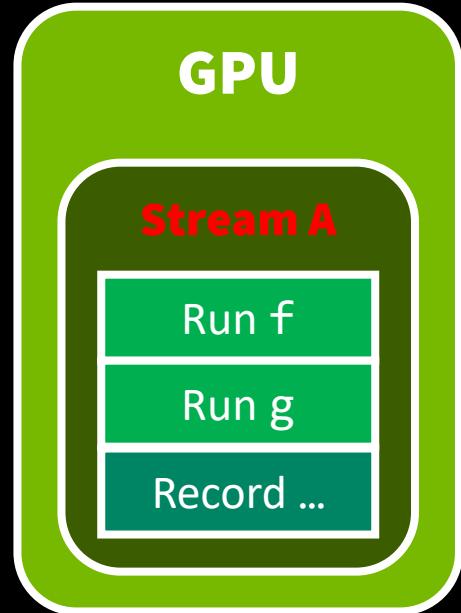
    for (auto compute_step : std::views::iota(0, compute_steps)) {
        heat_equation(thrust::cuda::par_nosync, U, V);
        U.swap(V);
    }
}
```

```
auto pol(thrust::cuda::par_nosync);  
  
cudax::stream A;  
  
thrust::for_each(pol.on(A), ..., f);  
thrust::for_each(pol.on(A), ..., g);
```



A stream is a chain of device work that is executed in order.

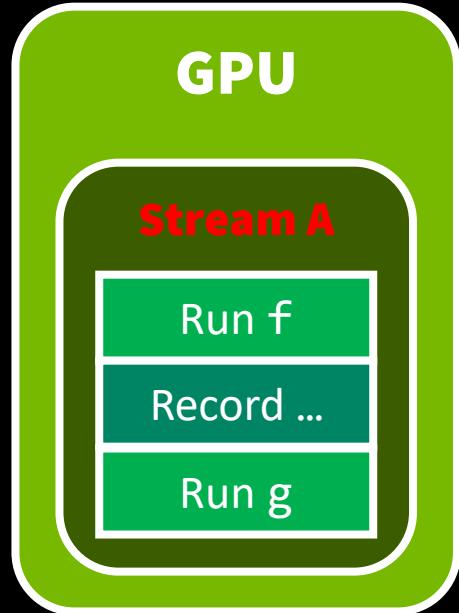
```
auto pol(thrust::cuda::par_nosync);  
  
cudax::stream A;  
  
thrust::for_each(pol.on(A), ..., f);  
thrust::for_each(pol.on(A), ..., g);  
A.wait();
```



A stream is a chain of device work that is executed in order.

They can record events that host code can sync with.

```
auto pol(thrust::cuda::par_nosync);  
  
cudax::stream A;  
  
thrust::for_each(pol.on(A), ..., f);  
A.wait();  
thrust::for_each(pol.on(A), ..., g);
```



A stream is a chain of device work that is executed in order.

They can record events that host code can sync with.

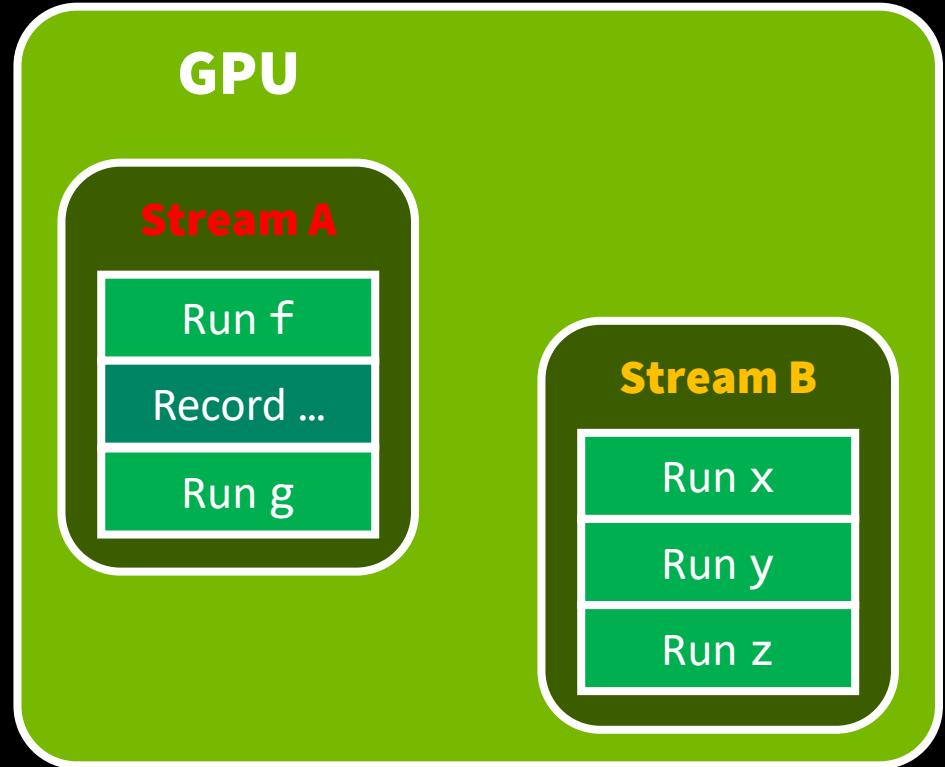
```
auto pol(thrust::cuda::par_nosync);
```

```
cudax::stream A;
```

```
thrust::for_each(pol.on(A), ..., f);  
A.wait();  
thrust::for_each(pol.on(A), ..., g);
```

```
cudax::stream B;
```

```
thrust::for_each(pol.on(B), ..., x);  
thrust::for_each(pol.on(B), ..., y);  
thrust::for_each(pol.on(B), ..., z);
```



```

auto pol(thrust::cuda::par_nosync);

cudax::stream A;

thrust::for_each(pol.on(A), ..., f);
A.wait();
thrust::for_each(pol.on(A), ..., g);

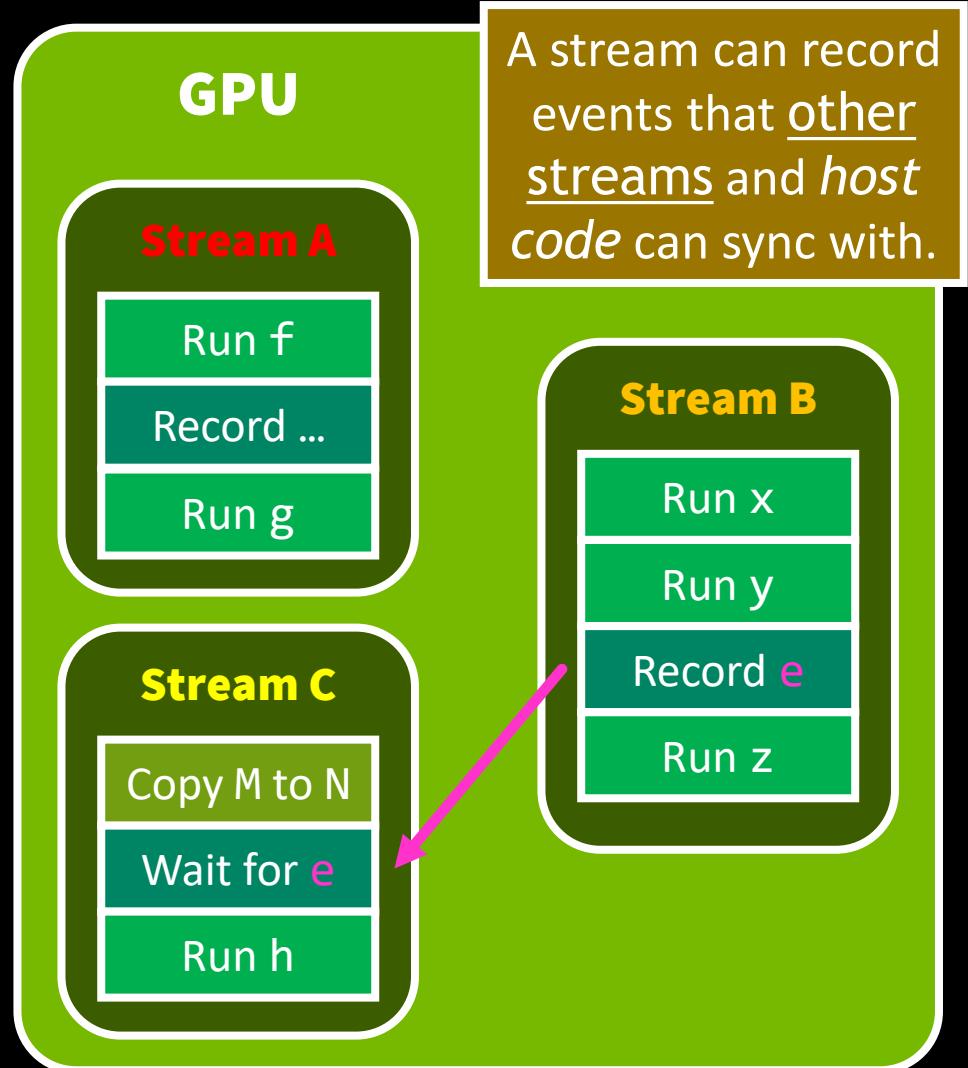
cudax::stream B;

thrust::for_each(pol.on(B), ..., x);
thrust::for_each(pol.on(B), ..., y);
auto e = B.record_event();
thrust::for_each(pol.on(B), ..., z);

cudax::stream C;

thrust::copy(pol.on(C), M, ..., N);
C.wait(e);
thrust::for_each(pol.on(C), ... h);

```



```

auto pol(thrust::cuda::par_nosync);

cudax::stream A(cudax::devices[0]);

thrust::for_each(pol.on(A), ..., f);
A.wait();
thrust::for_each(pol.on(A), ..., g);

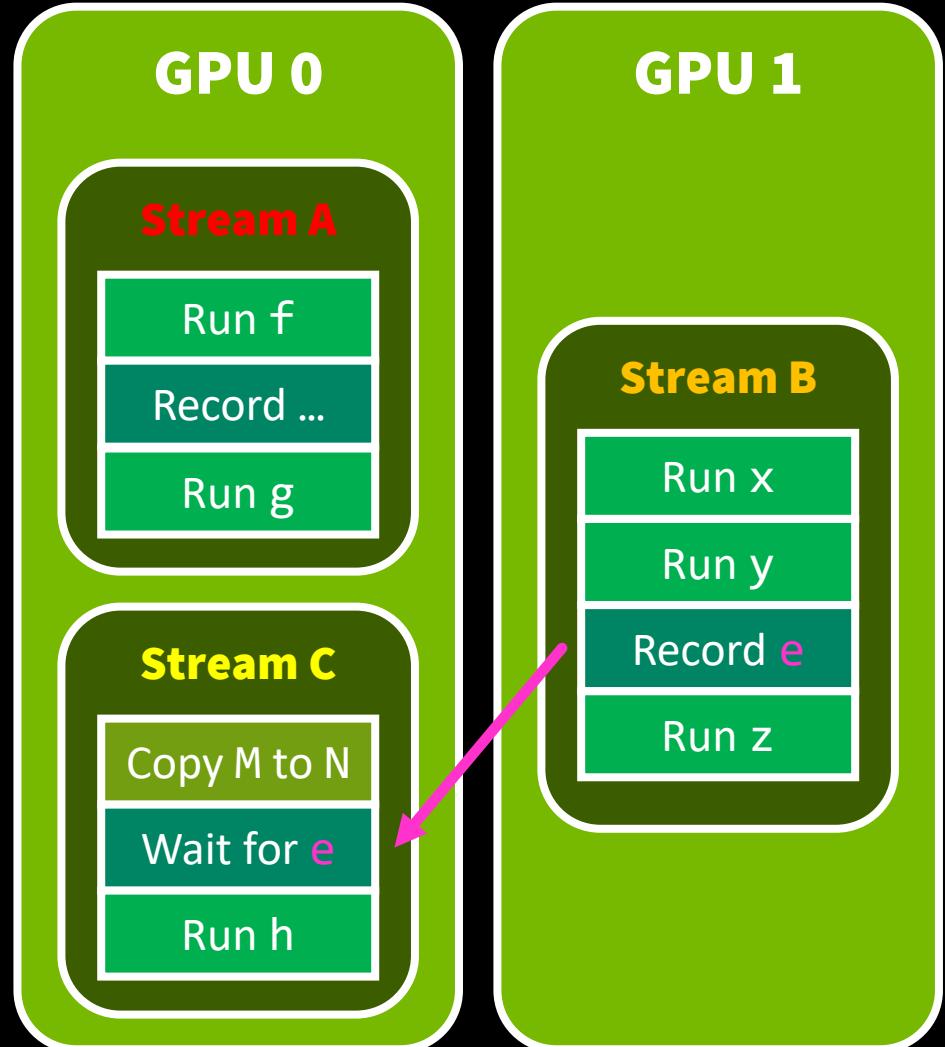
cudax::stream B(cudax::devices[1]);

thrust::for_each(pol.on(B), ..., x);
thrust::for_each(pol.on(B), ..., y);
auto e = B.record_event();
thrust::for_each(pol.on(B), ..., z);

cudax::stream C(cudax::devices[0]);

thrust::copy(pol.on(C), M, ..., N);
C.wait(e);
thrust::for_each(pol.on(C), ... h);

```



```
int nx(...), ny(...), write_steps(...), compute_steps(...);

thrust::universal_vector<float> U_data(nx * ny), V_data(nx * ny);
cuda::std::mdspan U(...), V(...);

initialize_oven(thrust::cuda::par_nosync, U);

for (auto write_step : std::views::iota(0, write_steps)) {
    cudaDeviceSynchronize();
    save_to_file(U);

    for (auto compute_step : std::views::iota(0, compute_steps)) {
        heat_equation(thrust::cuda::par_nosync, U, V);
        U.swap(V);
    }
}
```

```
int nx(...), ny(...), write_steps(...), compute_steps(...);

cudax::stream stream;
auto policy(thrust::cuda::par_nosync.on(stream.get()));

thrust::universal_vector<float> U_data(nx * ny), V_data(nx * ny);
cuda::std::mdspan U(...), V(...);

initialize_oven(policy, U);

for (auto write_step : std::views::iota(0, write_steps)) {
    stream.wait();
    save_to_file(U);

    for (auto compute_step : std::views::iota(0, compute_steps)) {
        heat_equation(policy, U, V);
        U.swap(V);
    }
}
```

```
auto policy = thrust::cuda::par_nosync;
thrust::universal_vector A(...), B(...);

thrust::for_each_n(
    policy, A.begin(), A.size(), f0);
// Doesn't return anything.

auto it1 = thrust::transform(policy,
    A.begin(), A.end(), B.begin(), f1);
// Returns an iterator to the end of the output (B).
// Doesn't depend on the computation; it's just B.begin() + A.size().
```

```
auto policy = thrust::cuda::par_nosync;  
thrust::universal_vector A(...), B(...);  
  
thrust::for_each_n(  
    policy, A.begin(), A.size(), f0);  
// Doesn't return anything.  
  
auto it1 = thrust::transform(policy,  
    A.begin(), A.end(), B.begin(), f1);  
// Returns an iterator to the end of the output (B).  
// Doesn't depend on the computation; it's just B.begin() + A.size().
```

Some Thrust algorithms either:

- Return a result that depends on the computation, or
- Allocate temporary storage for the computation.

These interfaces will always block.

```
auto policy = thrust::cuda::par_nosync;  
thrust::universal_vector A(...), B(...);
```

```
thrust::for_each_n(  
    policy, A.begin(), A.size(), f0);  
// Doesn't return anything.
```

```
auto it1 = thrust::transform(policy,  
    A.begin(), A.end(), B.begin(), f1);
```

```
// Returns an iterator to the end of the output (B).  
// Doesn't depend on the computation; it's just B.begin() + A.size().
```

```
auto it2 = thrust::inclusive_scan(policy, A.begin(), A.end(), B.begin(), f2);  
// Also returns the end of the output, but needs temporary storage.
```

Some Thrust algorithms either:

- Return a result that depends on the computation, or
- Allocate temporary storage for the computation.

These interfaces will always block.

```
auto policy = thrust::cuda::par_nosync;  
thrust::universal_vector A(...), B(...);
```

```
thrust::for_each_n(  
    policy, A.begin(), A.size(), f0);  
// Doesn't return anything.
```

```
auto it1 = thrust::transform(policy,  
    A.begin(), A.end(), B.begin(), f1);
```

```
// Returns an iterator to the end of the output (B).  
// Doesn't depend on the computation; it's just B.begin() + A.size().
```

```
auto it2 = thrust::inclusive_scan(policy, A.begin(), A.end(), B.begin(), f2);  
// Also returns the end of the output, but needs temporary storage.
```

```
auto it3 = thrust::copy_if(policy, A.begin(), A.end(), B.begin(), f3);  
// Returns the end of the output, but it depends on the computation.
```

Some Thrust algorithms either:

- Return a result that depends on the computation, or
- Allocate temporary storage for the computation.

These interfaces will always block.

```
auto policy = thrust::cuda::par_nosync;  
thrust::universal_vector A(...), B(...);
```

```
thrust::for_each_n(  
    policy, A.begin(), A.size(), f0);  
// Doesn't return anything.
```

```
auto it1 = thrust::transform(policy,  
    A.begin(), A.end(), B.begin(), f1);  
// Returns an iterator to the end of the output (B).  
// Doesn't depend on the computation; it's just B.begin() + A.size().
```

```
auto it2 = thrust::inclusive_scan(policy, A.begin(), A.end(), B.begin(), f2);  
// Also returns the end of the output, but needs temporary storage.
```

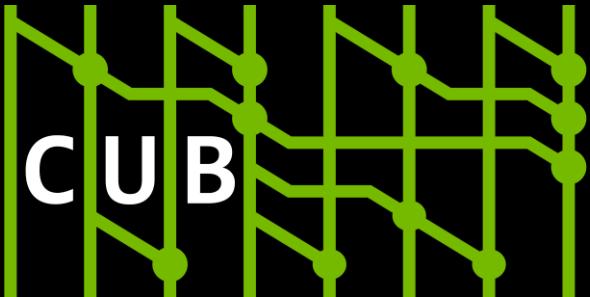
```
auto it3 = thrust::copy_if(policy, A.begin(), A.end(), B.begin(), f3);  
// Returns the end of the output, but it depends on the computation.
```

```
auto r4 = thrust::reduce(policy, A.begin(), A.end(), T{}, f4);  
auto r5 = thrust::find_if(policy, A.begin(), A.end(), f5);  
// Returns the result of the computation itself.
```

Some Thrust algorithms either:

- Return a result that depends on the computation, or
- Allocate temporary storage for the computation.

These interfaces will always block.



# The CUDA C++ Algorithm Authoring Toolkit

<https://nvidia.github.io/cccl/cub>

Learn More: [S72575 How You Should Write a CUDA C++ Kernel](#)

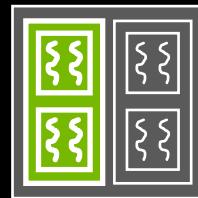
## Device-Wide Kernels

- `cub::DeviceSegmentedReduce`
- `cub::DeviceHistogram`
- `cub::DeviceRadixSort`
- `cub::DeviceSelect`
- ...



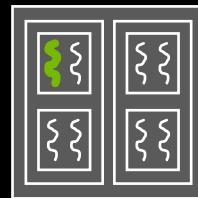
## Warp-Wide Cooperative Primitives

- `cub::WarpReduce`
- `cub::WarpExchange`
- `cub::WarpMergeSort`
- `cub::WarpLoad/Store`
- ...



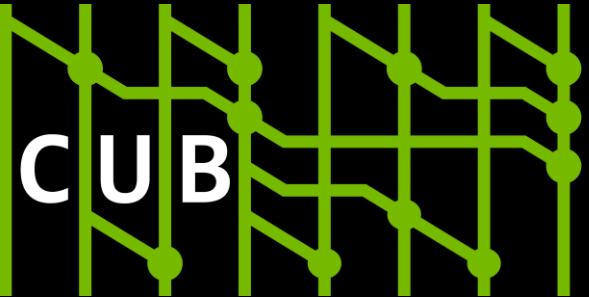
## Block-Wide Cooperative Primitives

- `cub::BlockReduce`
- `cub::BlockExchange`
- `cub::BlockRunLengthDecode`
- `cub::BlockLoad/Store`
- ...



## Per-Thread Primitives

- `cub::ThreadReduce`
- `cub::ThreadScan`
- `cub::MergePathSearch`
- `cub::ThreadLoad/Store`
- ...



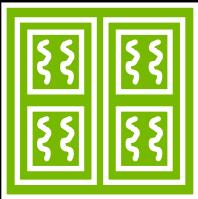
# The CUDA C++ Algorithm Authoring Toolkit

<https://nvidia.github.io/cccl/cub>

Learn More: [S72575 How You Should Write a CUDA C++ Kernel](#)

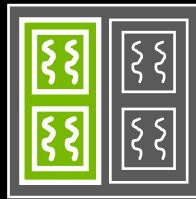
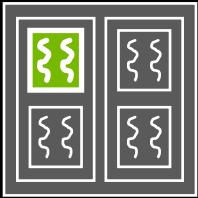
## Device-Wide Kernels

- `cub::DeviceSegmentedReduce`
- `cub::DeviceHistogram`
- `cub::DeviceRadixSort`
- `cub::DeviceSelect`
- ...



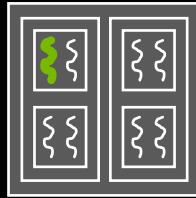
## Warp-Wide Cooperative Primitives

- `cub::WarpReduce`
- `cub::WarpExchange`
- `cub::WarpMergeSort`
- `cub::WarpLoad/Store`
- ...



## Block-Wide Cooperative Primitives

- `cub::BlockReduce`
- `cub::BlockExchange`
- `cub::BlockRunLengthDecode`
- `cub::BlockLoad/Store`
- ...



## Per-Thread Primitives

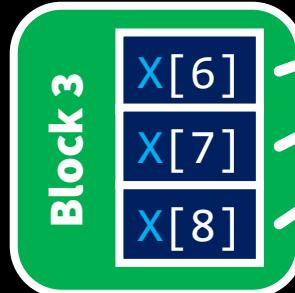
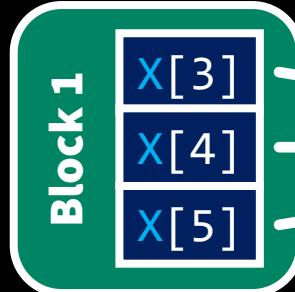
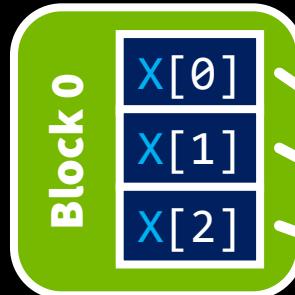
- `cub::ThreadReduce`
- `cub::ThreadScan`
- `cub::MergePathSearch`
- `cub::ThreadLoad/Store`
- ...

```
thrust::universal_vector X(...);

auto res = thrust::reduce(
    thrust::cuda::par,
    X.begin(), X.end(),
    T{}, cuda::maximum{});
```

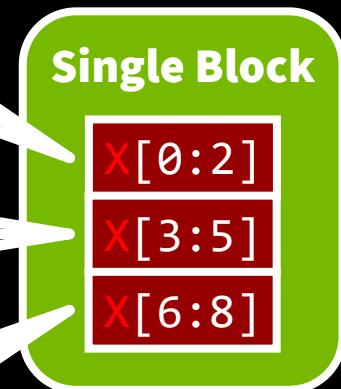
## 1<sup>st</sup> Kernel

$O(\text{input}) \rightarrow O(\text{blocks})$



## 2<sup>nd</sup> Kernel

$O(\text{blocks}) \rightarrow 1$



```
thrust::universal_vector X(...);  
  
auto res = thrust::reduce(  
    thrust::cuda::par,  
    X.begin(), X.end(),  
    T{}, cuda::maximum{});
```

$O(\text{blocks})$  temporary storage required for block sums.

```
thrust::universal_vector X(...);  
thrust::universal_vector res(1);  
  
thrust::universal_vector X(...);  
  
auto res = thrust::reduce(  
    thrust::cuda::par,  
    X.begin(), X.end(),  
    T{}, cuda::maximum{});
```

```
thrust::universal_vector X(...);  
thrust::universal_vector res(1);  
  
// Determine temporary storage size.  
// Doesn't Launch anything.  
int tmp_size = 0;  
cub::DeviceReduce::Reduce(  
    nullptr,  
    tmp_size,  
    X.begin(), res.begin(), X.size(),  
    cuda::maximum{}, T{});  
  
thrust::universal_vector X(...);  
  
auto res = thrust::reduce(  
    thrust::cuda::par,  
    X.begin(), X.end(),  
    T{}, cuda::maximum{});
```

```
thrust::universal_vector X(...);  
  
auto res = thrust::reduce(  
    thrust::cuda::par,  
    X.begin(), X.end(),  
    T{}, cuda::maximum{});
```

```
thrust::universal_vector X(...);  
thrust::universal_vector res(1);  
  
// Determine temporary storage size.  
// Doesn't Launch anything.  
int tmp_size = 0;  
cub::DeviceReduce::Reduce(  
    nullptr,  
    tmp_size,  
    X.begin(), res.begin(), X.size(),  
    cuda::maximum{}, T{});
```

```
thrust::universal_vector X(...);  
thrust::universal_vector res(1);  
  
// Determine temporary storage size.  
// Doesn't Launch anything.  
int tmp_size = 0;  
cub::DeviceReduce::Reduce(  
    nullptr,  
    tmp_size,  
    X.begin(), res.begin(), X.size(),  
    cuda::maximum{}, T{});  
  
thrust::device_vector<cuda::std::byte>  
tmp(tmp_size);
```

```
thrust::universal_vector X(...);
thrust::universal_vector res(1);

// Determine temporary storage size.
// Doesn't Launch anything.
int tmp_size = 0;
cub::DeviceReduce::Reduce(
    nullptr,
    tmp_size,
    X.begin(), res.begin(), X.size(),
    cuda::maximum{}, T{});

thrust::device_vector<cuda::std::byte>
tmp(tmp_size);

// Launch kernel.
cub::DeviceReduce::Reduce(
    thrust::raw_pointer_cast(tmp.data()),
    tmp_size,
    X.begin(), res.begin(), X.size(),
    cuda::maximum{}, T{});
```

```

void heat_equation(auto policy,
                    cuda::std::mdspan<float const, cuda::std::dims<2>> U,
                    cuda::std::mdspan<float, cuda::std::dims<2>> V) {
    thrust::for_each_n(policy, thrust::make_counting_iterator(0), U.size(),
                      [U, V] __host__ __device__ (int xy) {
                        int x = xy / U.extent(1);
                        int y = xy % U.extent(1);

                        if (x > 0 && y > 0 && x < I.extent(0) - 1 && y < U.extent(1) - 1) {
                            auto d2tdx2 = U(x, y - 1) - U(x, y) * 2 + U(x, y + 1);
                            auto d2tdy2 = U(x - 1, y) - U(x, y) * 2 + U(x + 1, y);

                            V(x, y) = U(x, y) + 0.2f * (d2tdx2 + d2tdy2);
                        } else {
                            V(x, y) = U(x, y);
                        }
                    });
}

```

```

void heat_equation(cudax::stream_ref stream,
                    cuda::std::mdspan<float const, cuda::std::dims<2>> U,
                    cuda::std::mdspan<float, cuda::std::dims<2>> V) {
    cub::DeviceFor::ForEachN(thrust::make_counting_iterator(0), U.size(),
        [U, V] __host__ __device__ (int xy) {
            int x = xy / U.extent(1);
            int y = xy % U.extent(1);

            if (x > 0 && y > 0 && x < I.extent(0) - 1 && y < U.extent(1) - 1) {
                auto d2tdx2 = U(x, y - 1) - U(x, y) * 2 + U(x, y + 1);
                auto d2tdy2 = U(x - 1, y) - U(x, y) * 2 + U(x + 1, y);

                V(x, y) = U(x, y) + 0.2f * (d2tdx2 + d2tdy2);
            } else {
                V(x, y) = U(x, y);
            }
        }, stream);
}

```

```

void heat_equation(cudax::stream_ref stream,
                    cuda::std::mdspan<float const, cuda::std::dims<2>> U,
                    cuda::std::mdspan<float, cuda::std::dims<2>> V) {
    cub::DeviceFor::ForEachN(thrust::make_counting_iterator(0), U.size(),
        [U, V] __host__ __device__ (int xy) {
            int x = xy / U.extent(1);
            int y = xy % U.extent(1);

            if (x > 0 && y > 0 && x < I.extent(0) - 1 && y < U.extent(1) - 1) {
                auto d2tdx2 = U(x, y - 1) - U(x, y) * 2 + U(x, y + 1);
                auto d2tdy2 = U(x - 1, y) - U(x, y) * 2 + U(x + 1, y);

                V(x, y) = U(x, y) + 0.2f * (d2tdx2 + d2tdy2);
            } else {
                V(x, y) = U(x, y);
            }
        }, stream);
}

```

```

void heat_equation(cudax::stream_ref stream,
                    cuda::std::mdspan<float const, cuda::std::dims<2>> U,
                    cuda::std::mdspan<float, cuda::std::dims<2>> V) {
    cub::DeviceFor::ForEachN(thrust::make_counting_iterator(0), U.size(),
        [U, V] __host__ __device__ (int xy) {
            int x = xy / U.extent(1);
            int y = xy % U.extent(1);

            if (x > 0 && y > 0 && x < I.extent(0) - 1 && y < U.extent(1) - 1) {
                auto d2tdx2 = U(x, y - 1) - U(x, y) * 2 + U(x, y + 1);
                auto d2tdy2 = U(x - 1, y) - U(x, y) * 2 + U(x + 1, y);

                V(x, y) = U(x, y) + 0.2f * (d2tdx2 + d2tdy2);
            } else {
                V(x, y) = U(x, y);
            }
        }, stream);
}

```

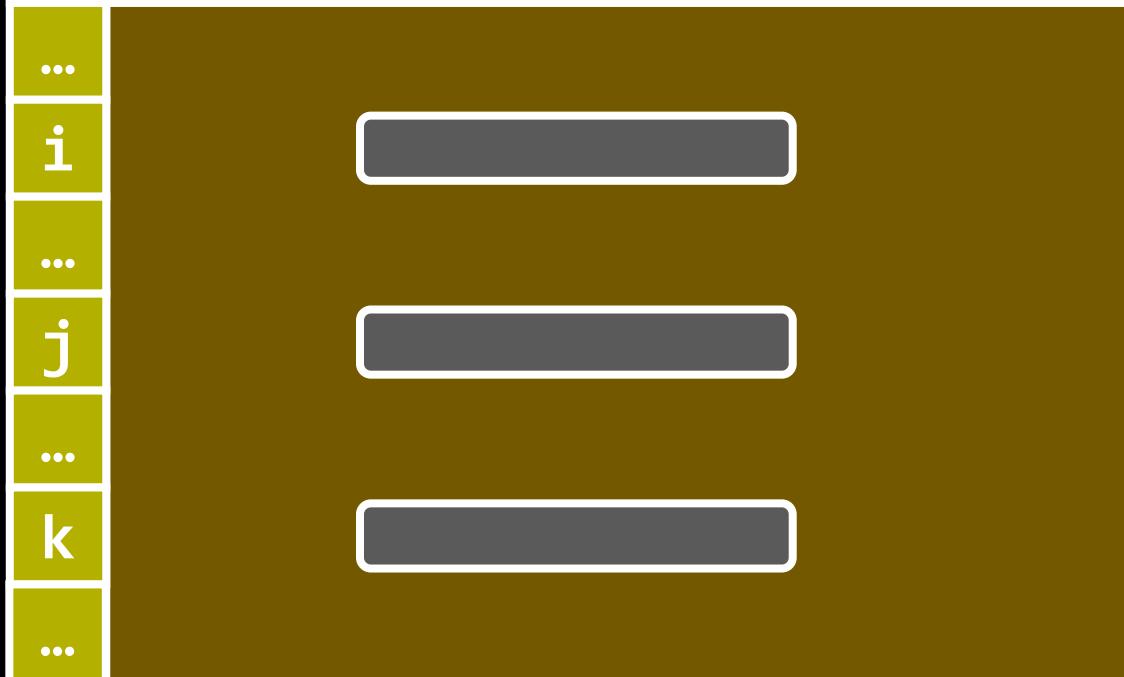
```

void heat_equation(cudax::stream_ref stream,
                    cuda::std::mdspan<float const, cuda::std::dims<2>> U,
                    cuda::std::mdspan<float, cuda::std::dims<2>> V) {
    cub::DeviceFor::ForEachInExtents(U.extents(),
        [U, V] __host__ __device__ (int idx, int x, int y) {
            if (x > 0 && y > 0 && x < I.extent(0) - 1 && y < U.extent(1) - 1) {
                auto d2tdx2 = U(x, y - 1) - U(x, y) * 2 + U(x, y + 1);
                auto d2tdy2 = U(x - 1, y) - U(x, y) * 2 + U(x + 1, y);

                V(x, y) = U(x, y) + 0.2f * (d2tdx2 + d2tdy2);
            } else {
                V(x, y) = U(x, y);
            }
        }, stream);
}

```

# thrust::universal\_vector



## Host Memory

## Device Memory

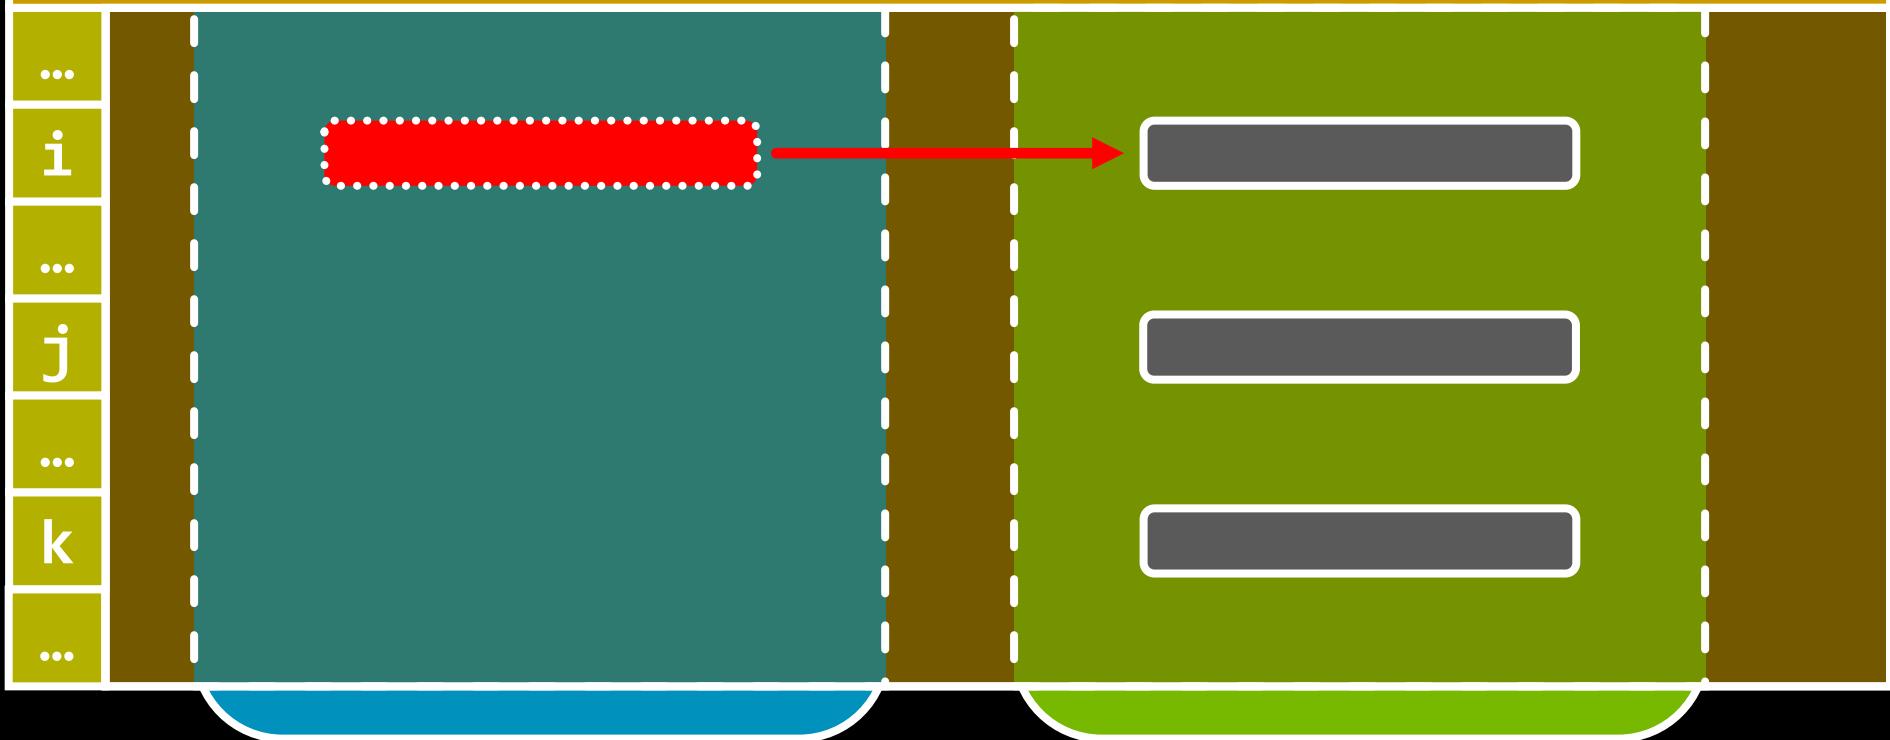
`thrust::universal_vector`



## Host Memory

## Device Memory

`thrust::universal_vector`



## Host Memory

## Device Memory

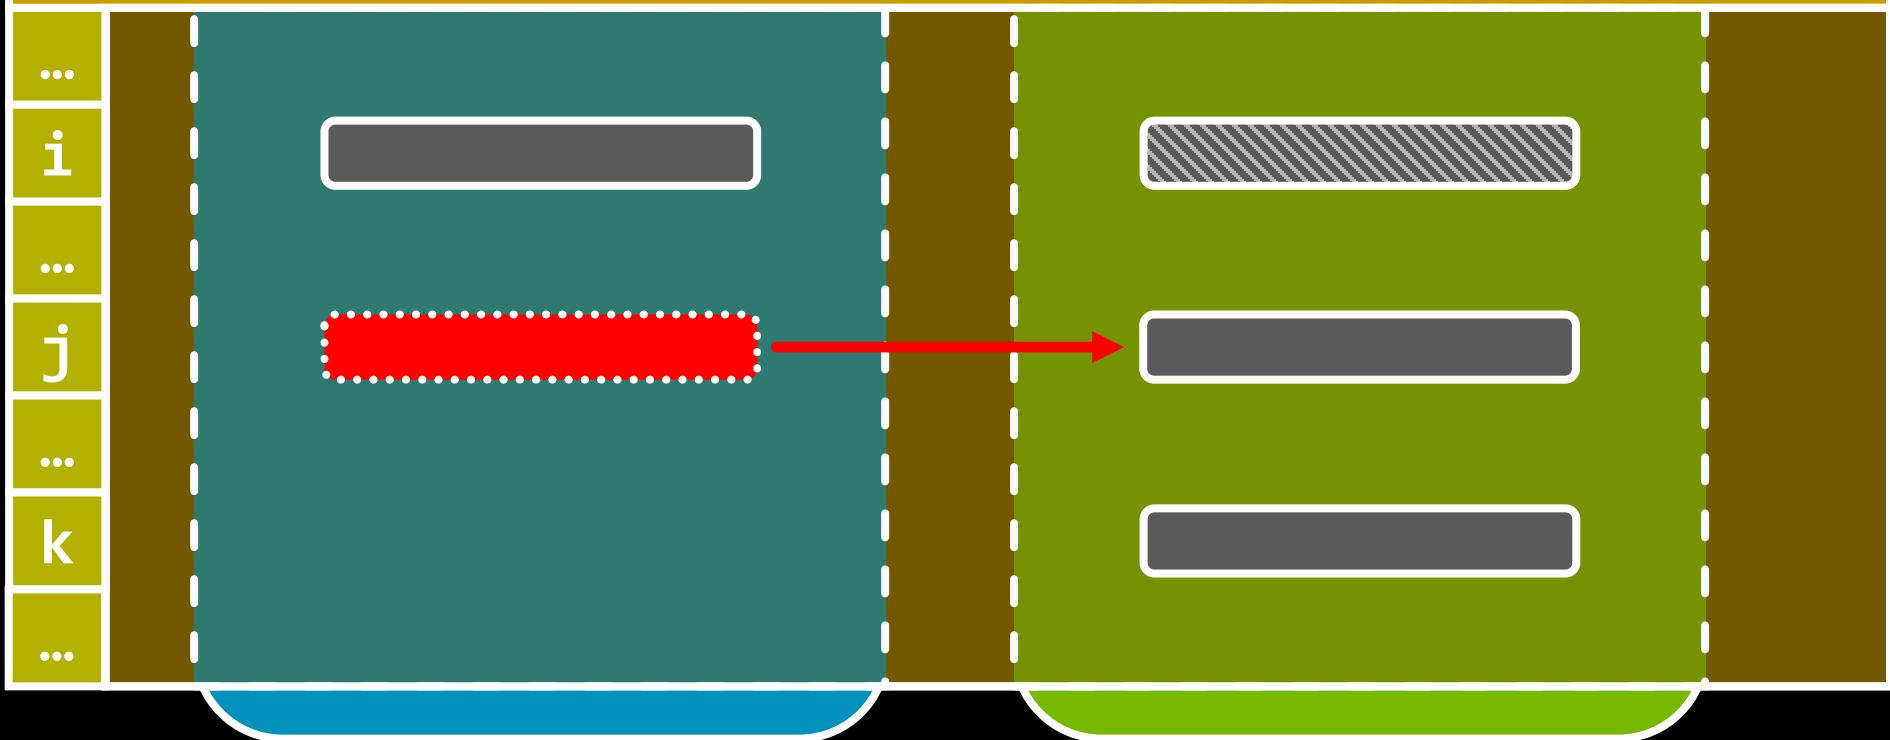
`thrust::universal_vector`



## Host Memory

## Device Memory

`thrust::universal_vector`



## Host Memory

## Device Memory

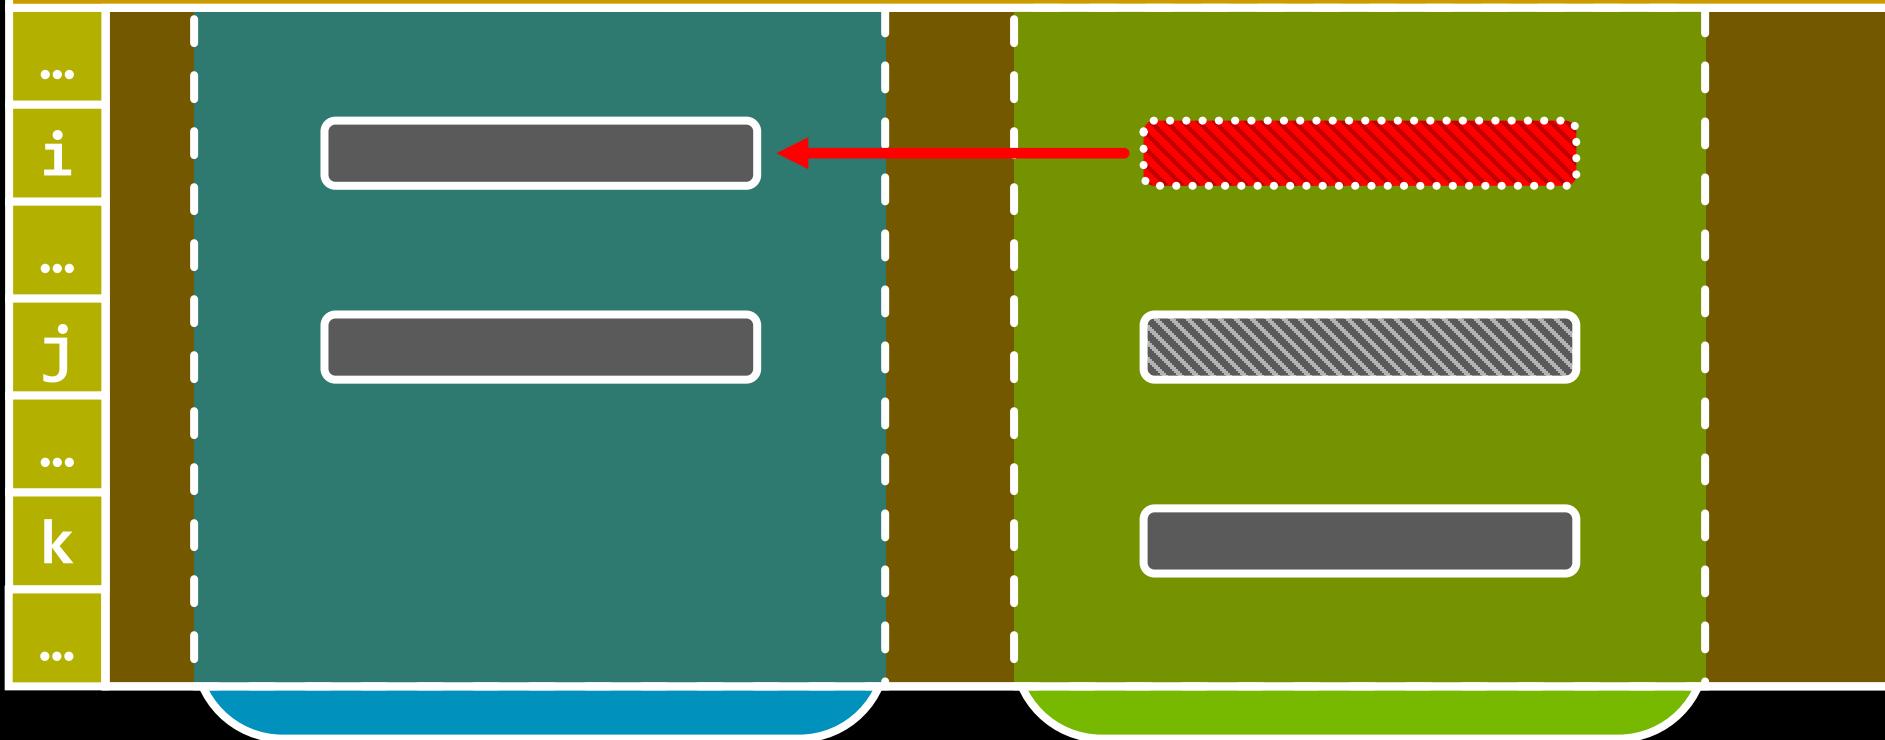
`thrust::universal_vector`



## Host Memory

## Device Memory

`thrust::universal_vector`



## Host Memory

## Device Memory

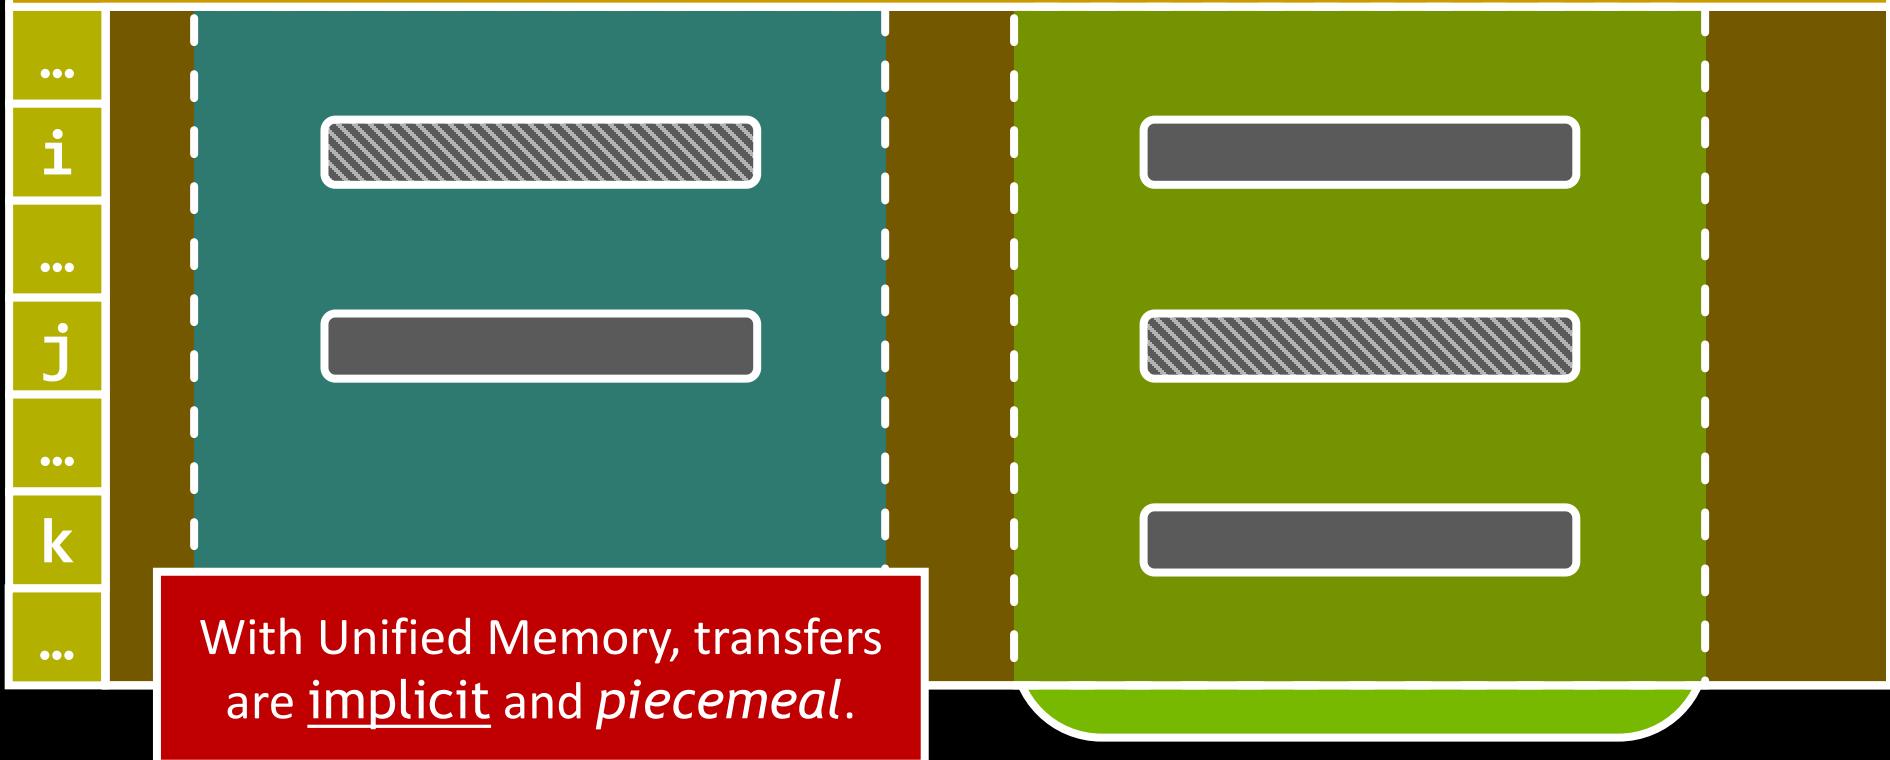
`thrust::universal_vector`



## Host Memory

## Device Memory

`thrust::universal_vector`



Thrust Container	Accessible in...	Transfers are...
universal_vector	Host & device code	Implicit upon access

Thrust Container	Accessible in...	Transfers are...
<code>universal_vector</code>	Host & device code	Implicit upon access
<code>host_vector</code>	Host code only	
<code>device_vector</code>	Device code only	

Thrust Container	Accessible in...	Transfers are...
universal_vector	Host & device code	Implicit upon access
host_vector	Host code only	
device_vector	Device code only	Explicit via <code>thrust::copy</code>

```
int nx(...), ny(...), write_steps(...), compute_steps(...);

cudax::stream stream;
auto policy(thrust::cuda::par_nosync.on(stream.get()));

thrust::universal_vector<float> U_data(nx * ny), V_data(nx * ny);
cuda::std::mdspan U(...), V(...);

initialize_oven(stream, U);

for (auto write_step : std::views::iota(0, write_steps)) {
    stream.wait();
    save_to_file(U);

    for (auto compute_step : std::views::iota(0, compute_steps)) {
        heat_equation(stream, U, V);
        U.swap(V);
    }
}
```

```

int nx(...), ny(...), write_steps(...), compute_steps(...);

cudax::stream stream;
auto policy(thrust::cuda::par_nosync.on(stream.get()));

thrust::universal_vector<float> U_data(nx * ny), V_data(nx * ny);
cuda::std::mdspan U(...), V(...);

initialize_oven(stream, U);

for (auto write_step : std::views::iota(0, write_steps)) {
    stream.wait();
    save_to_file(U); // Implicit device to host transfer.

    for (auto compute_step : std::views::iota(0, compute_steps)) {
        heat_equation(stream, U, V);
        U.swap(V);
    }
}

```

```
int nx(...), ny(...), write_steps(...), compute_steps(...);

cudax::stream stream;
auto policy(thrust::cuda::par_nosync.on(stream.get()));

thrust::universal_vector<float> U_data(nx * ny), V_data(nx * ny);
thrust::host_vector<float> buf_data(nx * ny);
cuda::std::mdspan U(...), V(...), buf(...);

initialize_oven(stream, U);

for (auto write_step : std::views::iota(0, write_steps)) {
    thrust::copy_n(policy, U.data(), U.size(), buf); // Device to host.
    stream.wait();

    for (auto compute_step : std::views::iota(0, compute_steps)) {
        heat_equation(stream, U, V);
        U.swap(V);
    }

    save_to_file(buf); // No transfer.
}
```

```
int nx(...), ny(...), write_steps(...), compute_steps(...);

cudax::stream stream;
auto policy(thrust::cuda::par_nosync.on(stream.get()));

thrust::universal_vector<float> U_data(nx * ny), V_data(nx * ny);
thrust::host_vector<float> buf_data(nx * ny);
cuda::std::mdspan U(...), V(...), buf(...);

initialize_oven(stream, U);

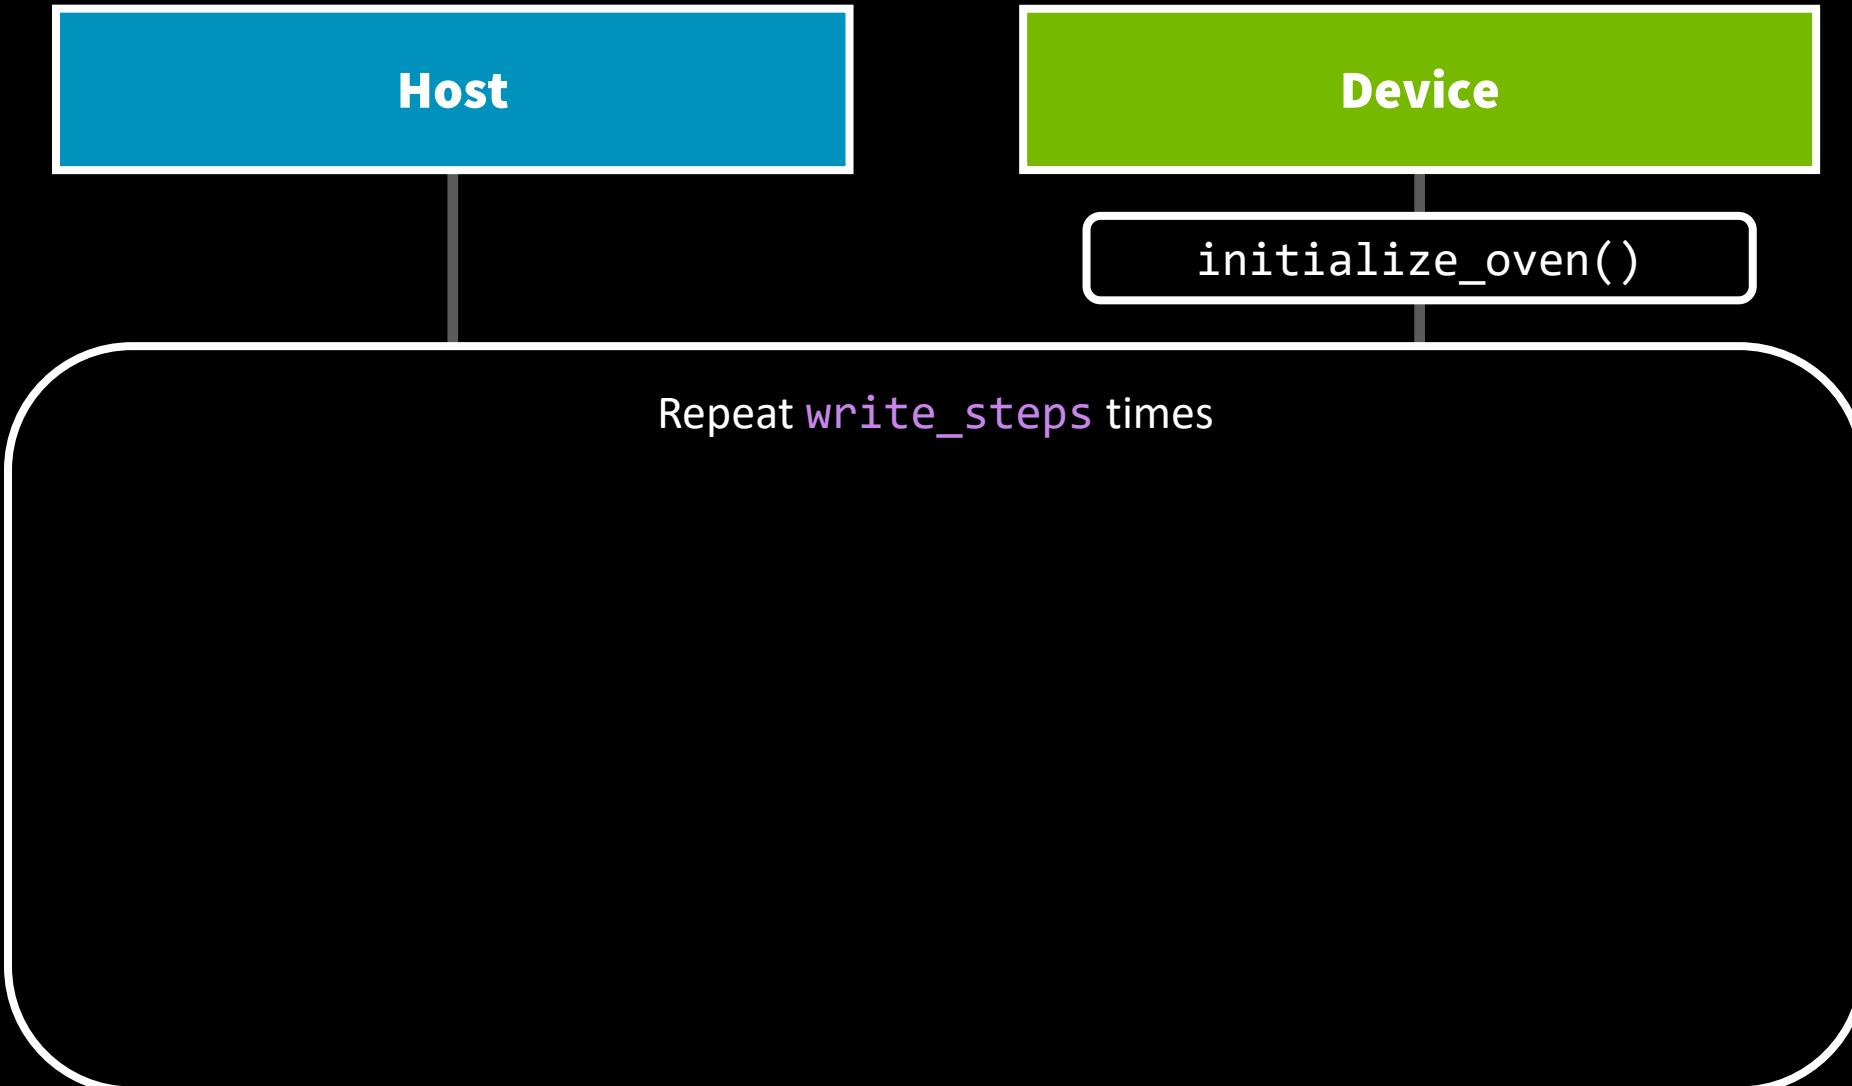
for (auto write_step : std::views::iota(0, write_steps)) {
    thrust::copy_n(policy, U.data(), U.size(), buf); // Device to host.
    stream.wait();

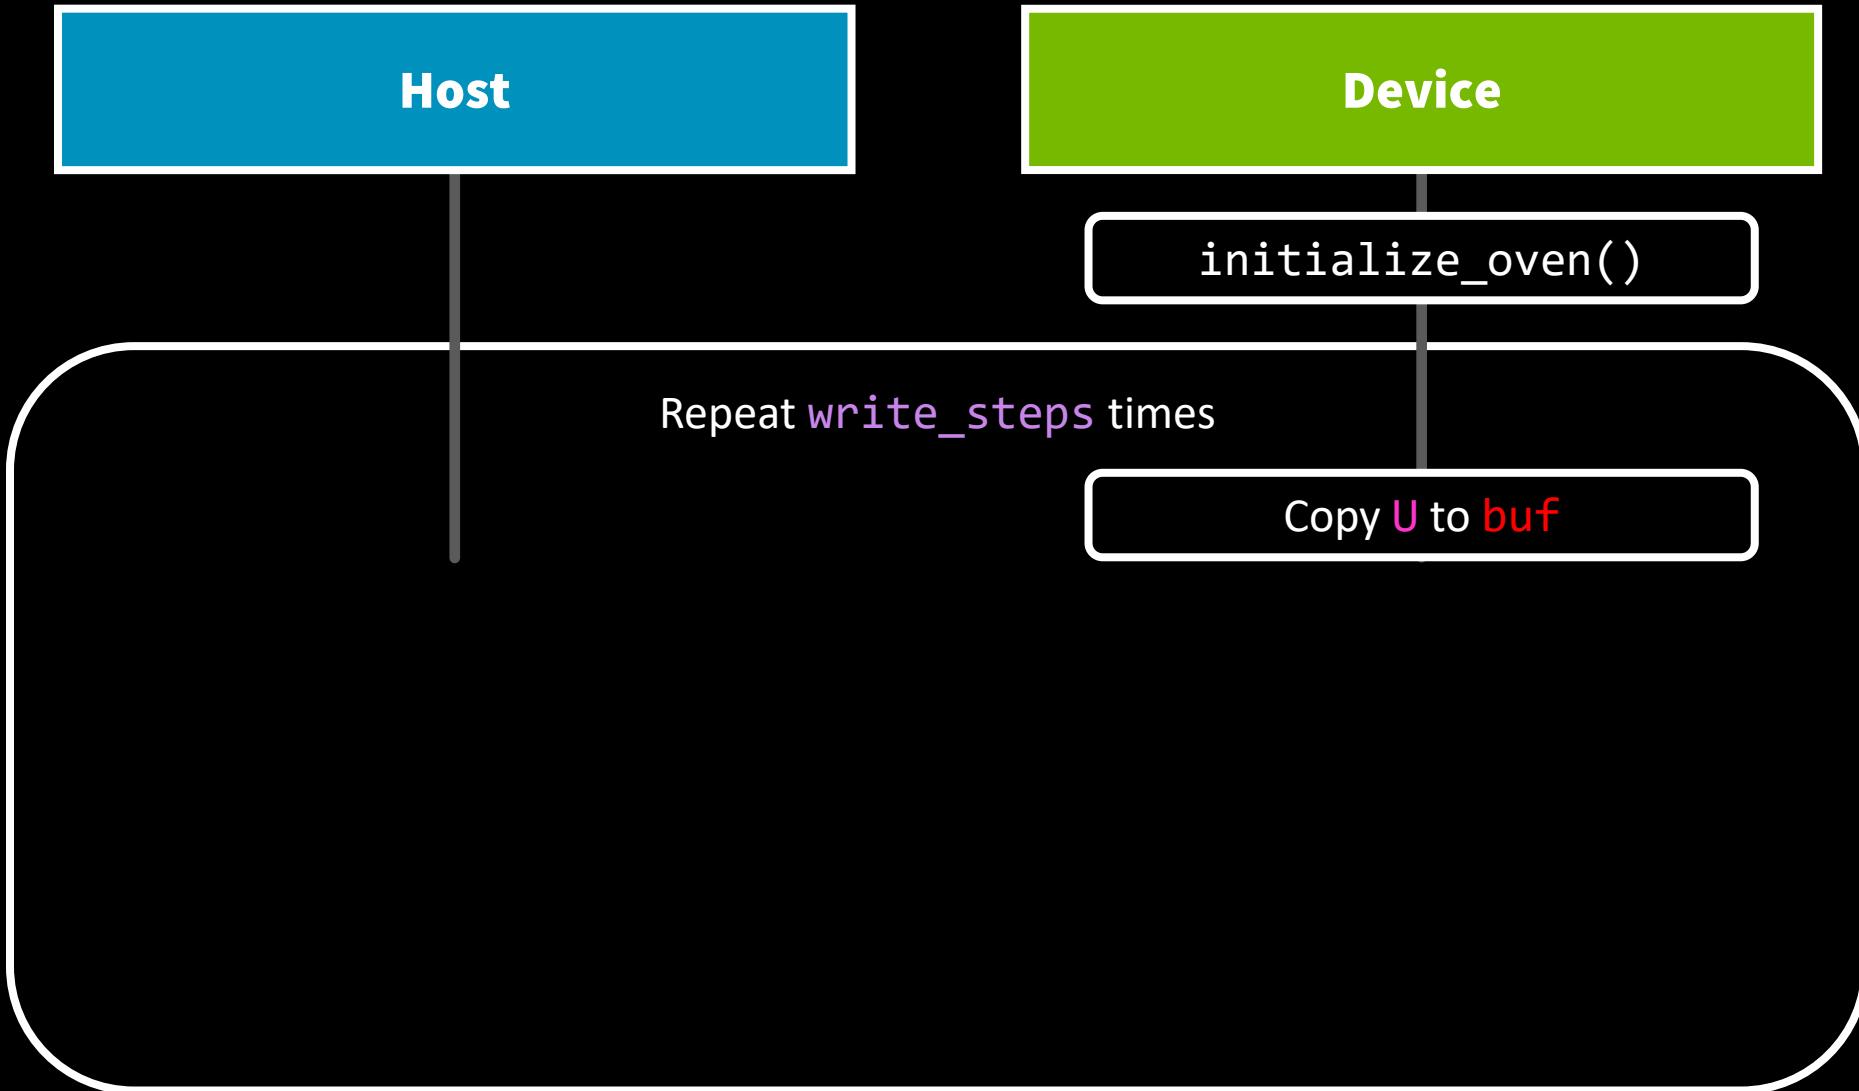
    for (auto compute_step : std::views::iota(0, compute_steps)) {
        heat_equation(stream, U, V);
        U.swap(V);
    }

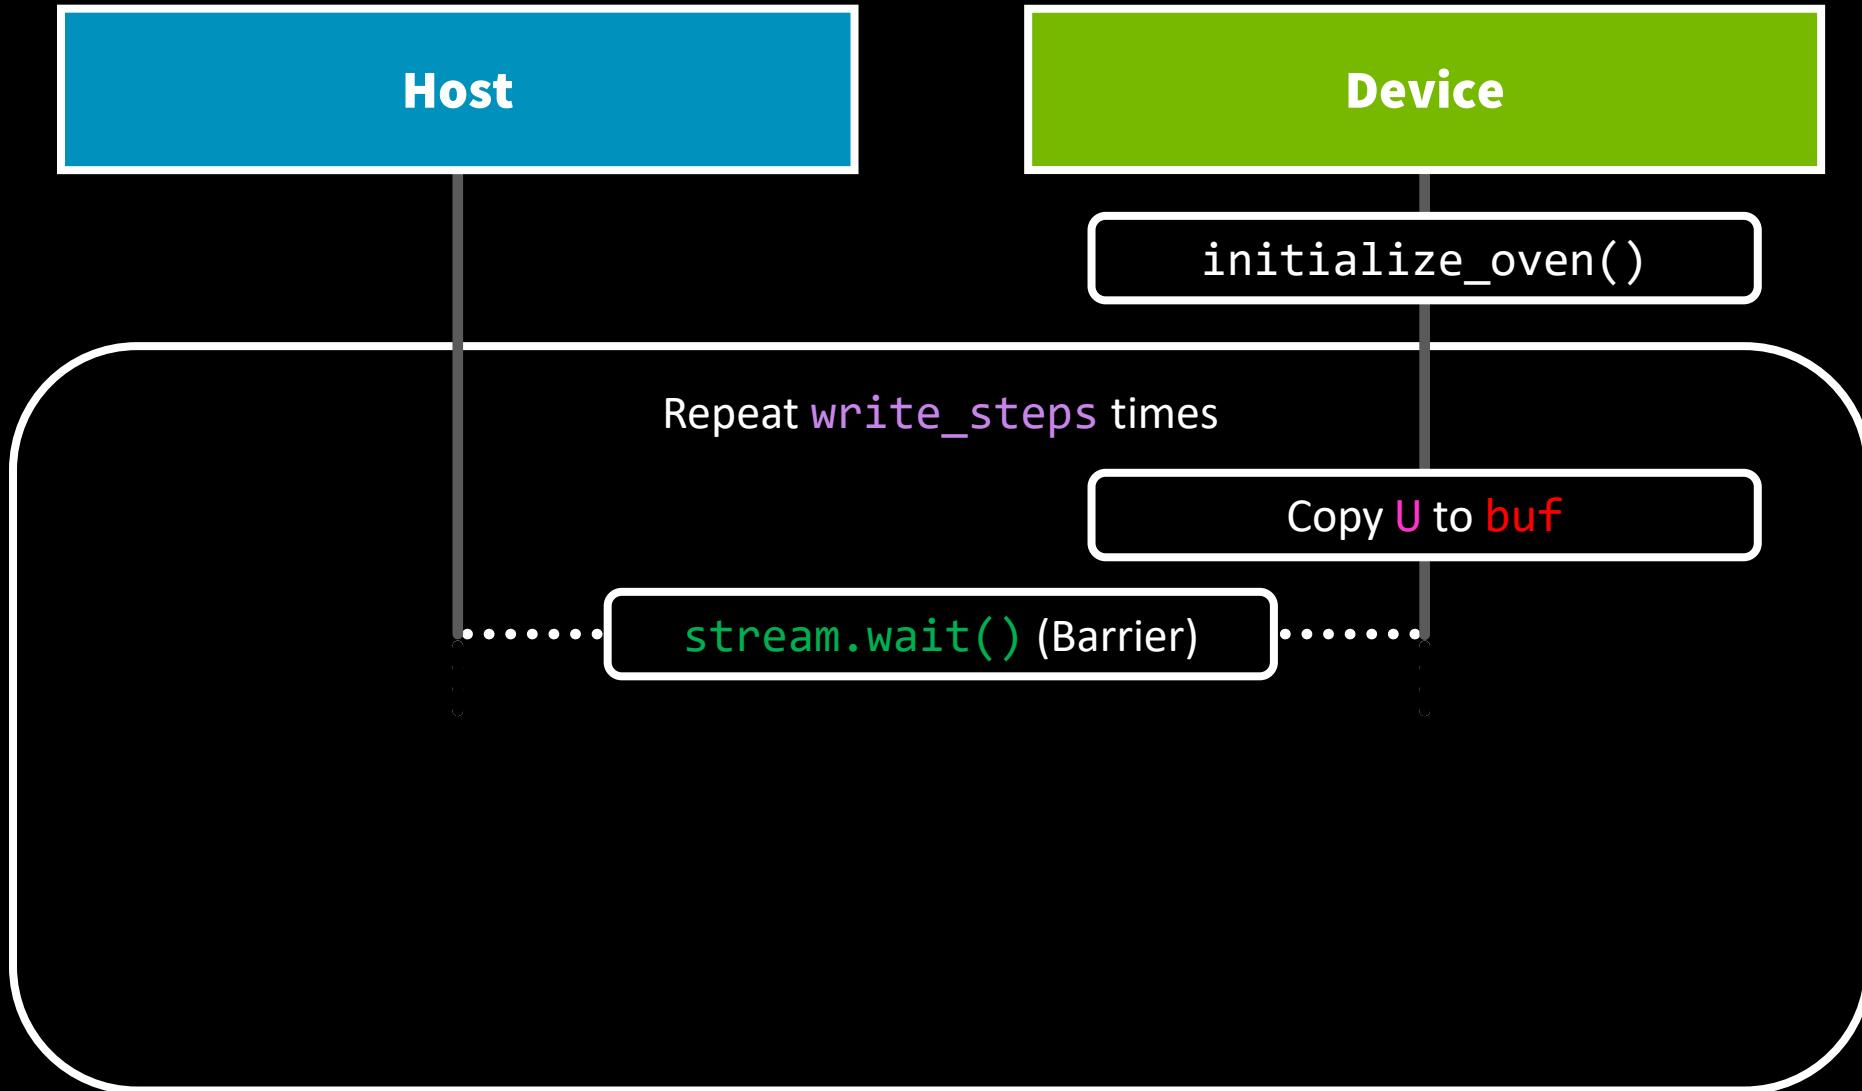
    save_to_file(buf); // No transfer.
}
```

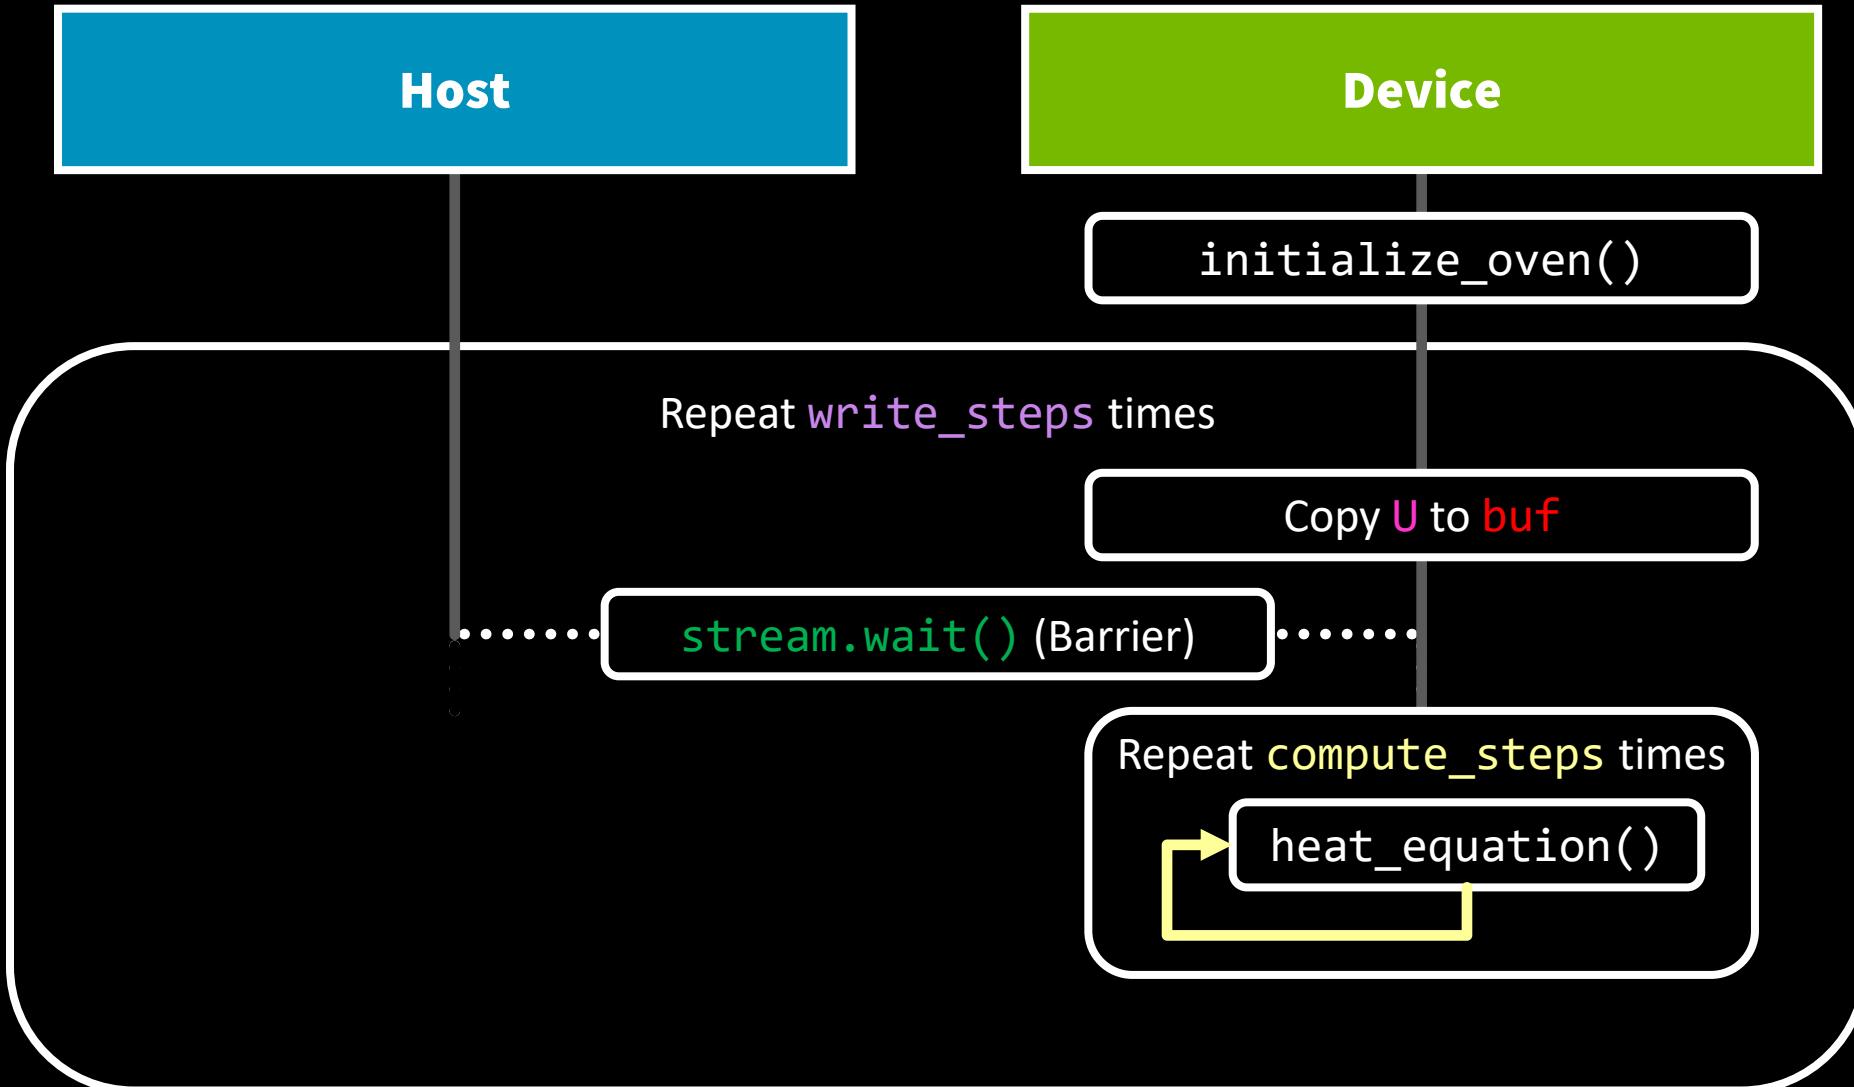


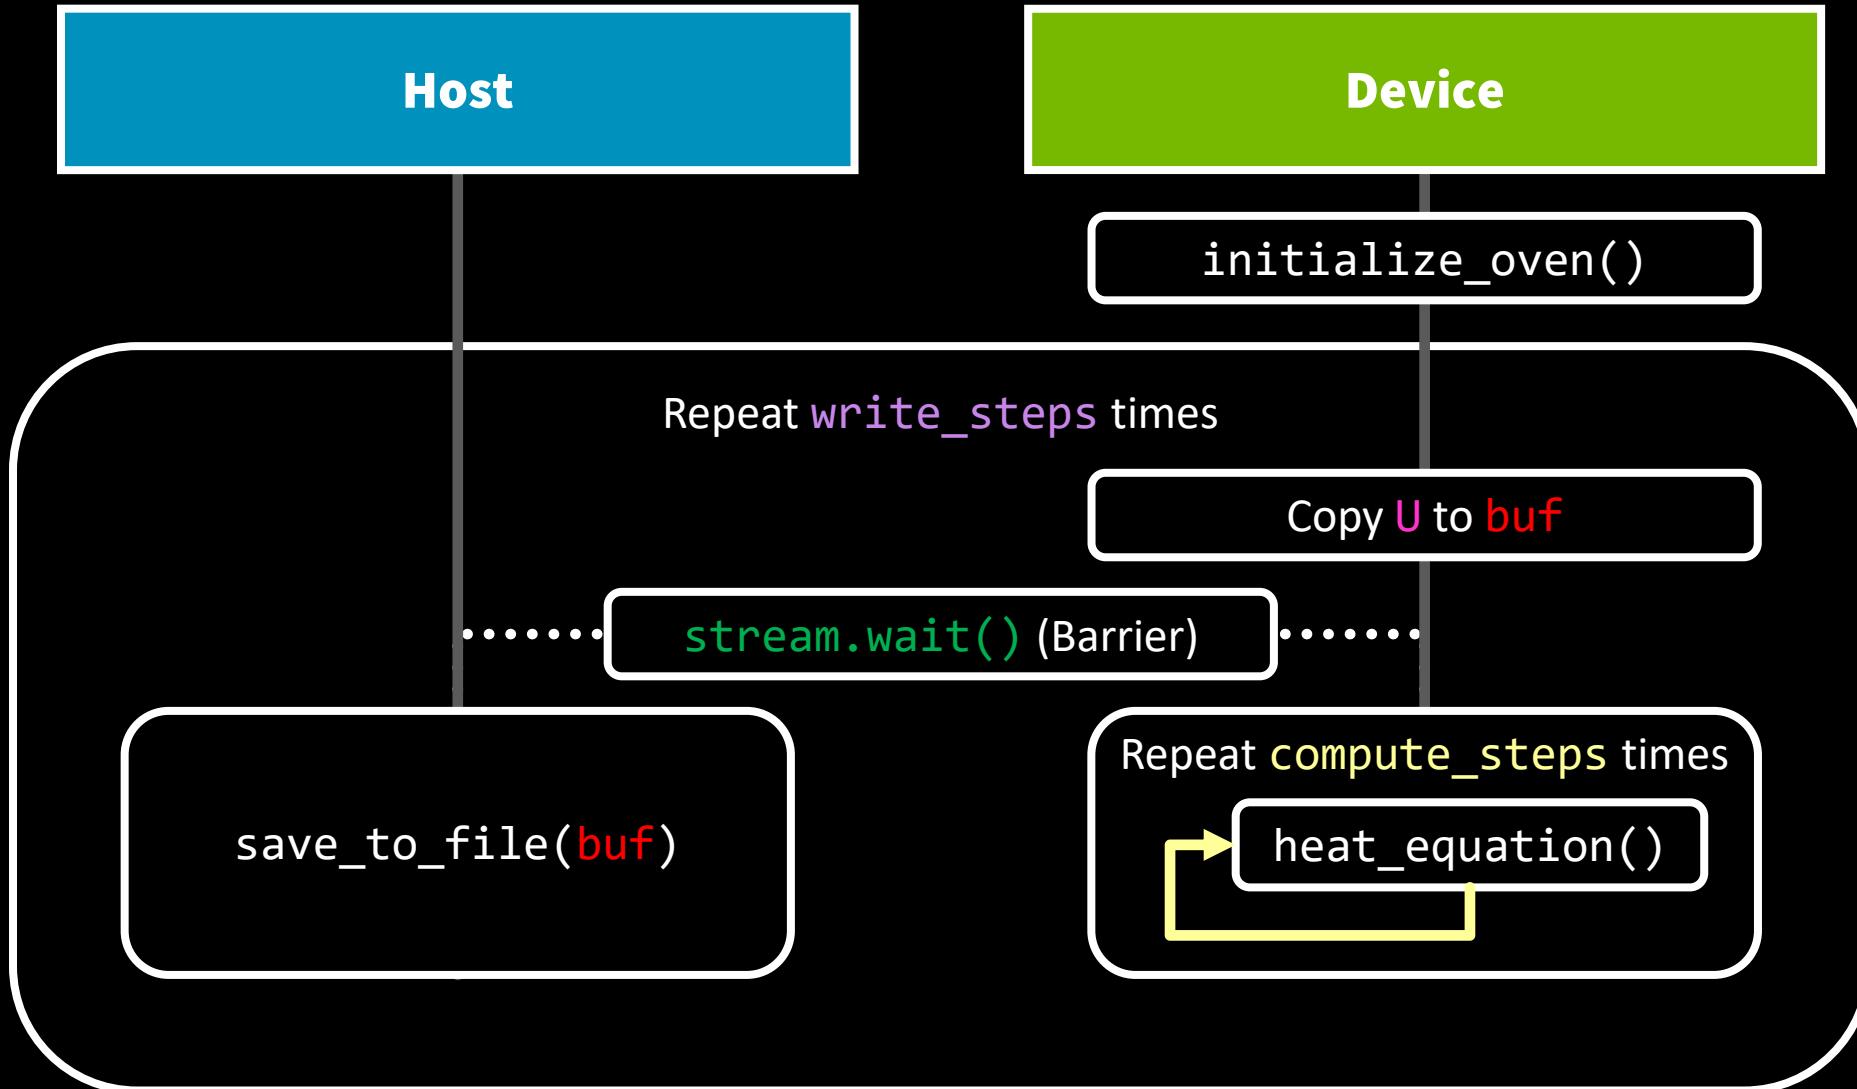
A callout box with a white background and a thin black border, containing the text "initialize\_oven()".

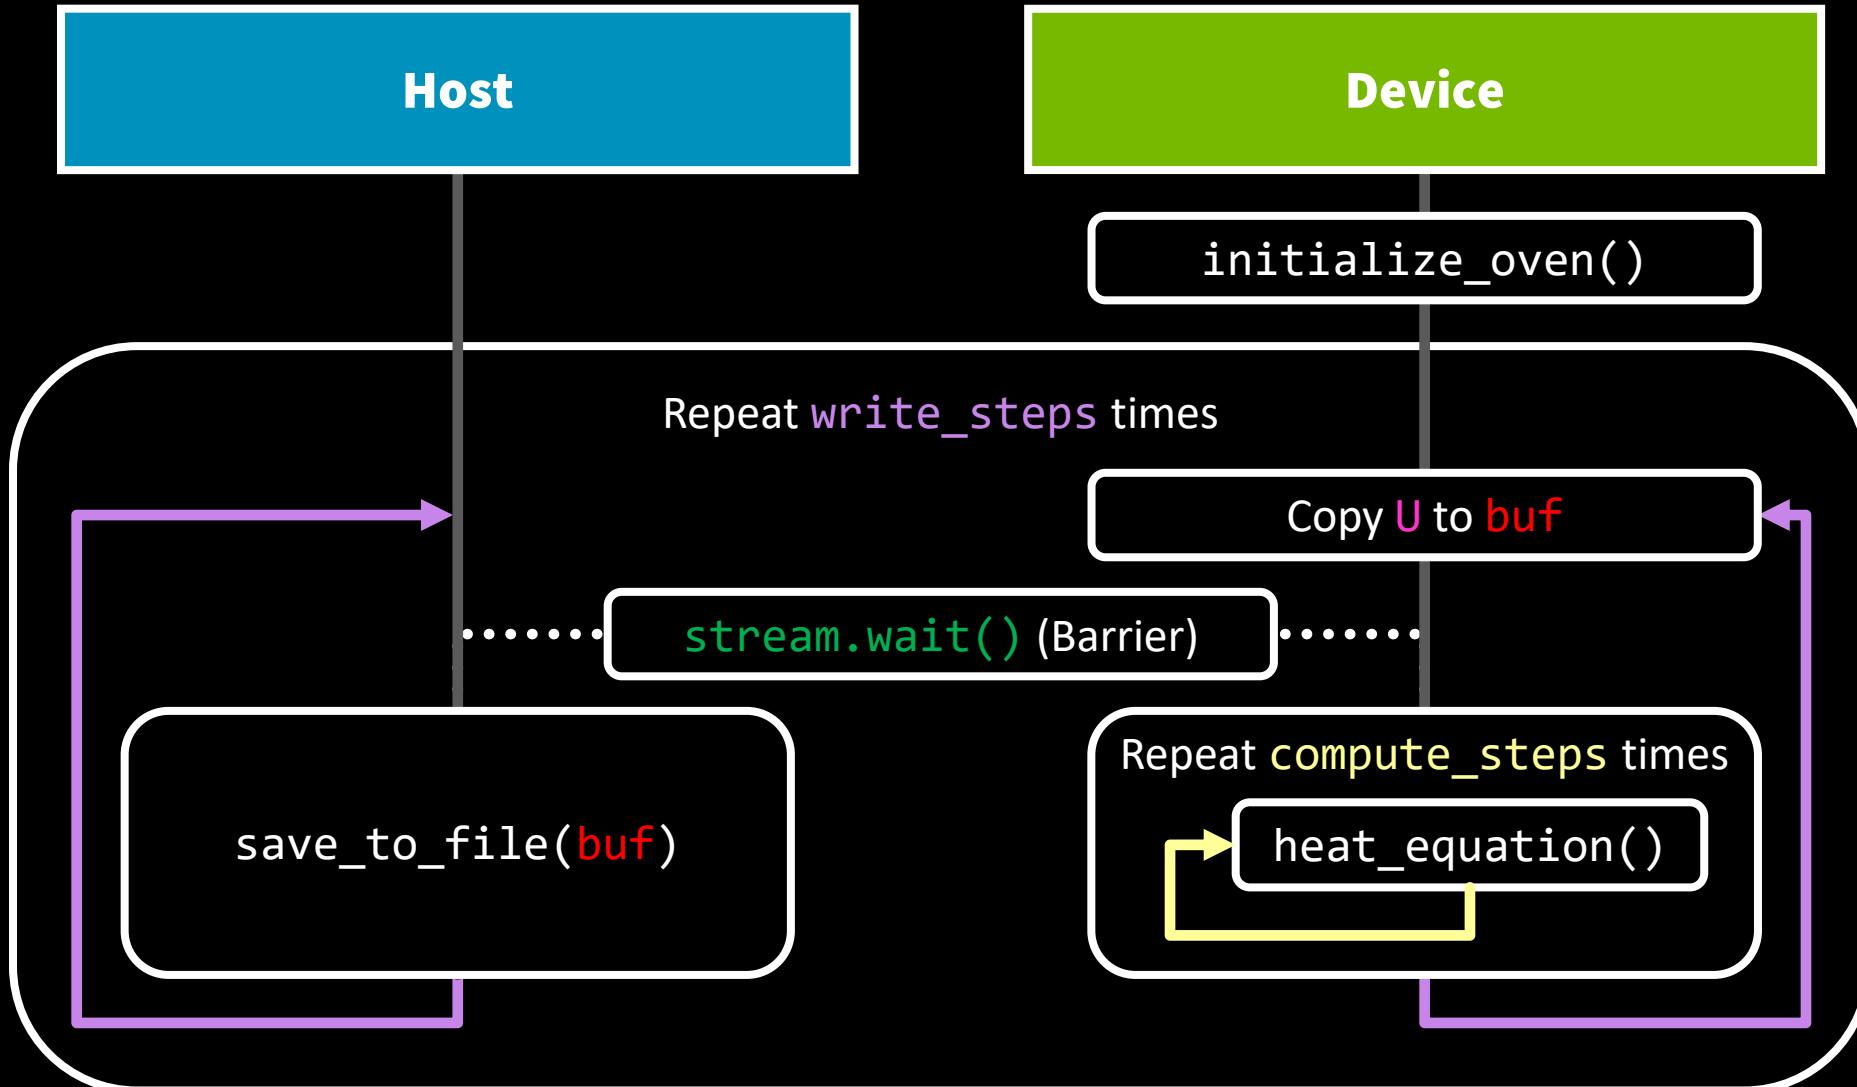






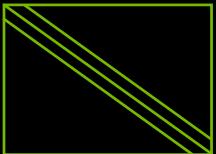








**cuBLAS**



**cuSPARSE**



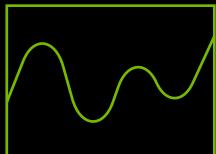
**cuDNN**



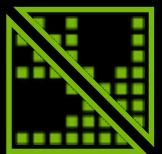
**cuFile  
(GDS)**



**NVBENCH**

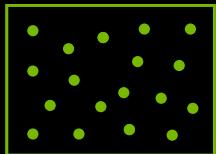


**cuFFT**

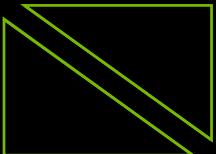


**cuDSS**

**CV  
CUDA**



**cuRAND**

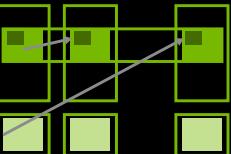


**cuSOLVER**

**RAPIDS**

**cuLitho**

**Cooperative  
Groups**

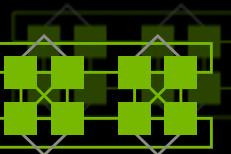


**NVSHMEM**

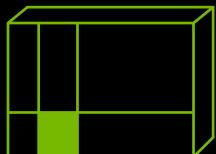


**CUDA-Q**

**CUTLASS**



**NCCL**



**cuTENSOR**



**CUDA Math**



**CUDA  
Runtime**



**GTC CUDA Developer Track**  
[nvda.ws/4ccb7Qh](https://nvda.ws/4ccb7Qh)



**Accelerated Computing Hub**  
[nvda.ws/3QZRuRI](https://nvda.ws/3QZRuRI)

## GENERAL CUDA

- [S72571](#) What's CUDA All About Anyways?
- [S72897](#) How To Write A CUDA Program
- [S72527](#) Debugging & Optimizing CUDA with Intelligent Developer Tools

## CUDA PYTHON

- [S72450](#) Accelerated Python: Tour of Community & Ecosystem
- [S72448](#) The CUDA Python Developer's Toolbox
- [S72449](#) 1001 Ways to Write CUDA Kernels in Python
- [S74639](#) Tensor Core Programming in Python with CUTLASS 4.0

## CUDA C++

- [S72574](#) Building CUDA Software at the Speed-of-Light
- [S72572](#) The CUDA C++ Developer's Toolbox
- [S72575](#) How You Should Write a CUDA C++ Kernel

## MULTI-GPU PROGRAMMING

- [S72576](#) Getting Started with Multi-GPU Scaling: Distributed Libraries
- [S72579](#) Going Deeper with Multi-GPU Scaling: Task-based Runtimes
- [S72578](#) Advanced Multi-GPU Scaling: Communication Libraries

## PERFORMANCE OPTIMIZATION

- [S72683](#) CUDA Techniques to Maximize Memory Bandwidth & Hide Latency
- [S72685](#) CUDA Techniques to Maximize Compute & Instruction Throughput
- [S72686](#) CUDA Techniques to Maximize Concurrency & System Utilization
- [S72687](#) Get the Most Performance From Grace Hopper



**GTC CUDA Developer Track**  
[nvda.ws/4ccb7Qh](https://nvda.ws/4ccb7Qh)



**Accelerated Computing Hub**  
[nvda.ws/3QZRuRI](https://nvda.ws/3QZRuRI)

## GENERAL CUDA

- [S72571](#) What's CUDA All About Anyways?
- [S72897](#) How To Write A CUDA Program
- [S72527](#) Debugging & Optimizing CUDA with Intelligent Developer Tools

## CUDA PYTHON

- [S72450](#) Accelerated Python: Tour of Community & Ecosystem
- [S72448](#) The CUDA Python Developer's Toolbox
- [S72449](#) 1001 Ways to Write CUDA Kernels in Python
- [S74639](#) Tensor Core Programming in Python with CUTLASS 4.0

## CUDA C++

- [S72574](#) Building CUDA Software at the Speed-of-Light
- [S72572](#) The CUDA C++ Developer's Toolbox
- [S72575](#) How You Should Write a CUDA C++ Kernel

## MULTI-GPU PROGRAMMING

- [S72576](#) Getting Started with Multi-GPU Scaling: Distributed Libraries
- [S72579](#) Going Deeper with Multi-GPU Scaling: Task-based Runtimes
- [S72578](#) Advanced Multi-GPU Scaling: Communication Libraries

## PERFORMANCE OPTIMIZATION

- [S72683](#) CUDA Techniques to Maximize Memory Bandwidth & Hide Latency
- [S72685](#) CUDA Techniques to Maximize Compute & Instruction Throughput
- [S72686](#) CUDA Techniques to Maximize Concurrency & System Utilization
- [S72687](#) Get the Most Performance From Grace Hopper