



Lightning Talk: Toying with constexpr and type inference

Xavier Roche

Search Core Team

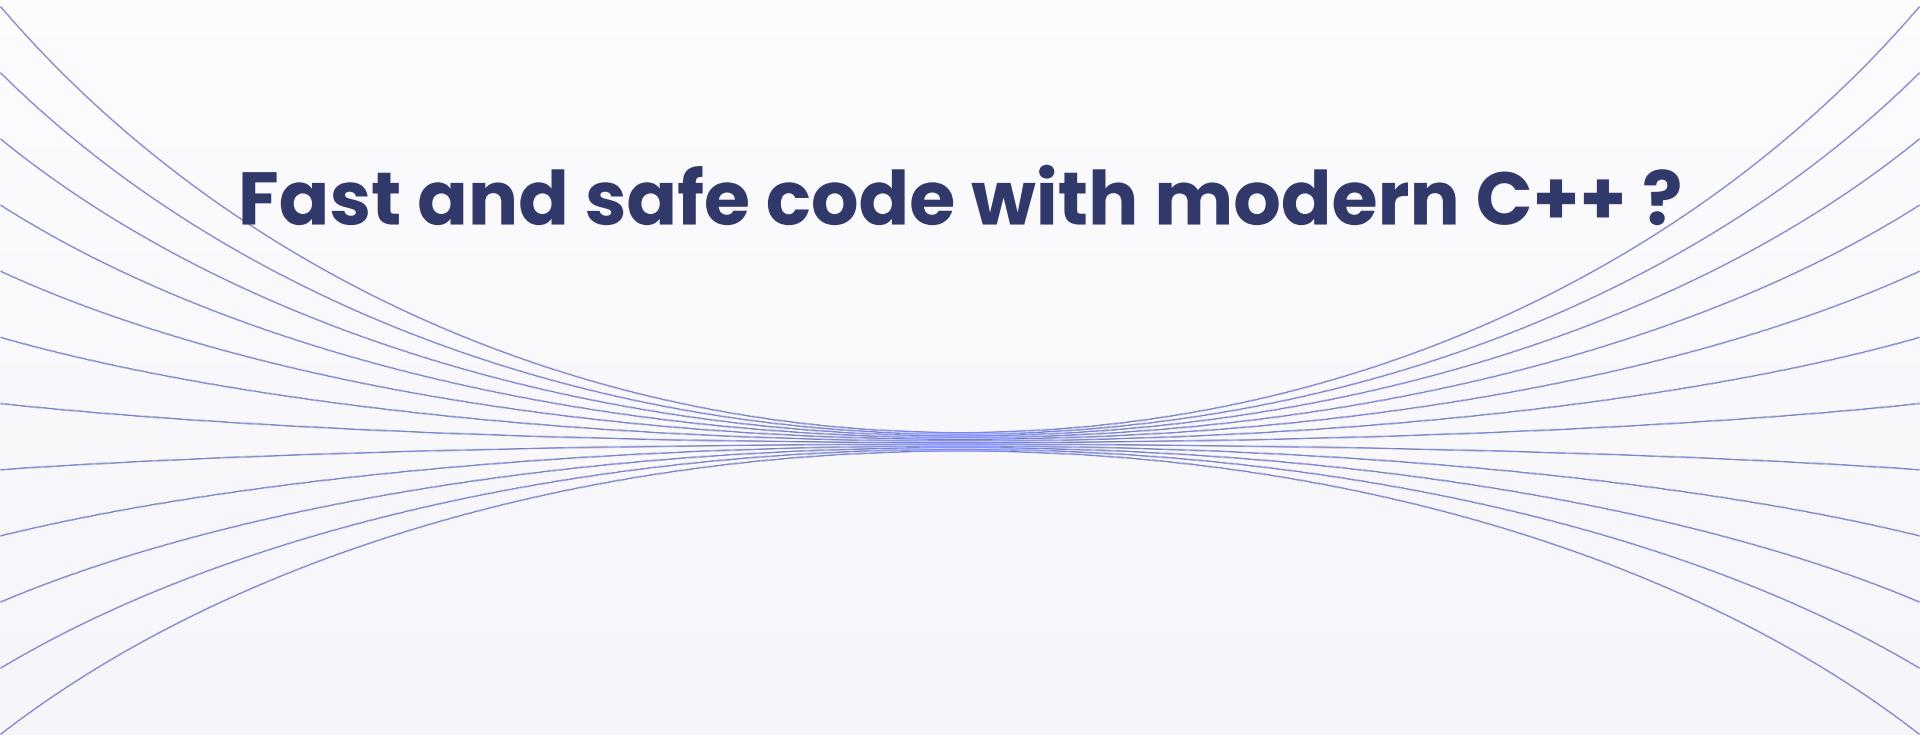


A vertical decorative element on the left side of the slide features a series of thin, light blue lines forming a wavy, organic pattern that tapers towards the top.

Hello!

Xavier Roche

Algolia
Search/Core



Fast and safe code with modern C++ ?



Overview

Toying with `constexpr` & type containers!

1. **switch/case for strings**
2. **runtime function arg validator**

switch/case for strings!

The problem

```
if (key == "poney") {  
} else if (key == "elephant") {  
} else if (key == "dog") {  
} else if (key == "kitten") {
```

- **Linear complexity $O(\text{chars})$**
- **Not so readable**

The problem

```
if (len == 5 && key == "poney") {  
} else if (len == 7 && key == "elephant") {  
} else if (len == 3 && key == "dog") {  
} else if (len == 6 && key == "kitten") {
```

- **Still linear complexity (but better)**
- **Even less readable and bug prone
(did you spot it?)**

What we want

```
switch(key) {  
case "poney":  
    break;  
case "elephant":  
    break;  
case "dog":  
    break;  
case "kitten":  
    break;  
}
```

- **Readable**

- **Performances**

What we can do: hash everything!

```
switch(hash(key)) {  
case "poney"_hash:  
    break;  
case "elephant"_hash:  
    break;  
case "dog"_hash:  
    break;  
case "kitten"_hash:  
    break;  
}
```

- **Readable**

- **Performances
(hopefully)**

What we need: `constexpr fnv-1a/128 hash`

```
algorithm fnv-1 is
    hash := FNV_offset_basis do
        for each byte of data to be hashed
            hash := hash × FNV_prime
            hash := hash XOR byte_of_data
    return hash
```



- “Good enough” (speed/diffusion)
- Collisions extremely unlikely
- Attacks are possible though

What we need: `constexpr` hash

```
template <size_t Bits>
struct fnv1a {
    using Type = typename fnv1a_traits<Bits>::Type;

    /**
     * Compute the Fowler-Noll-Vo hash
     * @param s The string
     * @param l The string size
     * @return The fnv-1a hash
     */
    template <typename C>
    static constexpr Type hash(const C* s, const std::size_t l, Type hash = fnv1a_traits<Bits>::Offset)
    {
        // See <https://en.wikipedia.org/wiki/Fowler%20%93Noll%20%93Vo\_hash\_function>
        for (std::size_t j = 0; j < l; j++) {
            const uint8_t byte = s[j];
            hash ^= byte;
            hash *= fnv1a_traits<Bits>::Prime;
        }
        return hash;
    }
}
```

What we need: traits for different hash sizes

```
template <size_t Bits>
struct fnv1a_traits { };

// Traits for 64-bit FNV1a
template <>
struct fnv1a_traits<64> {
    using Type = uint64_t;
    static constexpr Type Prime = 0x100000001b3;
    static constexpr Type Offset = 0xcbf29ce484222325;
};

// Traits for 128-bit FNV1a
template <>
struct fnv1a_traits<128> {
    using Type = __uint128_t; ← 128-bit integer (no std::uint128_t though)
    static constexpr Type Prime = Pack128(0x1000000, 0x00000000000000013b);
    static constexpr Type Offset = Pack128(0x6c62272e07bb0142, 0x62b821756295c58d);
};
```

What we need: custom string literal

```
using fnv1a128 = fnv1a<128>;\n\nconstexpr auto operator"" _hash(const char* s, const std::size_t l)\n{\n    return fnv1a128::hash(s, l);\n}\n\n// Static unit tests: <https://fnvhash.github.io/fnv-calculator-online/>\nstatic_assert("hello"_hash == Pack128(0xe3e1efd54283d94f, 0x7081314b599d31b3));
```

Cosmetic, but nice!

Hash & dispatch example

```
constexpr auto hash(const char* s, const std::size_t l)          const char* dispatch(const fnv1a_traits<128>::Type hash)
{
    return fnv1a128::hash(s, l);
}                                                               {  
    switch (hash) {  
        case "poney"_hash:  
            return "I want one, too!";  
        case "elephant"_hash:  
            return "Not in my apartment please!";  
        case "dog"_hash:  
            return "Good puppy!";  
        case "kitten"_hash:  
            return "Aawwwwww! ";  
        default:  
            return "Don't know this animal!";
    }
}
```

Dispatch example : hash (extract)

```

label3:
movzbl (%rdi,%r10,1),%r8d
xor %rax,%r8
mov %r8,%rax
mul %r11
shl $0x18,%r8      ←
add %rdx,%r8
imul $0x13b,%rcx,%rbx ← prime.1
add %r8,%rbx
movzbl 0x1(%rdi,%r10,1),%ecx
xor %rax,%rcx
mov %rcx,%rax
mul %r11
shl $0x18,%rcx      ←prime.2
add %rdx,%rcx
imul $0x13b,%rbx,%rbx ←
add %rcx,%rbx

```

```

movzbl 0x2(%rdi,%r10,1),%ecx
xor %rax,%rcx
mov %rcx,%rax
mul %r11
→ shl $0x18,%rcx
add %rdx,%rcx
→ imul $0x13b,%rbx,%rdx
add %rcx,%rdx
movzbl 0x3(%rdi,%r10,1),%ecx
xor %rax,%rcx
→ imul $0x13b,%rdx,%rbx
add %rcx,%rax
mul %r11
→ shl $0x18,%rcx
add %rdx,%rcx
add %rbx,%rcx
add $0x4,%r10
...

```

Dispatch example : binary search (extract)

```
movabs $0x6efcb17ab10f2a3d,%rax # "kitten"_fnv1a128.1
cmp    %rdi,%rax
```

```
movabs $0xfcfd05704e13c64bf,%rax # "kitten"_fnv1a128.2
movq    %rdi,%xmm0
movq    %rsi,%xmm1
punpcklqdq %xmm1,%xmm0
sbb     %rsi,%rax
jl     0x400bba <test_kitten>
```

```
movdqa 0x17bce(%rip),%xmm1      # "dog"_fnv1a128
pcmpeqb %xmm0,%xmm1
pmovmskb %xmm1,%eax
cmp    $0xffff,%eax
je     0x400bf0 <puppy>
```

```
pcmpeqb 0x17bc7(%rip),%xmm0      # "poney"_fnv1a128
pmovmskb %xmm0,%eax
```

```
cmp    $0xffff,%eax
jne   0x400bea <unknown>
```

```
mov    $0x41c6a0,%eax          # "I want one, too!"
retq
```

test_kitten:

```
movdqa 0x17b7e(%rip),%xmm1      # "kitten"_fnv1a128
pcmpeqb %xmm0,%xmm1
pmovmskb %xmm1,%eax
cmp    $0xffff,%eax
je     0x400bf6 <kitten>
```

...

Want to have a look ?

<https://github.com/xroche/stringswitch>

runtime function arg validator

The problem: size/offset overflow

```
char foo_write(const void*, char size, Foo*);
```

- Function taking non-standard sizes (ie. not std::size_t)
- Real-world is using standard sizes

The problem: size/offset overflow

```
char foo_write(const void*, char size, Foo*);  
...  
const char str[142] = "Hello!\n";  
std::size_t size = 142;  
...  
const char b = foo_write(str, size, foo);
```

Oops (trivial case here)

What we can do

```
char foo_write(const void*, char size, Foo*);  
...  
const char str[142] = "Hello!\n";  
std::size_t size = 142;  
...  
assert(size <= std::numeric_limits<char>::max());  
const char b = foo_write(str, size, foo);
```

Cumbersome

What we want: automatic cast assert

```
char foo_write(const void*, char size, Foo*);  
...  
const char b =  
    checked_cast_call<foo_write>(str, size, foo);
```

I am lazy:

→ **sed -e 's/foo_write/checked_cast_call<foo_write>/g'**

Step 1: cast

```
/**  
 * Template numerical conversion.  
 * @param v The source numerical value  
 * @return The destination numerical value, or the original pointer for non-numerical types  
 * @comment Inspired by Bjarne Stroustrup's in "The C++ Programming Language 4th Edition"  
**/  
  
template<class Target, class Source>  
auto inline checked_cast(Source v)  
{  
    if constexpr (std::is_same<Target, Source>::value) {  
        return v;  
    } else if constexpr (!std::is_integral<Source>::value) {  
        return v;  
    } else {  
        const auto r = static_cast<Target>(v);  
        if (static_cast<Source>(r) != v)  
            throw std::runtime_error(std::string(__PRETTY_FUNCTION__));  
        return r;  
    }  
}
```

Step 1

```
char foo_write(const void*, char size, Foo*);  
...  
const char b =  
    foo_write(str, checked_cast<char>(size), foo);
```

Need to validate each argument type

Step 2: argument type inference

```
template<typename T>
class checked_cast_call_container
{
public:
    inline constexpr checked_cast_call_container(T result)
        : _result(result)
    {}
}
```

```
template<typename U>
inline constexpr operator U() const
{
    return checked_cast<U>(_result);
}

private:
    const T _result;
};
```

- T is the “source”
- U is the “target”

← Types inferred

Step 2: argument type inference

```
char foo_write(const void*, char size, Foo*);  
...  
const char b =  
    foo_write(str, checked_cast_call_container(size),  
              foo);
```

Need to decorate each argument

Step 3: templated function caller

```
/**  
 * Wrapped call to a function, with runtime-checked casted input and output values.  
 * @example checked_cast_call<my_function>(str, 1, size, output)  
 * @comment This version is using templated function pointer.  
 */  
template<auto fn, typename... Args>  
auto constexpr checked_cast_call(Args... args)  
{  
    return checked_cast_call_container(fn(checked_cast_call_container(args)...));  
}
```

← Non-type template
C++17
(template <typename T, T fn>) ← Fold expression C++17

We also return a container to infer target

Step 3: templated function caller

```
char foo_write(const void*, char size, Foo*);  
...  
const char b =  
    checked_cast_call<foo_write>(str, size, foo);
```

Here we are!

Want to have a look ?

<https://github.com/xroche/checkedcast>

<https://github.com/xroche/stringswitch>
<https://github.com/xroche/checkedcast>

Questions ?



Xavier Roche

Search Core Team

